

# CSE 351 Spring 2019

# Ruth Anderson

# Gavin Cai

# Jack Eggleston

# John Feltrup

Britt Henderson

# Richard Jiang

# Jack Skalitzky

Sophie Tian

# Connie Wang

# Sam Wolfson

# Casey Xing

# Chin Yeoh



# Administrivia

- ❖ Lab 0 due TODAY @ 11:59 pm
  - *You will be revisiting this program throughout this class!*
- ❖ Homework 1 due Wednesday
  - Reminder: autograded, 20 tries, no late submissions
- ❖ Lab 1a released
  - Workflow:
    - 1) Edit `pointer.c`
    - 2) Run the Makefile (`make`) and check for compiler errors & warnings
    - 3) Run `pctest (. /pctest)` and check for correct behavior
    - 4) Run rule/syntax checker (`python dlc.py`) and check output
  - Due Monday 4/15, will overlap a bit with Lab 1b
    - We grade just your *last* submission

# Lab Reflections

- ❖ All subsequent labs (after Lab 0) have a “reflection” portion
  - The Reflection questions can be found on the lab specs and are intended to be done *after* you finish the lab
  - You will type up your responses in a `.txt` file for submission on Canvas
  - These will be graded “by hand” (read by TAs)
- ❖ Intended to check your understand of what you should have learned from the lab
  - Also great practice for short answer questions on the exams

# Memory, Data, and Addressing

- ❖ Hardware - High Level Overview
- ❖ Representing information as bits and bytes
  - Memory is a byte-addressable array
  - Machine “word” size = address size = register size
- ❖ Organizing and addressing data in memory
  - Endianness – ordering bytes in memory
- ❖ Manipulating data in memory using C
- ❖ **Boolean algebra and bit-level manipulations**

# Boolean Algebra

- ❖ Developed by George Boole in 19th Century
  - Algebraic representation of logic (True  $\rightarrow$  1, False  $\rightarrow$  0)
  - AND:  $A \& B = 1$  when both A is 1 and B is 1
  - OR:  $A | B = 1$  when either A is 1 or B is 1
  - XOR:  $A \wedge B = 1$  when either A is 1 or B is 1, but not both
  - NOT:  $\sim A = 1$  when A is 0 and vice-versa
  - DeMorgan's Law:
 
$$\sim (A | B) = \sim A \& \sim B$$

$$\sim (A \& B) = \sim A | \sim B$$

AND			OR			XOR			NOT		
$\&$	0	1		0	1	$\wedge$	0	1	$\sim$		
0	0	0	0	0	1	0	0	1	0	1	
1	0	1	1	1	1	1	1	0	1	0	

# General Boolean Algebras

## ❖ Operate on bit vectors

- Operations applied bitwise
- All of the properties of Boolean algebra apply

$$\begin{array}{cccc}
 \begin{array}{r} 01101001 \\ \& 01010101 \\ \hline 01000001 \end{array} & 
 \begin{array}{r} 01101001 \\ | 01010101 \\ \hline 01111101 \end{array} & 
 \begin{array}{r} 01101001 \\ \wedge 01010101 \\ \hline 00111100 \end{array} & 
 \begin{array}{r} 01101001 \\ \sim 01010101 \\ \hline 10101010 \end{array}
 \end{array}$$

## ❖ Examples of useful operations:

$$x \wedge x = 0$$

"sets to 1"

$$\begin{array}{l}
 x | 1 = 1, \\
 0 | 1 = 1 \\
 1 | 1 = 1
 \end{array}$$

"leaves as is"

$$\begin{array}{l}
 x | 0 = x \\
 0 | 0 = 0 \\
 1 | 0 = 1
 \end{array}$$

$$\begin{array}{r}
 01010101 \\
 \wedge 01010101 \\
 \hline
 00000000
 \end{array}$$

← creates 0

$$\begin{array}{r}
 01010101 \\
 | 11110000 \\
 \hline
 11110101
 \end{array}$$

← data of interest  
← bit mask (specifically chosen)

set left as is

# Bit-Level Operations in C

bitwise & is the same as an  
bitwise "product"

## ❖ & (AND), | (OR), ^ (XOR), ~ (NOT)

- View arguments as bit vectors, apply operations bitwise

- Apply to any "integral" data type

(8 bytes) (4 bytes) (2 bytes) (1 byte)  
• long, int, short, char, unsigned

bit vector will be  
width of datatype

## ❖ Examples with char a, b, c;

	<u>C code</u>		<u>Internally</u>	<u>Result</u>
■	a = (char) 0x41;	//	0x41 → 0b 0100 0001	
	b = ~a;	//	0b 1011 1110 → 0x BE	
■	a = (char) 0x69;	//	0x69 → 0b 0110 1001	
	b = (char) 0x55;	//	0x55 → 0b 0101 0101	
	c = a & b;	//	0b 0100 0001 → 0x 41	
■	a = (char) 0x41;	//	0x41 → 0b 0100 0001	
	b = a;	//	0b 0100 0001	
	c = a ^ b;	//	0b 0000 0000 → 0x 00	

# Contrast: Logic Operations

❖ Logical operators in C: `&&` (AND), `||` (OR), `!` (NOT)

■ **0** is False, **anything nonzero** is True

■ **Always** return 0 or 1

$0xCC = 0b1100\ 1100$   
 $0x33 = 0b0011\ 0011$

■ **Early termination** (a.k.a. short-circuit evaluation) of `&&`, `||`

❖ Examples (char data type)  $0xCC \ \& \ 0x33 \ \rightarrow \ 0x00$

■  $!0x41 \xrightarrow{T} \xrightarrow{F} 0x00$

■  $0xCC \xrightarrow{T} \&\& \ 0x33 \xrightarrow{T} \rightarrow 0x01$

■  $!0x00 \xrightarrow{F} \xrightarrow{T} 0x01$

■  $0x00 \xrightarrow{F} || \ 0x33 \xrightarrow{T} \rightarrow 0x01$

■  $!(0x41) \xrightarrow{T} \rightarrow 0x01$

■  $\textcircled{1} \ p \ \&\& \ * \ \textcircled{2} \ p$

- If  $p$  is the **null pointer** (0x0), then  $p$  is never dereferenced! normally trying to dereference a null pointer would cause an error

If  $\textcircled{1}$  determines output of logical operator, then  $\textcircled{2}$  is never evaluated



# Roadmap

C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

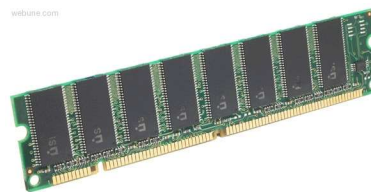
Assembly  
language:

```
get_mpg:  
    pushq    %rbp  
    movq     %rsp, %rbp  
    ...  
    popq     %rbp  
    ret
```

Machine  
code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
110000011111101000011111
```

Computer  
system:



Memory & data

**Integers & floats**

x86 assembly

Procedures & stacks

Executables

Arrays & structs

Memory & caches

Processes

Virtual memory

Memory allocation

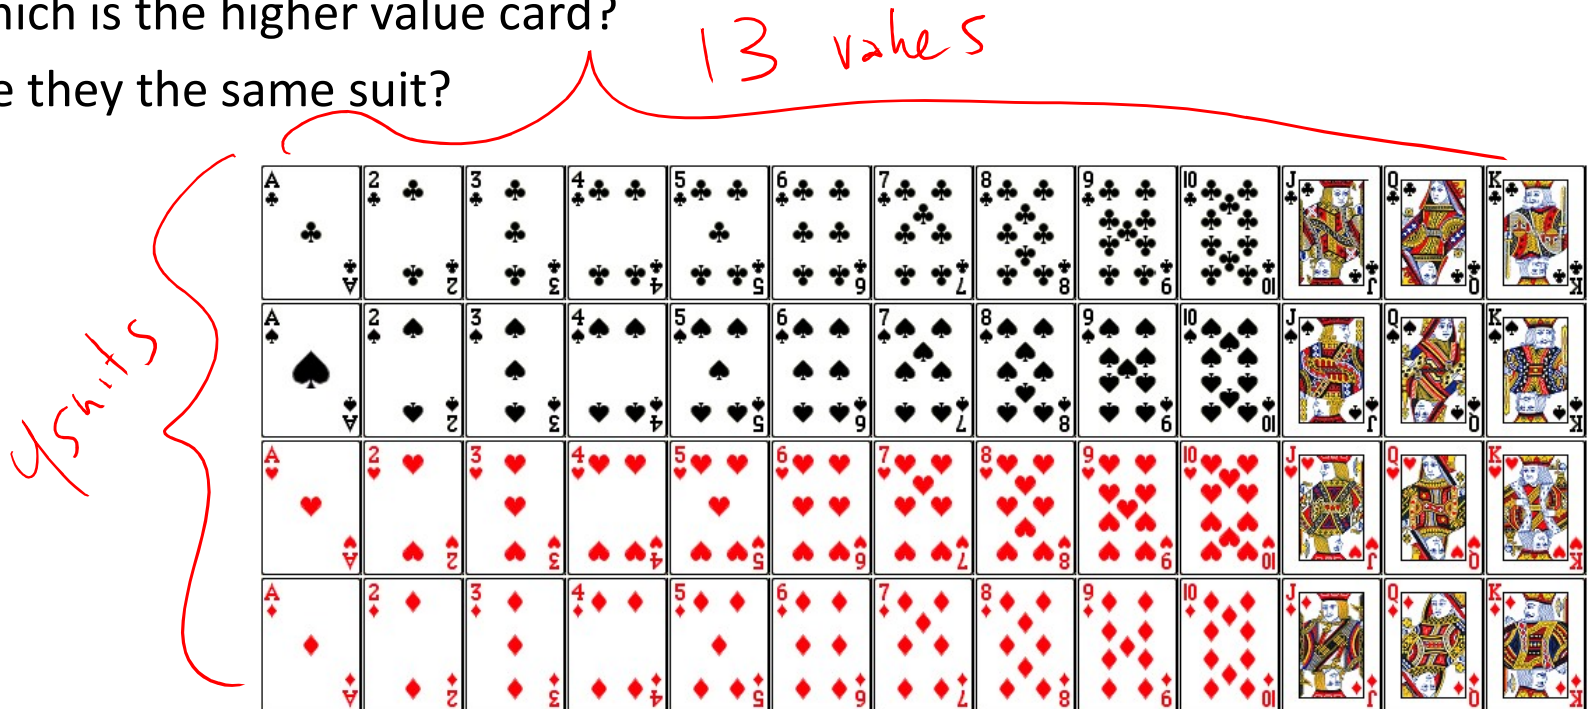
Java vs. C

OS:

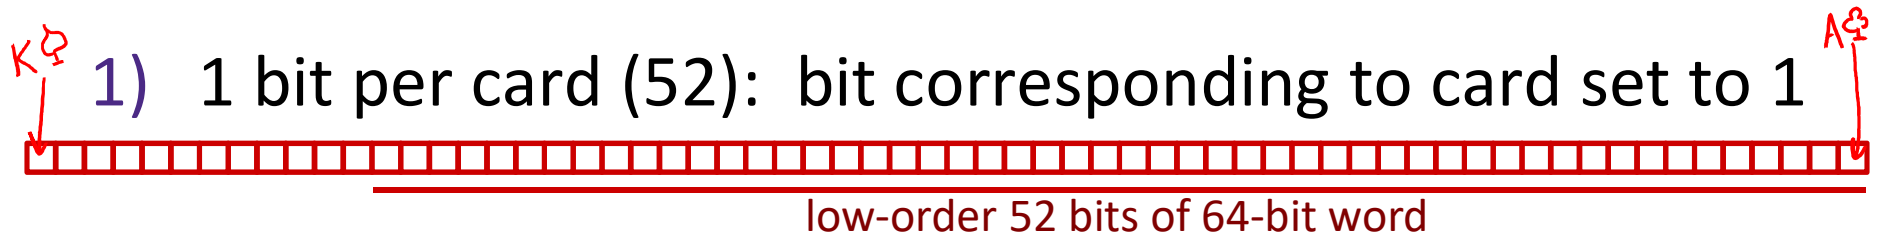


# But before we get to integers....

- ❖ Encode a standard deck of playing cards
- ❖ 52 cards in 4 suits
  - How do we encode suits, face cards?
- ❖ What operations do we want to make easy to implement?
  - Which is the higher value card?
  - Are they the same suit?



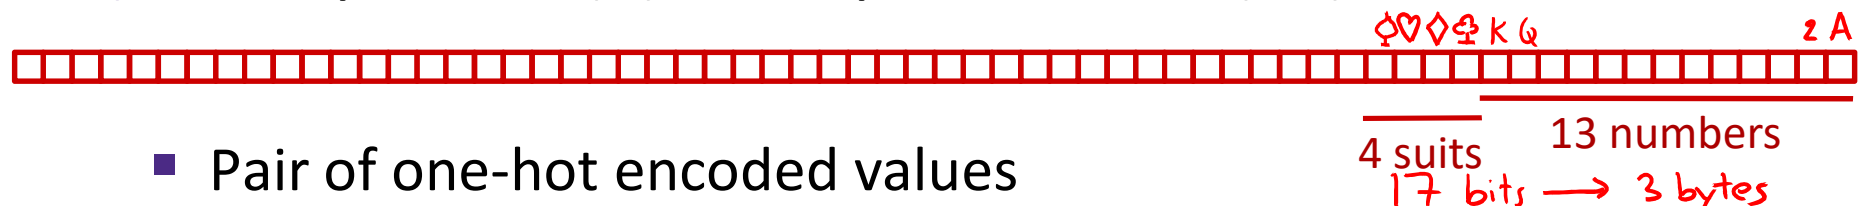
# Two possible representations



- “One-hot” encoding (similar to set notation)
- Drawbacks:

- Hard to compare values and suits
  - Large number of bits required
- 52 bits  $\xrightarrow{\text{fits in}}$  7 bytes (56 bits)

2) 1 bit per suit (4), 1 bit per number (13): 2 bits set



- Pair of one-hot encoded values
- Easier to compare suits and values, but still lots of bits used

# Two better representations

## 3) Binary encoding of all 52 cards – only 6 bits needed

- $2^6 = 64 \geq 52$   
 $2^5 = 32 < 52$



low-order 6 bits of a byte

- Fits in one byte (smaller than one-hot encodings)
- How can we make value and suit comparisons easier?

## 4) Separate binary encodings of suit (2 bits) and value (4 bits)



suit

value

- Also fits in one byte, and easy to do comparisons

K	Q	J	...	3	2	A
1101	1100	1011	...	0011	0010	0001

13

...

1

C	♣	00
D	♦	01
H	♥	10
S	♠	11

12

e.g. a filter

# Compare Card Suits

**mask:** a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector  $v$ .

Here we turn all *but* the bits of interest in  $v$  to 0.

```
char hand[5];           // represents a 5-card hand
char card1, card2;      // two cards to compare
card1 = hand[0];
card2 = hand[1];
...
if ( sameSuitP(card1, card2) ) { ... }
```

text substitution

```
#define SUIT_MASK 0x30
```

```
int sameSuitP(char card1, char card2) {
    return (!((card1 & SUIT_MASK) ^ (card2 & SUIT_MASK)));
    //return (card1 & SUIT_MASK) == (card2 & SUIT_MASK);
}
```

returns int

SUIT\_MASK = 0x30 =

0	0	1	1	0	0	0	0
---	---	---	---	---	---	---	---

 $x \& 0 = 0$  $x \& 1 = x$ 

suit  
(keep)

value  
(discard)

equivalent

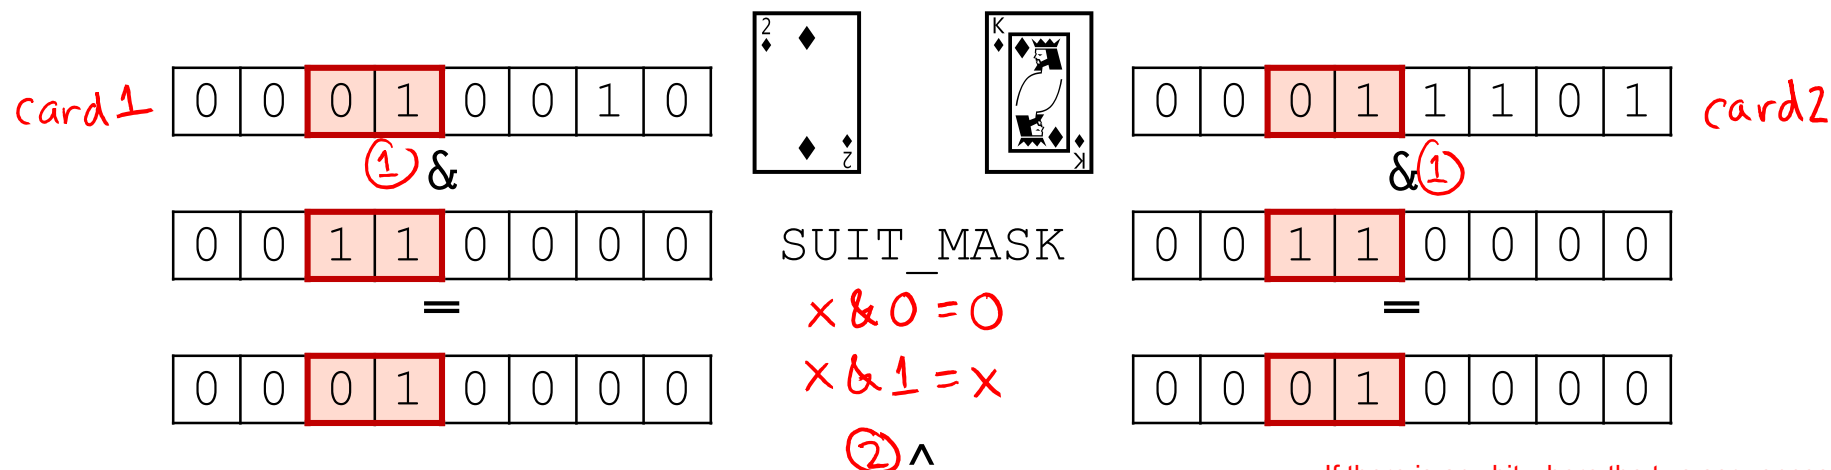
# Compare Card Suits

**mask:** a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector  $v$ .

Here we turn all *but* the bits of interest in  $v$  to 0.

```
#define SUIT_MASK 0x30
```

```
int sameSuitP(char card1, char card2) {
    return (!((card1 & SUIT_MASK) ^ (card2 & SUIT_MASK)));
    //return (card1 & SUIT_MASK) == (card2 & SUIT_MASK);
}
```



If there is any bit where the two sequences don't agree, the XOR will produce a one in that bit. Hence the XOR is identically zero IFF the two bit sequences are the same.

! (x ^ y) equivalent to x == y

③! ← logical

# Compare Card Values

**mask:** a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector  $v$ .

```
char hand[5];           // represents a 5-card hand
char card1, card2;      // two cards to compare
card1 = hand[0];
card2 = hand[1];

...

if ( greaterValue(card1, card2) ) { ... }
```

```
#define VALUE_MASK 0x0F

int greaterValue(char card1, char card2) {
    return ((unsigned int)(card1 & VALUE_MASK) >
            (unsigned int)(card2 & VALUE_MASK));
}
```

casting to an unsigned int will use the two-encoding formula  $\sum x_i 2^i$  to convert to a nonnegative integer

VALUE\_MASK = 0x0F =

0	0	0	0	1	1	1	1
---	---	---	---	---	---	---	---

suit  
 (discard)

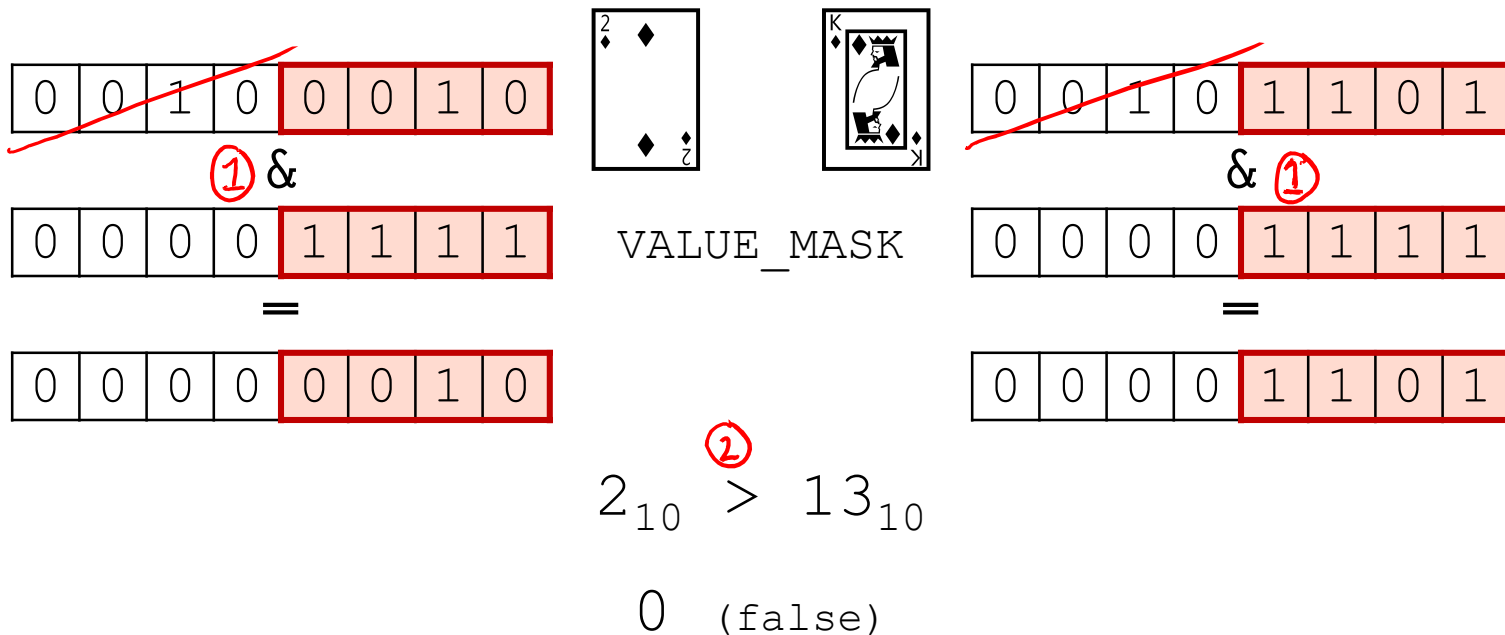
value  
 (keep)

# Compare Card Values

**mask:** a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector  $v$ .

```
#define VALUE_MASK 0x0F

int greaterValue(char card1, char card2) {
    return ((unsigned int) (card1① & VALUE_MASK) ②
           (unsigned int) (card2① & VALUE_MASK));
}
```





# Integers

- ❖ **Binary representation of integers**
  - **Unsigned and signed**
  - Casting in C
- ❖ Consequences of finite width representation
  - Overflow, sign extension
- ❖ Shifting and arithmetic operations

# Encoding Integers

- ❖ The hardware (and C) supports two flavors of integers
  - *unsigned* – only the non-negatives
  - *signed* – both negatives and non-negatives

- ❖ Cannot represent all integers with  $w$  bits

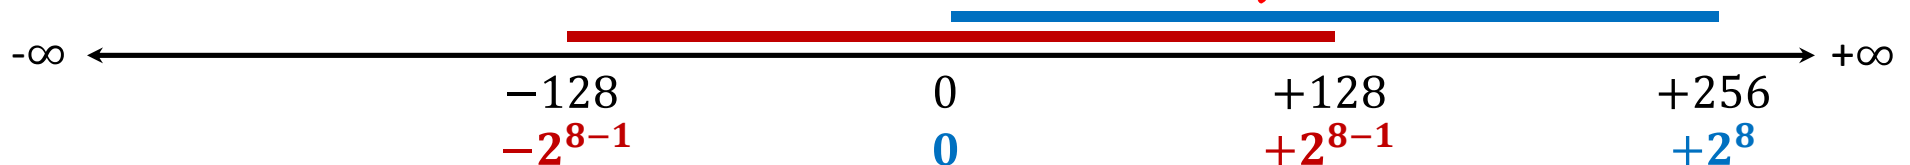
- Only  $2^w$  distinct bit patterns

- Unsigned values:  $0 \dots 2^w - 1$

- Signed values:  $-2^{w-1} \dots 2^{w-1} - 1$

$w \rightarrow 8 \text{ bits}$   
 $0 \dots 255$   
 $-128 \dots 127$   
 same widths, just shifted

- ❖ **Example:** 8-bit integers (e.g. char)



# Unsigned Integers

- ❖ Unsigned values follow the standard base 2 system

- $b_7b_6b_5b_4b_3b_2b_1b_0 = b_72^7 + b_62^6 + \dots + b_12^1 + b_02^0$

- ❖ Add and subtract using the normal “carry” and “borrow” rules, just in binary

$$\begin{array}{r} 63 \\ + 8 \\ \hline 71 \end{array}$$

$$\begin{array}{r} 00111111 \\ + 00001000 \\ \hline 01000111 \end{array}$$

← 6 1's in a row

This “adding binary” only makes sense in the interpretation of the binary numbers as unsigned integers. Here, the addition is compatible with the above encoding. So instead of converting both the binary numbers to unsigned integers, finding the sum, and converting it back this gives a convenient way to just add the binary numbers

- ❖ Useful formula:  $2^{N-1} + 2^{N-2} + \dots + 2 + 1 = 2^N - 1$

- i.e. N ones in a row  $= 2^N - 1$

- ❖ How would you make *signed* integers?

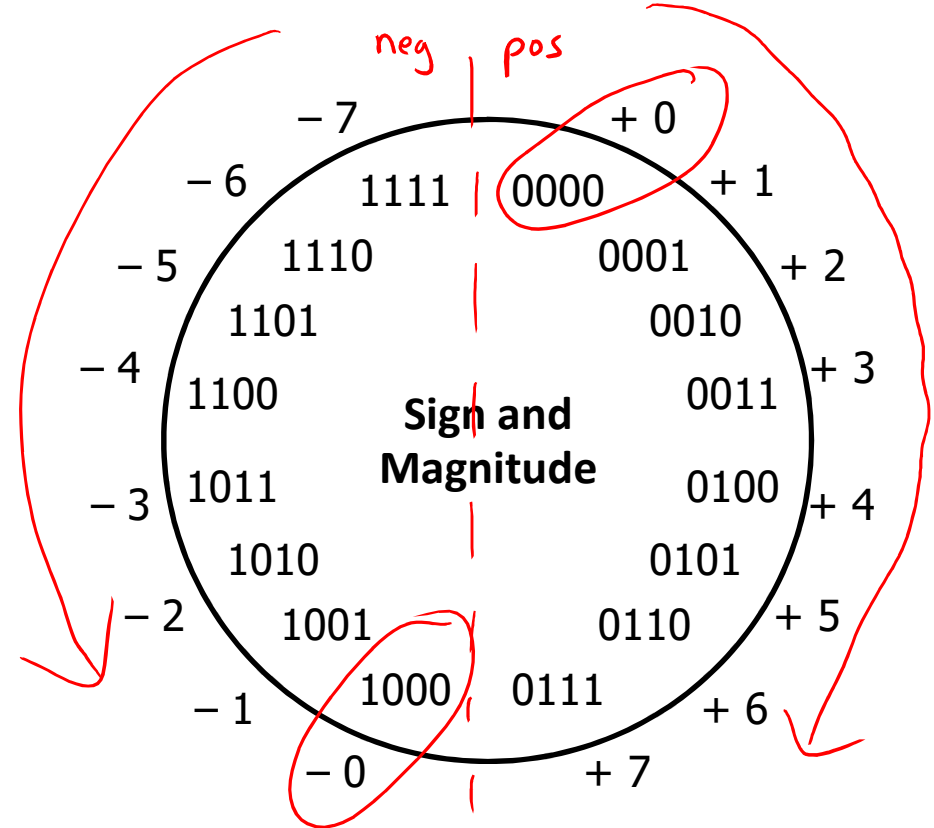
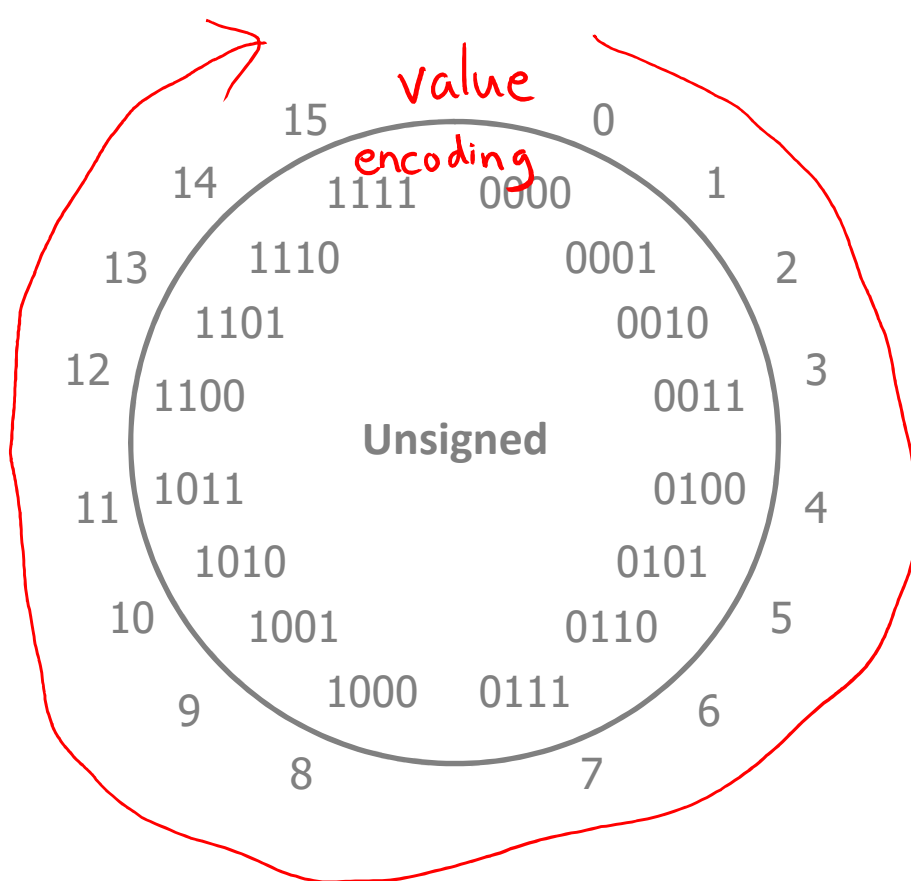
# Sign and Magnitude

Most Significant Bit

- ❖ Designate the high-order bit (MSB) as the “sign bit”
  - $\text{sign}=0$ : positive numbers;  $\text{sign}=1$ : negative numbers
- ❖ Benefits:
  - Using MSB as sign bit matches positive numbers with unsigned *unsigned:  $0b\ 0010 = 2^1 = 2$  ; sign + mag:  $0b\ 0010 = +2^1 = 2$  ✓*
  - All zeros encoding is still  $= 0$
- ❖ Examples (8 bits):
  - ✓ ■  $0x00 = \underline{00000000}_2$  is non-negative, because the sign bit is 0
  - $0x7F = \underline{01111111}_2$  is non-negative ( $+127_{10}$ )
  - $0x85 = 10000101_2$  is negative ( $-5_{10}$ )
  - $0x80 = 10000000_2$  is negative... zero???

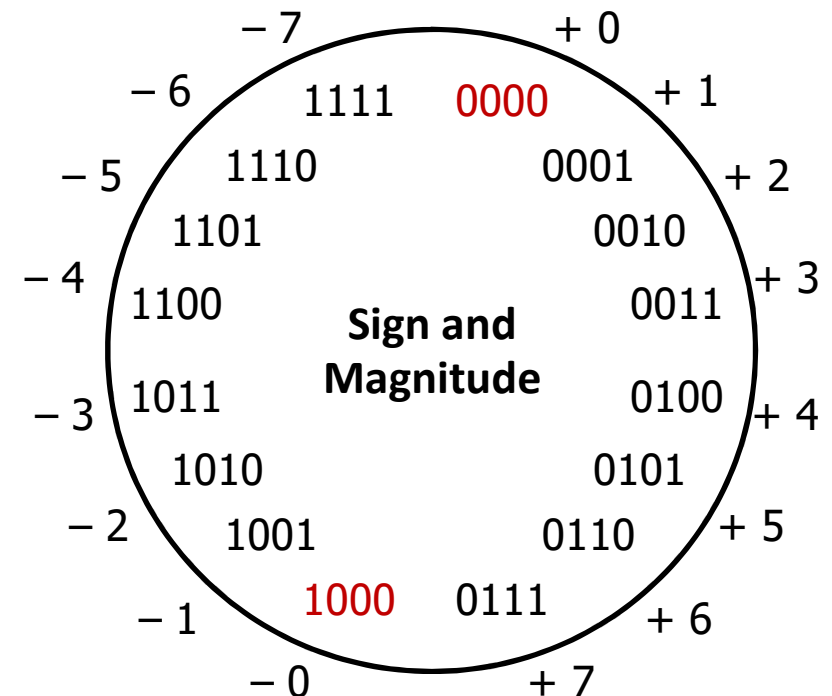
# Sign and Magnitude

- ❖ MSB is the sign bit, rest of the bits are magnitude
- ❖ Drawbacks?



# Sign and Magnitude

- ❖ MSB is the sign bit, rest of the bits are magnitude
- ❖ Drawbacks:
  - Two representations of 0 (bad for checking equality)



# Sign and Magnitude

❖ MSB is the sign bit, rest of the bits are magnitude

❖ Drawbacks:

■ Two representations of 0 (bad for checking equality)

■ Arithmetic is cumbersome

• Example:  $4 - 3 \neq 4 + (-3)$

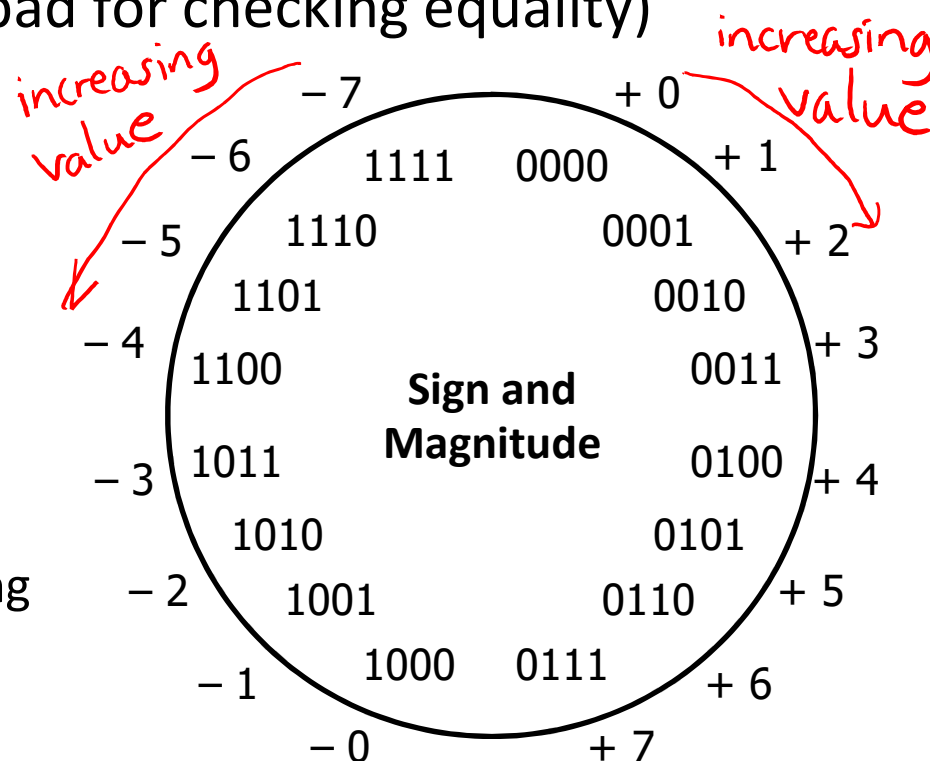
4	0100
- 3	- 0011
1	0001



4	0100
+ -3	+ 1011
-7	1111



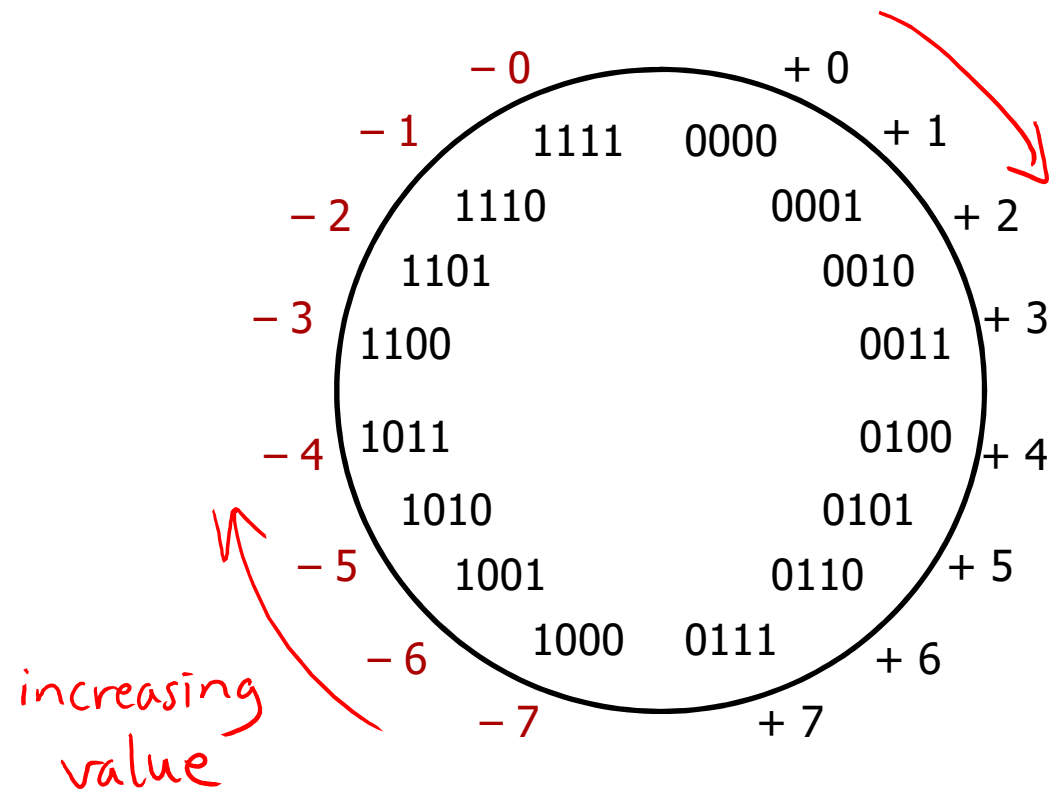
• Negatives “increment” in wrong direction!



# Two's Complement

❖ Let's fix these problems:

1) "Flip" negative encodings so incrementing works





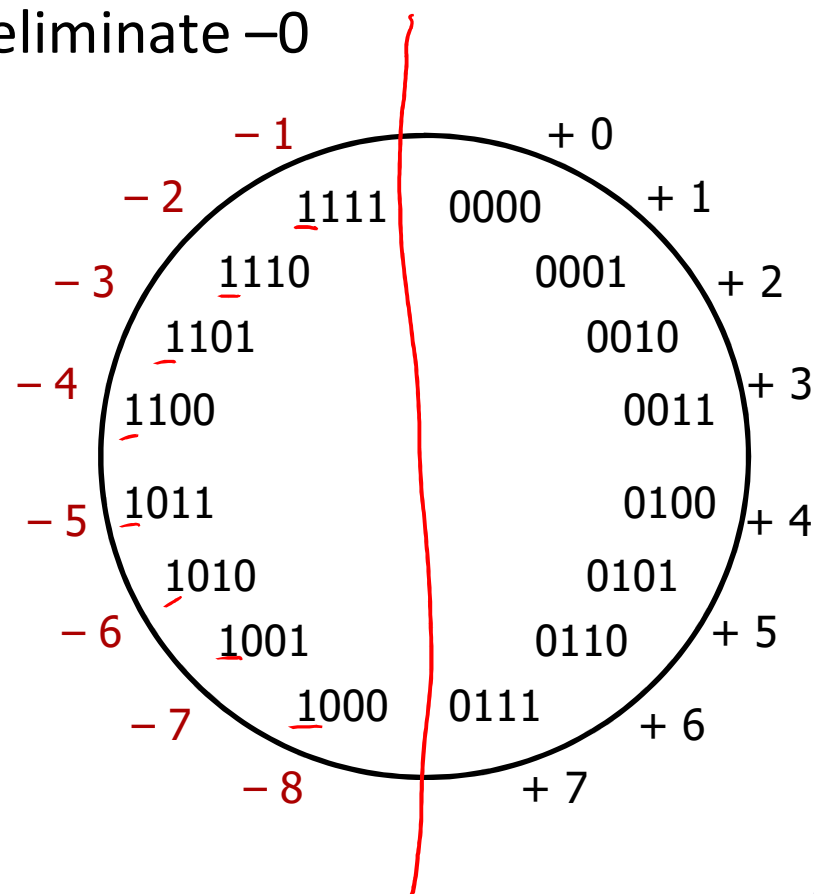
# Two's Complement

## ❖ Let's fix these problems:

- 1) "Flip" negative encodings so incrementing works
- 2) "Shift" negative numbers to eliminate  $-0$

## ❖ MSB *still* indicates sign!

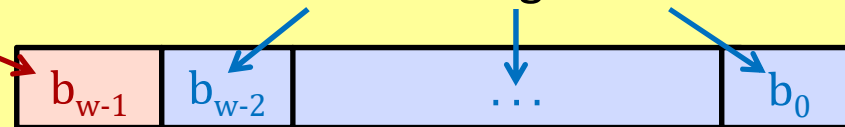
- This is why we represent one more negative than positive number ( $-2^{N-1}$  to  $2^{N-1} - 1$ )



# Two's Complement Negatives

- ❖ Accomplished with one neat mathematical trick!

$b_{w-1}$  has weight  $-2^{w-1}$ , other bits have usual weights  $+2^i$



## 4-bit Examples:

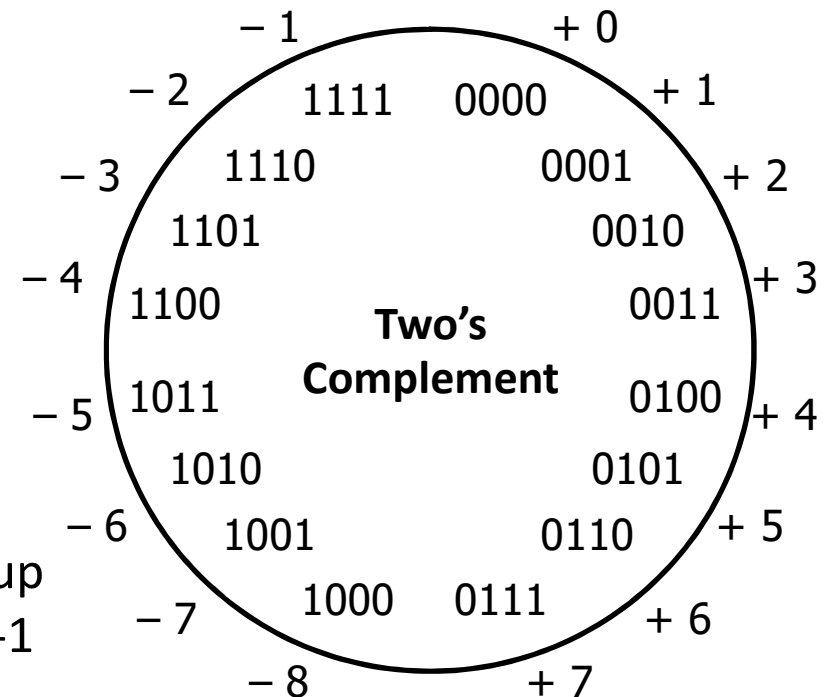
- $1010_2$  unsigned:  
 $1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 0 \cdot 2^0 = 10$
- $1010_2$  two's complement:  
 $-1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 0 \cdot 2^0 = -6$

## -1 represented as:

3 one's in a row

$$1111_2 = -(2^3) + (2^3 - 1)$$

- MSB makes it super negative, add up all the other bits to get back up to -1



# Why Two's Complement is So Great

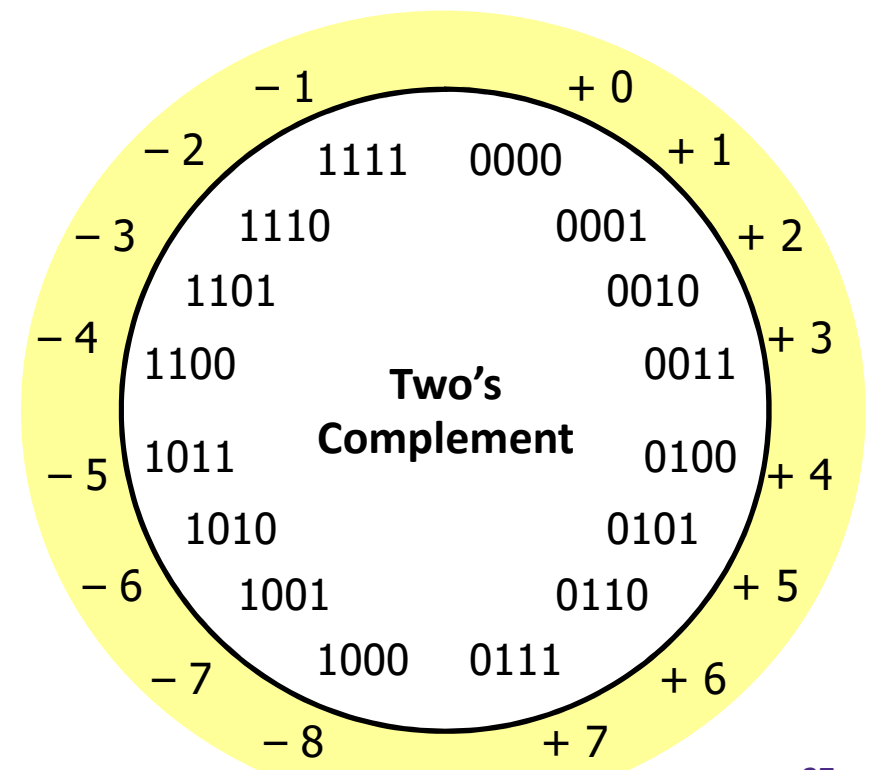
- ❖ Roughly same number of (+) and (−) numbers
- ❖ Positive number encodings match unsigned
- ❖ Single zero
- ❖ All zeros encoding = 0

- ❖ Simple negation procedure:

- Get negative representation of any integer by taking bitwise complement and then adding one!

$$(\sim x + 1 == -x)$$

complement the bit encoding of the number x, convert it to an integer via the two's complement encoding, and then add one. This gives the negative of the original integer value



# Peer Instruction Question

- ❖ Take the 4-bit number encoding  $x = 0b\overset{\text{MSB}}{\underset{\downarrow}{1}}011$
- ❖ Which of the following numbers is NOT a valid interpretation of  $x$  using any of the number representation schemes discussed today?
  - Unsigned, Sign and Magnitude, Two's Complement
  - Vote at <http://PollEv.com/rea>

A. -4

~~B. -5~~

~~C. 11~~

~~D. -3~~

E. We're lost...

unsigned:  $8 + 2 + 1 = 11$

sign + mag:  $1011 \rightarrow -(2+1) = -3$

two's:  $-8 + 2 + 1 = -5$

$-x = 0b\ 0100 + 1 = 5 \rightarrow x = -5$

# Summary

- ❖ Bit-level operators allow for fine-grained manipulations of data
  - Bitwise AND (`&`), OR (`|`), and NOT (`~`) different than logical AND (`&&`), OR (`||`), and NOT (`!`)
  - Especially useful with bit masks
- ❖ Choice of *encoding scheme* is important
  - Tradeoffs based on size requirements and desired operations
- ❖ Integers represented using unsigned and two's complement representations
  - Limited by fixed bit width
  - We'll examine arithmetic operations next lecture