

Integers II

CSE 351 Spring 2019

Instructor: Teaching Assistants:

Ruth Anderson

Gavin Cai

Britt Henderson

Sophie Tian

Casey Xing

Jack Eggleston

Richard Jiang

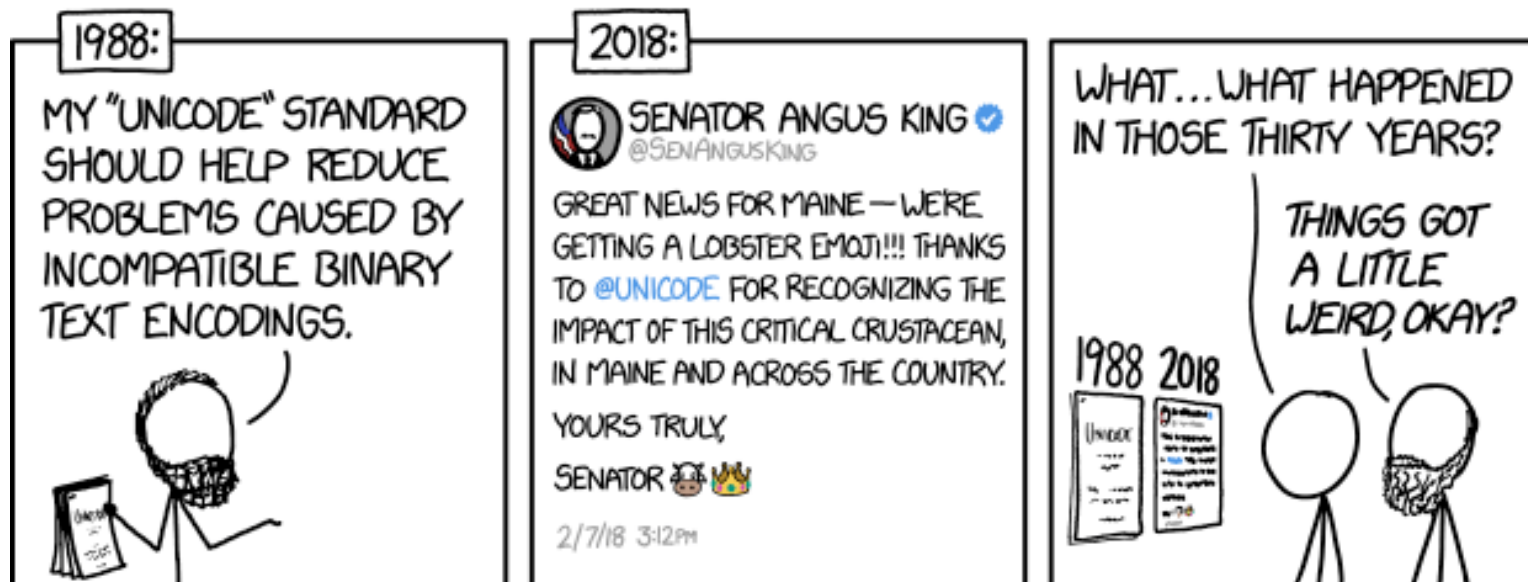
Connie Wang

Chin Yeoh

John Feltrup

Jack Skalitzy

Sam Wolfson



<http://xkcd.com/1953/>

Administrivia

- ❖ Homework 1 due TONIGHT (4/10)
 - Reminder: autograded, 20 tries, no late submissions
- ❖ Lab 1a due Monday (4/15)
 - Submit `pointer.c` and `lab1Areflect.txt` to Canvas
- ❖ Lab 1b released soon, due Monday 4/22
 - Bit puzzles on number representation
 - Have much of what you need after today, will need floating point, coming soon
 - Section tomorrow will be useful!
 - Bonus slides at the end of today's lecture have relevant examples

Extra Credit

- ❖ All labs starting with Lab 1b have extra credit portions
 - These are meant to be fun extensions to the labs
- ❖ Extra credit points *don't* affect your lab grades
 - From the course policies: “they will be accumulated over the course and will be used to bump up borderline grades at the end of the quarter.”
 - Make sure you finish the rest of the lab before attempting any extra credit

Integers

- ❖ **Binary representation of integers**
 - Unsigned and signed
 - Casting in C
- ❖ Consequences of finite width representations
 - Overflow, sign extension
- ❖ Shifting and arithmetic operations

Two's Complement Arithmetic

- ❖ The same addition procedure works for both unsigned and two's complement integers
 - **Simplifies hardware:** only one algorithm for addition
 - **Algorithm:** simple addition, **discard the highest carry bit**
 - Called modular addition: result is sum *modulo* 2^w

Can implement subtraction in a similar way;
see slide 15

❖ 4-bit Examples:

4	0100	-4	1100	4	0100
+3	+0011	+3	+0011	-3	+1101
=7	0111	=-1	1111	=1	0001

Why Does Two's Complement Work?

- ❖ For all representable positive integers x , we want:

$$\text{additive inverse} \left\{ \begin{array}{l} \text{bit representation of } x \\ + \text{ bit representation of } -x \end{array} \right. \underline{\hspace{1cm}} 0 \quad (\text{ignoring the carry-out bit})$$

- What are the 8-bit negative encodings for the following?

$$\begin{array}{r} 00000001 \\ + \quad ?\,?\,?\,?\,?\,?\,?\,? \\ \hline \cancel{X} 00000000 \end{array}$$

$$\begin{array}{r} 00000010 \\ + \quad ?\,?\,?\,?\,?\,?\,?\,? \\ \hline \cancel{X} 00000000 \end{array}$$

$$\begin{array}{r} 11000011 \\ + \quad ?\,?\,?\,?\,?\,?\,?\,? \\ \hline \cancel{X} 00000000 \end{array}$$

Why Does Two's Complement Work?

- ❖ For all representable positive integers x , we want:

$$\begin{array}{r} \text{bit representation of } x \\ + \text{ bit representation of } -x \\ \hline 0 \end{array}$$
 (ignoring the carry-out bit)

$$x + (\sim x) = 0b1\dots1$$

$$x + (\sim x) = -1$$

$$x + (\sim x + 1) = 0$$

$$-x = \sim x + 1$$

- What are the 8-bit negative encodings for the following?

$$\begin{array}{r} 00000001 \\ + 11111111 \\ \hline 100000000 \end{array}$$

$$\begin{array}{r} 00000010 \\ + 11111110 \\ \hline 100000000 \end{array}$$

$$\begin{array}{r} 11000011 \\ + 00111101 \\ \hline 100000000 \end{array}$$


These are the bitwise complement plus 1!

$$-x == \sim x + 1$$

Signed/Unsigned Conversion Visualized

❖ Two's Complement → Unsigned

- Ordering Inversion
- Negative → Big Positive



Values To Remember

❖ Unsigned Values

- UMin = 0b00...0
= 0
- UMax = 0b11...1
= $2^w - 1$

❖ Two's Complement Values

- TMin = 0b10...0
= -2^{w-1}
- TMax = 0b01...1
= $2^{w-1} - 1$
- -1 = 0b11...1

❖ Example: Values for $w = 64$

	Decimal	Hex
UMax	18,446,744,073,709,551,615	FF FF FF FF FF FF FF FF
TMax	9,223,372,036,854,775,807	7F FF FF FF FF FF FF FF
TMin	-9,223,372,036,854,775,808	80 00 00 00 00 00 00 00
-1	-1	FF FF FF FF FF FF FF FF
0	0	00 00 00 00 00 00 00 00

In C: Signed vs. Unsigned

❖ Casting

- Bits are unchanged, just interpreted differently!

- `int tx, ty;`
- `unsigned int ux, uy;`
- *Explicit* casting
 - `tx = (int) ux;`
 - `uy = (unsigned int) ty;`
- *Implicit* casting can occur during assignments or function calls
cast to target variable/parameter type
 - `tx = ux;`
 - `uy = ty;` *(also implicitly occurs with printf format specifiers)*



Casting Surprises

❖ Integer literals (constants)

- By default, integer constants are considered signed integers
 - Hex constants already have an explicit binary representation
- Use “U” (or “u”) suffix to explicitly force *unsigned*
 - Examples: 0U, 4294967259u

❖ Expression Evaluation

- When you mix unsigned and signed in a single expression, then **signed values are implicitly cast to unsigned** (unsigned “dominates”)
- Including comparison operators <, >, ==, <=, >=



Casting Surprises

❖ 32-bit examples:

- TMin = -2,147,483,648, TMax = 2,147,483,647

Left Constant	Order	Right Constant	Interpretation
0 0000 0000 0000 0000 0000 0000 0000 0000	=	0U 0000 0000 0000 0000 0000 0000 0000 0000	unsigned
-1 1111 1111 1111 1111 1111 1111 1111 1111	<	0 0000 0000 0000 0000 0000 0000 0000 0000	signed
-1 1111 1111 1111 1111 1111 1111 1111 1111	>	0U 0000 0000 0000 0000 0000 0000 0000 0000	unsigned
2147483647 0111 1111 1111 1111 1111 1111 1111 1111	>	-2147483648 1000 0000 0000 0000 0000 0000 0000 0000	signed
2147483647U 0111 1111 1111 1111 1111 1111 1111 1111	<	-2147483648 1000 0000 0000 0000 0000 0000 0000 0000	unsigned
-1 1111 1111 1111 1111 1111 1111 1111 1111	>	-2 1111 1111 1111 1111 1111 1111 1111 1110	signed
(unsigned) -1 1111 1111 1111 1111 1111 1111 1111 1111	>	-2 1111 1111 1111 1111 1111 1111 1111 1110	unsigned
2147483647 0111 1111 1111 1111 1111 1111 1111 1111	<	2147483648U 1000 0000 0000 0000 0000 0000 0000 0000	unsigned
2147483647 0111 1111 1111 1111 1111 1111 1111 1111	>	(int) 2147483648U 1000 0000 0000 0000 0000 0000 0000 0000	signed

Integers

- ❖ Binary representation of integers
 - Unsigned and signed
 - Casting in C
- ❖ **Consequences of finite width representations**
 - **Overflow, sign extension**
- ❖ Shifting and arithmetic operations

Arithmetic Overflow

Bits	Unsigned	Signed
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

- ❖ When a calculation produces a result that can't be represented in the current encoding scheme
 - Integer range limited by fixed width
 - Can occur in both the positive and negative directions
 - e.g. larger than U_{max} or smaller than U_{min} ; cannot be represented by the current number of bits
- ❖ C and Java ignore overflow exceptions
 - You end up with a bad value in your program and no warning/indication... oops!

Assigned arithmetic

Overflow: Unsigned

- ❖ **Addition:** drop carry bit (-2^N)

$$\begin{array}{r} 15 \\ + 2 \\ \hline \cancel{17} \\ 1 \end{array} \qquad \begin{array}{r} 1111 \\ + 0010 \\ \hline \cancel{1}0001 \end{array}$$


- ❖ **Subtraction:** borrow ($+2^N$)

$$\begin{array}{r} 1 \\ - 2 \\ \hline \cancel{-1} \\ 15 \end{array} \qquad \begin{array}{r} \cancel{1}0001 \\ - 0010 \\ \hline 1111 \end{array}$$

subtract and carry just like normal. Since unsigned integers are nonnegative, we always treat the first number like its larger (if necessary, we add 2^N by adding the grey one to the front)

$\pm 2^N$ because of modular arithmetic

$$2^4 = 16$$



i.e. addition and subtraction are modulo 2^w where w is the number of bits

Overflow: Two's Complement

- ❖ **Addition:** $(+) + (+) = (-)$ result?

sum is going above T_{max}

$$\begin{array}{r} 6 \\ + 3 \\ \hline \cancel{9} \\ -7 \end{array}$$

$$\begin{array}{r} 0110 \\ + 0011 \\ \hline 1001 \end{array}$$


- ❖ **Subtraction:** $(-) + (-) = (+)$?

sum is going below T_{min}

$$\begin{array}{r} -7 \\ - 3 \\ \hline \cancel{-10} \\ 6 \end{array}$$

$$\begin{array}{r} 1001 \\ - 0011 \\ \hline 0110 \end{array}$$

Same addition/subtraction as before; add/subtract and carry. Again, everything is modulo 2^N



For signed: overflow if operands have same sign and result's sign is different

Sign Extension

- ❖ What happens if you convert a *signed* integral data type to a larger one?

■ e.g. ^{1 byte}char → ^{2 bytes}short → ^{4 bytes}int → ^{8 bytes}long

- ❖ **4-bit → 8-bit Example:**

- Positive Case

✓ • Add 0's?

4-bit: 0010 = +2

8-bit: 00000010 = +2

- Negative Case?

Peer Instruction Question

- ❖ Which of the following 8-bit numbers has the same *signed* value as the 4-bit number **0b1100**? ^{-8 4 2 1} $-8+4 = -4$

- Underlined digit = MSB
- Vote at <http://pollev.com/rea>

$$-x = \begin{array}{r} 0011 \\ +1 \\ \hline 0100 = 4 \Rightarrow x = -4 \end{array}$$

~~A.~~ 0b 0000 1100

~~B.~~ 0b 1000 1100

C. 0b 1111 1100

~~D.~~ 0b 1100 1100

E. We're lost...

positive!

much too negative: $-2^7 + 2^3 + 2^2 = -116$

correct! $-y = 0b\ 0000\ 0011 + 1 = 4, \ y = -4$

$-2^7 + 2^6 + 2^3 + 2^2 = -52$


Sign Extension

❖ **Task:** Given a w -bit signed integer X , convert it to $w+k$ -bit signed integer X' *with the same value*

❖ **Rule:** Add k copies of sign bit

■ Let x_i be the i -th digit of X in binary

■ $X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, \underbrace{x_{w-1}, x_{w-2}, \dots, x_1, x_0}_{\text{original } X}$



Sign Extension Example

- ❖ Convert from smaller to larger integral data types
- ❖ C automatically performs sign extension
 - Java too

```
short int x = 12345;
int      ix = (int) x;
short int y = -12345;
int      iy = (int) y;
```

Var	Decimal	Hex	Binary
x	12345	30 39	00110000 00111001
ix	12345	00 00 30 39	00000000 00000000 00110000 00111001
y	-12345	CF C7	11001111 11000111
iy	-12345	FF FF CF C7	11111111 11111111 11001111 11000111

Handwritten annotations:
 - Red arrow from '30 39' to '0b 0011'
 - Red arrow from 'CF' to '0b 1100'

Integers

- ❖ Binary representation of integers
 - Unsigned and signed
 - Casting in C
- ❖ Consequences of finite width representations
 - Overflow, sign extension
- ❖ **Shifting and arithmetic operations**

Shift Operations

- ❖ Left shift ($x \ll n$) bit vector x by n positions
 - Throw away (drop) extra bits on left
 - Fill with 0s on right
- ❖ Right shift ($x \gg n$) bit-vector x by n positions
 - Throw away (drop) extra bits on right
 - Logical shift (for unsigned values)
 - Fill with 0s on left
 - Arithmetic shift (for signed values)
 - Replicate most significant bit on left
 - Maintains sign of x

Shift Operations

❖ Left shift ($x \ll n$)

- Fill with 0s on right

❖ Right shift ($x \gg n$)

- Logical shift (for **unsigned** values)
 - Fill with 0s on left
- Arithmetic shift (for **signed** values)
 - Replicate most significant bit on left

❖ Notes:

- Shifts by $n < 0$ or $n \geq w$ (w is bit width of x) are **undefined** *behavior not guaranteed*
- **In C:** behavior of \gg is determined by compiler
 - In gcc / C lang, depends on data type of x (signed/unsigned) *arithmetic / logical*
- **In Java:** logical shift is \ggg and arithmetic shift is \gg

8-bit example:

x	0010 0010
$x \ll 3$	0001 0 000
logical: $x \gg 2$	00 00 1000
arithmetic: $x \gg 2$	00 00 1000

x	1010 0010
$x \ll 3$	0001 0 000
logical: $x \gg 2$	00 10 1000
arithmetic: $x \gg 2$	11 10 1000

Shifting Arithmetic?

❖ What are the following computing?

- $x \gg n$
 - $0b \underbrace{0100}_4 \gg 1 = 0b \overbrace{0010}^2$
 - $0b \underbrace{0100}_4 \gg 2 = 0b \underbrace{0001}_1$
 - Divide by 2^n
- $x \ll n$
 - $0b \underbrace{0001}_1 \ll 1 = 0b \overbrace{0010}^2$
 - $0b \underbrace{0001}_1 \ll 2 = 0b \underbrace{0100}_4$
 - Multiply by 2^n

❖ Shifting is faster than general multiply and divide operations

Left Shifting Arithmetic 8-bit Example

- ❖ No difference in left shift operation for unsigned and signed numbers (just manipulates bits)

- Difference comes during interpretation: $x * 2^n$?

		Signed	Unsigned
$x = 25;$	00011001 =	25	25
$L1 = x << 2;$	00 01100100 =	100	100
$L2 = x << 3;$	000 11001000 =	-56	200
$L3 = x << 4;$	0001 10010000 =	-112	144

signed overflow
 unsigned overflow

Handwritten notes: $200 - 256 \rightarrow 2^8$ (for L2), $400 - 256 \rightarrow 2^8$ (for L3)

Right Shifting 8-bit Examples

- ❖ **Reminder:** C operator `>>` does *logical* shift on **unsigned** values and *arithmetic* shift on **signed** values

- **Logical** Shift: $x / 2^n$?

`xu = 240u;` `11110000` `= 240` $/8 = 30$

`R1u=xu>>3;` `00011110` `= 30` $/4 = 7.5$

`R2u=xu>>5;` `00000111` `= 7`

rounding (down)

Right Shifting Arithmetic 8-bit Examples

- ❖ **Reminder:** C operator `>>` does *logical* shift on **unsigned** values and *arithmetic* shift on **signed** values

- **Arithmetic** Shift: $x / 2^n$?

`xs = -16;` 11110000 = -16

`R1s=xs>>3;` 11111110 = -2 $\frac{1}{4} = -0.5$

`R2s=xs>>5;` 11111111 = -1

rounding (down)

Question

$uMin = 0, uMax = 255$
 8-bits, so $TMin = -128, TMax = 127$

For the following expressions, find a value of **signed char** x , if there exists one, that makes the expression TRUE. Compare with your neighbor(s)!

❖ Assume we are using 8-bit arithmetic:

■ x ^{unsigned} $==$ (unsigned char) x	<u>Example:</u> $x = 0$	<u>All solutions:</u> works for all x
■ x ^{unsigned} $>=$ 128U <small>0b10000000</small>	$x = -1$	any $x < 0$
■ $x \neq (x >> 2) << 2$	$x = 3$	any x where lowest two bits are not 0b00
■ $x == -x$ • Hint: there are two solutions	$x = 0$	① $x = 0b0...0 = 0$ ② $x = 0b10...0 = -128$
■ $(x < 128U) \ \&\& \ (x > 0x3F)$	$x = 64$	any x where upper two bits are exactly 0b01

Summary

- ❖ Sign and unsigned variables in C
 - Bit pattern remains the same, just *interpreted* differently
 - Strange things can happen with our arithmetic when we convert/cast between sign and unsigned numbers
 - Type of variables affects behavior of operators (shifting, comparison)
- ❖ We can only represent so many numbers in w bits
 - When we exceed the limits, *arithmetic overflow* occurs
 - *Sign extension* tries to preserve value when expanding
- ❖ Shifting is a useful bitwise operator
 - Right shifting can be arithmetic (sign) or logical (0)
 - Can be used in multiplication with constant or bit masking

BONUS SLIDES

Some examples of using shift operators in combination with bitmasks, which you may find helpful for Lab 1.

- ❖ Extract the 2nd most significant byte of an `int`
- ❖ Extract the sign bit of a signed `int`
- ❖ Conditionals as Boolean expressions

Using Shifts and Masks

❖ Extract the 2nd most significant *byte* of an `int`:

- First shift, then mask: $(x \gg 16) \& 0xFF$

x	00000001	00000010	00000011	00000100
x >> 16	00000000	00000000	00000001	00000010
0xFF	00000000	00000000	00000000	11111111
(x >> 16) & 0xFF	00000000	00000000	00000000	00000010

- Or first mask, then shift: $(x \& 0xFF0000) \gg 16$

x	00000001	00000010	00000011	00000100
0xFF0000	00000000	11111111	00000000	00000000
x & 0xFF0000	00000000	00000010	00000000	00000000
(x & 0xFF0000) >> 16	00000000	00000000	00000000	00000010

Using Shifts and Masks

- ❖ Extract the *sign bit* of a signed `int`:
 - First shift, then mask: $(x \gg 31) \ \& \ 0x1$
 - Assuming arithmetic shift here, but this works in either case
 - Need mask to clear 1s possibly shifted in

x	00000001 00000010 00000011 00000100
x >> 31	00000000 00000000 00000000 00000000
0x1	00000000 00000000 00000000 00000001
(x >> 31) & 0x1	00000000 00000000 00000000 00000000

x	10000001 00000010 00000011 00000100
x >> 31	11111111 11111111 11111111 11111111
0x1	00000000 00000000 00000000 00000001
(x >> 31) & 0x1	00000000 00000000 00000000 00000001

Using Shifts and Masks

❖ Conditionals as Boolean expressions

- For `int x`, what does `(x<<31)>>31` do?

<code>x=!!123</code>	00000000 00000000 00000000 00000000 1
<code>x<<31</code>	1 00000000 00000000 00000000 00000000
<code>(x<<31)>>31</code>	11111111 11111111 11111111 11111111
<code>!x</code>	00000000 00000000 00000000 00000000 0
<code>!x<<31</code>	0 00000000 00000000 00000000 00000000
<code>(!x<<31)>>31</code>	00000000 00000000 00000000 00000000

- Can use in place of conditional:

- In C: `if(x) {a=y;} else {a=z;} equivalent to a=x?y:z;`
- `a=(((x<<31)>>31)&y) | (((!x<<31)>>31)&z);`