

Floating Point I

CSE 351 Spring 2019

Instructor: **Teaching Assistants:**

Ruth Anderson

Gavin Cai

Britt Henderson

Sophie Tian

Casey Xing

Jack Eggleston

Richard Jiang


Connie Wang

Chin Yeoh

John Feltrup

Jack Skalitzy

Sam Wolfson




<http://xkcd.com/571/>

Administrivia

- ❖ Lab 1a due Monday 4/15 at 11:59 pm
 - Submit `pointer.c` and `lab1Areflect.txt`
- ❖ Lab 1b due Monday (4/22)
 - Submit `bits.c` and `lab1Breflect.txt`
- ❖ Homework 2 coming soon, due Wednesday (4/24)
 - On Integers, Floating Point, and x86-64

Unsigned Multiplication in C



❖ Standard Multiplication Function

- Ignores high order w bits

When we multiply two unsigned w-bit numbers, it could possibly take 2w bits to store the result (e.g. if we multiply $2^w - 1$ by $2^w - 1$, the max number representable by a 2bit unsigned int, then the product is $2^{2w} - 2^{w+1} + 1$ which requires 2^{2w} bits to represent

❖ Implements Modular Arithmetic

- $\text{UMult}_w(u, v) = u \cdot v \bmod 2^w$

Again, instead of allocating more memory, we just throw away the higher order bits and treat multiplication like its mod 2^w (same as addition)


Multiplication with shift and add

This just appends zeros on the right hand side
and throws away higher order bits

❖ Operation $u \ll k$ gives $u * 2^k$ (as long as no overflow occurs;
operations are still modulo 2^w)

- Both signed and unsigned


Operands: w bits



True Product: $w + k$ bits



Discard k bits: w bits



❖ Examples:

- $u \ll 3 == u * 8$
- $u \ll 5 - u \ll 3 == u * 24$
 $u \ll 4 + u \ll 3$
 $2^5 - 2^3 \rightarrow 24 = 32 - 8$
 $\rightarrow 24 = 16 + 8$
- Most machines shift and add faster than multiply
 - Compiler generates this code automatically

Number Representation Revisited

- ❖ We know how to represent:
 - Signed and Unsigned Integers
 - Characters (ASCII)
 - Addresses
- ❖ How do we encode the following:
 - Real numbers (*e.g.* 3.14159)
 - Very large numbers (*e.g.* 6.02×10^{23})
 - Very small numbers (*e.g.* 6.626×10^{-34})
 - Special numbers (*e.g.* ∞ , NaN)

} **Floating
Point**

Floating Point Topics

- ❖ **Fractional binary numbers**
- ❖ IEEE floating-point standard
- ❖ Floating-point operations and rounding
- ❖ Floating-point in C



- ❖ There are many more details that we won't cover
 - It's a 58-page standard...


Floating Point Summary

- ❖ Floats also suffer from the fixed number of bits available to represent them
 - Can get overflow/underflow, just like `ints`
 - “Gaps” produced in representable numbers means we can lose precision, unlike `ints`
 - We can only represent a finite number of numbers with our finite number of bits; so theres no way we can represent any kind of interval of real numbers
 - Some “simple fractions” have no exact representation (e.g. 0.2)
 - “Every operation gets a slightly wrong result”
- ❖ Floating point arithmetic not associative or distributive
 - *Mathematically* equivalent ways of writing an expression may compute different results
- ❖ **Never** test floating point values for equality!
- ❖ **Careful** when converting between `ints` and `floats`!

Representation of Fractions

- ❖ “Binary Point,” like decimal point, signifies boundary between integer and fractional parts:

Example 6-bit representation:




- ❖ Example: 10.1010₂ = 1× 2^1 + 1×2^{-1} + 1×2^{-3} = 2.625₁₀

Representation of Fractions

- ❖ “Binary Point,” like decimal point, signifies boundary between integer and fractional parts:

Example 6-bit representation:



- ❖ In this 6-bit representation:
 - What is the encoding and value of the smallest (most negative) number?
 - What is the encoding and value of the largest (most positive) number?
 - What is the smallest number greater than 2 that we can represent?

$$00.0000_2 = 0$$

$$11.111\textcircled{1}_2 = 4 - 2^{-4}$$

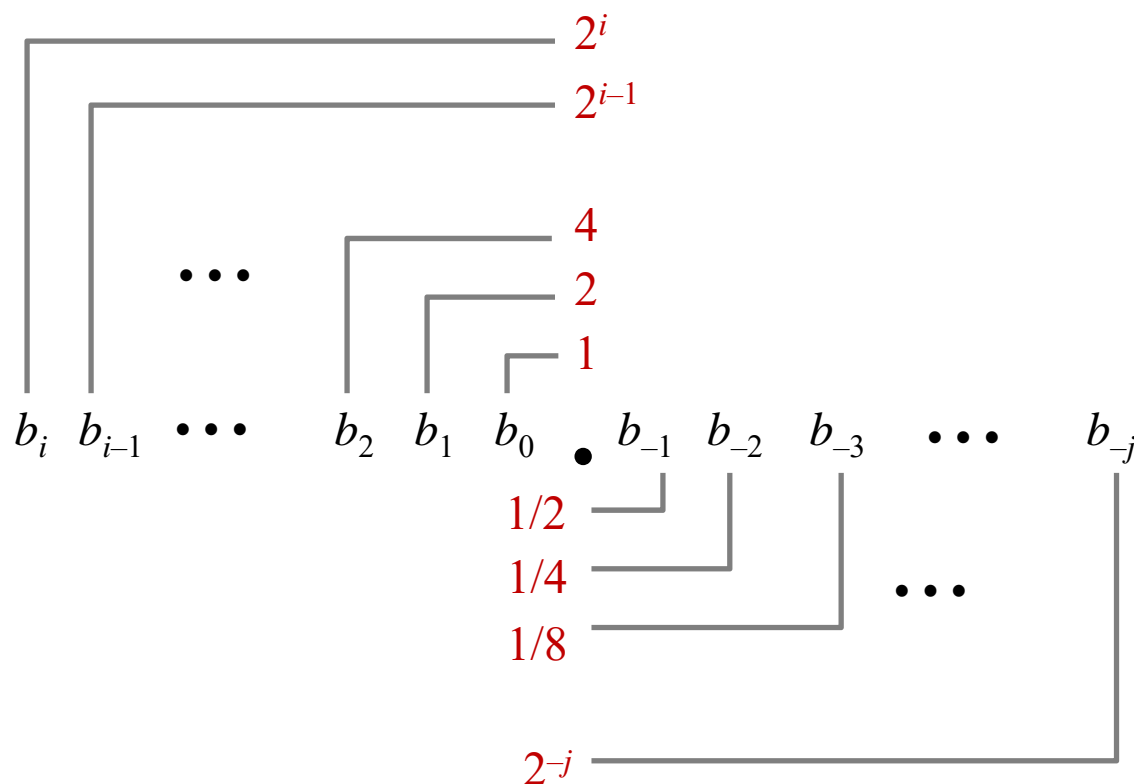
\uparrow
 2^{-4}

$$2^1 = 10.0000_2$$

$$10.0001 = 2 + 2^{-4}$$

can't represent anything in-between!
☹️

Fractional Binary Numbers



❖ Representation

- Bits to right of “binary point” represent fractional powers of 2
- Represents rational number:

$$\sum_{k=-j}^i b_k \cdot 2^k$$

Fractional Binary Numbers

❖ Value Representation

- 5 and 3/4 101.11_2
- 2 and 7/8 10.111_2
- 47/64 0.101111_2

❖ Observations

- Shift left = multiply by power of 2
- Shift right = divide by power of 2
- Numbers of the form $0.111111..._2$ are just below 1.0
 - $1/2 + 1/4 + 1/8 + ... + 1/2^i + ... \rightarrow 1.0$
 - Use notation $1.0 - \epsilon$

Limits of Representation

❖ Limitations:

- Even given an arbitrary number of bits, can only exactly represent numbers of the form $x * 2^y$ (y can be negative)
where x is an integer
- Other rational numbers have repeating bit representations

Value:	Binary Representation:
• $1/3 = 0.333333..._{10}$	$0.01010101[01]..._2$
• $1/5 = 0.2_{10}$	$0.001100110011[0011]..._2$
• $1/10 = 0.1_{10}$	$0.0001100110011[0011]..._2$

Fixed Point Representation

e.g how many bits do we allocate to the decimal part and how many bits do we allocate to the integer part

- ❖ Implied binary point. Two example schemes:
 - #1: the binary point is between bits 2 and 3
 $b_7 b_6 b_5 b_4 b_3 \text{ [.] } b_2 b_1 b_0$
 - #2: the binary point is between bits 4 and 5
 $b_7 b_6 b_5 \text{ [.] } b_4 b_3 b_2 b_1 b_0$
- ❖ Wherever we put the binary point, with fixed point representations there is a trade off between the amount of range and precision we have
- ❖ Fixed point = fixed *range* and fixed *precision*
 - range: difference between largest and smallest numbers possible
 - precision: smallest possible difference between any two numbers
- ❖ Hard to pick how much you need of each!

Floating Point Representation

As opposed to fractional binary numbers, we now try to represent numbers of the form decimal * 2^{exponent}. So we've given more flexibility to the first term essentially.

❖ Analogous to scientific notation

■ In Decimal:

- Not 12000000, but 1.2 x 10⁷ In C: 1.2e7
- Not 0.0000012, but 1.2 x 10⁻⁶ In C: 1.2e-6

■ In Binary:

- Not 11000.000, but 1.1 x 2⁴
- Not 0.000101, but 1.01 x 2⁻⁴

❖ We have to divvy up the bits we have (e.g., 32) among:

- the sign (1 bit)
- the mantissa (significand) this is the number being multiplied by 2^k
- the exponent The power on 2^k

Scientific Notation (Decimal)

(significand)
mantissa

exponent

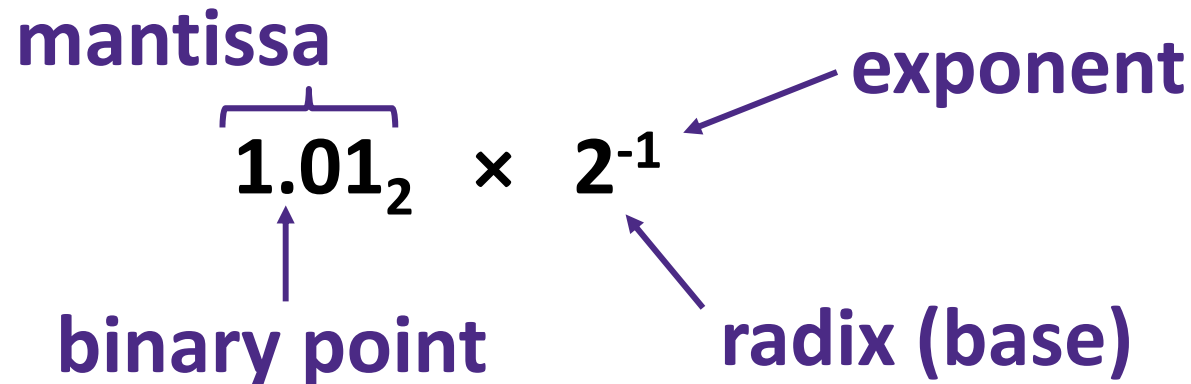
decimal point

radix (base)

$$6.02_{10} \times 10^{23}$$
A diagram illustrating the components of the scientific notation expression 6.02_10 × 10^23. A bracket above '6.02' is labeled '(significand)' in red and 'mantissa' in purple. An arrow points from 'exponent' to '10^23'. An arrow points from 'radix (base)' to '10'. An arrow points from 'decimal point' to the dot in '6.02'.

- ❖ Normalized form: exactly one digit (non-zero) to left of decimal point
- ❖ Alternatives to representing 1/1,000,000,000
 - Normalized: 1.0×10^{-9}
 - Not normalized: $0.1 \times 10^{-8}, 10.0 \times 10^{-10}$

Scientific Notation (Binary)



The diagram illustrates the components of the binary scientific notation $1.01_2 \times 2^{-1}$. The mantissa is 1.01_2 , with a bracket above it labeled "mantissa" and an arrow pointing to the decimal point labeled "binary point". The exponent is 2^{-1} , with an arrow pointing to the 2 labeled "radix (base)" and an arrow pointing to the -1 labeled "exponent".

- ❖ Computer arithmetic that supports this called **floating point** due to the “floating” of the binary point
 - Declare such variable in C as `float` (or `double`)

The question still is what is the best way to encode the "mantissa" and the exponent, and how many bits should we allocate to represent each.

Scientific Notation Translation

$$\begin{aligned} 2^{-1} &= 0.5 \\ 2^{-2} &= 0.25 \\ 2^{-3} &= 0.125 \\ 2^{-4} &= 0.0625 \end{aligned}$$

- ❖ Convert from scientific notation to binary point
 - Perform the multiplication by shifting the decimal until the exponent disappears
 - Example: $1.011_2 \times 2^4 = 10110_2 = 22_{10}$
 - Example: $1.011_2 \times 2^{-2} = 0.01011_2 = 0.34375_{10}$
- ❖ Convert from binary point to normalized scientific notation
 - Distribute out exponents until binary point is to the right of a single digit
 - Example: $1101.001_2 = 1.101001_2 \times 2^3$
- ❖ **Practice:** Convert 11.375_{10} to binary scientific notation

$$\begin{aligned} &8 + 2 + 1 + 0.25 + 0.125 \\ &2^3 + 2^1 + 2^0 + 2^{-2} + 2^{-3} = \underbrace{1011}_2 . 011_2 = \boxed{1.011011 \times 2^3} \end{aligned}$$

IEEE Floating Point

❖ IEEE 754

- Established in 1985 as uniform standard for floating point arithmetic
- Main idea: make numerically sensitive programs portable
- Specifies two things: representation and result of floating operations
- Now supported by all major CPUs

❖ Driven by numerical concerns

- **Scientists**/numerical analysts want them to be as **real** as possible
- **Engineers** want them to be **easy to implement** and **fast**
- In the end:
 - Scientists mostly won out
 - Nice standards for rounding, overflow, underflow, but...
 - Hard to make fast in hardware
 - Float operations can be an order of magnitude slower than integer ops

} competing goals

Floating Point Encoding

- ❖ Use normalized, base 2 scientific notation:

- Value: $\pm 1 \times \text{Mantissa} \times 2^{\text{Exponent}}$

- Bit Fields: $(-1)^S \times 1.M \times 2^{(E-\text{bias})}$


- ❖ Representation Scheme: (3 separate fields within 32 bits)

- Sign bit (0 is positive, 1 is negative)


values

- Mantissa (a.k.a. significand) is the fractional part of the number in normalized form and encoded in bit vector **M**

- Exponent weights the value by a (possibly negative) power of 2 and encoded in the bit vector **E**



The Exponent Field



❖ Use **biased notation**

- Read exponent as unsigned, but with **bias** of $2^{w-1}-1 = 127$
- Representable exponents roughly $\frac{1}{2}$ positive and $\frac{1}{2}$ negative
- Exponent 0 ($\text{Exp} = 0$) is represented as $\text{E} = 0b\ 0111\ 1111 = 2^7 - 1$
 $\text{E} - \text{bias} = 0 = \text{Exp}$

$w=8$, can encode $2^8=256$ exponents

❖ Why biased?


- Makes floating point arithmetic easier
- Makes somewhat compatible with two's complement

❖ **Practice:** To encode in biased notation, add the bias then encode in unsigned:

- $\text{Exp} = 1 \xrightarrow{+\text{bias}} 128 \xrightarrow{\text{encode}} \text{E} = 0b\ 1000\ 0000$
- $\text{Exp} = 127 \rightarrow 254 \rightarrow \text{E} = 0b\ 1111\ 1110$ ($254 = 255 - 1 = (2^8 - 1) - 1$)
- $\text{Exp} = -63 \rightarrow 64 \rightarrow \text{E} = 0b\ 0100\ 0000$

conflicting notation with what is used in the textbook. The textbook calls $\text{E} = \text{exp}$ and $\text{Exp} = \text{E}$

The Mantissa (Fraction) Field



$$(-1)^S \times (1.M) \times 2^{(E-\text{bias})}$$

1.1×2^0

- ❖ Note the implicit 1 in front of the M bit vector

■ Example: 0b 0011 1111 11100 0000 0000 0000 0000 0000

is read as $1.1_2 = 1.5_{10}$, *not* $0.1_2 = 0.5_{10}$

- Gives us an extra bit of *precision*

- ❖ Mantissa “limits”

- Low values near $M = 0b0...0$ are close to 2^{Exp}

- High values near $M = 0b1...1$ are close to $2^{\text{Exp}+1}$

$$\rightarrow 2^{\text{Exp}} \times 1.0...0 = 2^{\text{Exp}}$$

$$\hookrightarrow 2^{\text{Exp}} \times 1.1...1 = 2^{\text{Exp}}(2 - 2^{-23}) = 2^{\text{Exp}+1} - 2^{\text{Exp}-23}$$

Peer Instruction Question

Bias = 127

- ❖ What is the correct value encoded by the following floating point number?

■ 0b 0 ^S 10000000 ^E 11000000000000000000000000000000 ^M

- Vote at <http://pollev.com/rea> [⊕] ¹²⁸⁻¹²⁷ Exp = 1 ^{Man = 1.110...0} ^{implicit}

A. + 0.75

B. + 1.5

C. + 2.75

D. + 3.5

E. We're lost...

$$+ 1.\underline{1}1_2 \times 2^1$$


$$11.1_2 = 2^1 + 2^0 + 2^{-1} = 3.5$$

Precision and Accuracy

- ❖ **Precision** is a count of the number of bits in a computer word used to represent a value
 - Capacity for accuracy
- ❖ **Accuracy** is a measure of the difference between the *actual value of a number* and its computer representation
 - *High precision permits high accuracy but doesn't guarantee it. It is possible to have high precision but low accuracy.*
 - **Example:** `float pi = 3.14;`
 - `pi` will be represented using all 24 bits of the mantissa (highly precise), but is only an approximation (not accurate)

Need Greater Precision?

❖ Double Precision (vs. Single Precision) in 64 bits



- C variable declared as double vs float which is only single precision
- Exponent bias is now $2^{10}-1 = 1023$, $\text{bias} = 2^{w-1}-1$
- **Advantages:** greater precision (larger mantissa), greater range (larger exponent)
- **Disadvantages:** more bits used, slower to manipulate

Representing Very Small Numbers

❖ But wait... what happened to zero?


- Using standard encoding $0x00000000 = 1.0 \times 2^{-127} \neq 0$
 $S=0, E=0, M=0 \Rightarrow \text{Exp} = -127, \text{Man} = 1.0 \dots 0$
- *Special case:* **E** and **M** all zeros = 0
 - Two zeros! But at least $0x00000000 = 0$ like integers
 $0x80000000 = -0$

❖ New numbers closest to 0:

$(E = 0x01, \text{Exp} = -126)$

- $a = 1.0 \dots 0_2 \times 2^{-126} = 2^{-126}$
- $b = 1.0 \dots 01_2 \times 2^{-126} = 2^{-126} + 2^{-149}$
 $\xrightarrow{2^3}$

- Normalization and implicit 1 are to blame
- *Special case:* **E** = 0, **M** \neq 0 are **denormalized numbers**



This is extra
(non-testable)
material

Denorm Numbers

❖ Denormalized numbers

- No leading 1
- Uses implicit exponent of -126 even though $E = 0x00$

Because we have taken out the implicit 1 at the front of M
Namely we use $0.M \times 2^{\text{exp}}$ instead of $1.M \times 2^{\text{exp}}$

❖ Denormalized numbers close the gap between zero and the smallest normalized number

- Smallest norm: $\pm 1.0\dots0_{\text{two}} \times 2^{-126} = \pm 2^{-126}$
- Smallest denorm: $\pm 0.0\dots01_{\text{two}} \times 2^{-126} = \pm 2^{-149}$

- There is still a gap between zero and the smallest denormalized number

So much
closer to 0

Other Special Cases


- ❖ $E = 0xFF, M = 0$: $\pm \infty$
 - *e.g.* division by 0
 - Still work in comparisons!
- ❖ $E = 0xFF, M \neq 0$: Not a Number (NaN)
 - *e.g.* square root of negative number, $0/0, \infty - \infty$
 - NaN propagates through computations
 - Value of M can be useful in debugging
- ❖ New largest value (besides ∞)?
 - $E = 0xFF$ has now been taken!
 - $E = 0xFE$ has largest: $1.1...1_2 \times 2^{127} = 2^{128} - 2^{104}$

Floating Point Encoding Summary

E	M	Meaning
0x00	0	± 0
0x00	non-zero	\pm denorm num
0x01 – 0xFE	anything	\pm norm num
0xFF	0	$\pm \infty$
0xFF	non-zero	NaN

Summary

❖ Floating point approximates real numbers:



- Handles large numbers, small numbers, special numbers
- Exponent in biased notation (bias = $2^{w-1}-1$)
 - Outside of representable exponents is *overflow* and *underflow*
- Mantissa approximates fractional portion of binary point
 - Implicit leading 1 (normalized) except in special cases
 - Exceeding length causes *rounding*


E	M	Meaning
0x00	0	± 0
0x00	non-zero	\pm denorm num
0x01 – 0xFE	anything	\pm norm num
0xFF	0	$\pm \infty$
0xFF	non-zero	NaN

BONUS SLIDES

An example that applies the IEEE Floating Point concepts to a smaller (8-bit) representation scheme.

These slides expand on material covered today, so while you don't need to read these, the information is "fair game."

Tiny Floating Point Example



❖ 8-bit Floating Point Representation

- The sign bit is in the most significant bit (MSB)
- The next four bits are the exponent, with a bias of $2^{4-1}-1 = 7$
- The last three bits are the mantissa

❖ Same general form as IEEE Format

- Normalized binary scientific point notation
- Similar special cases for 0, denormalized numbers, NaN, ∞

Dynamic Range (Positive Only)

	S	E	M	Exp	Value	
Denormalized numbers	0	0000	000	-6	0	
	0	0000	001	-6	$1/8 * 1/64 = 1/512$	closest to zero
	0	0000	010	-6	$2/8 * 1/64 = 2/512$	
	...					
	0	0000	110	-6	$6/8 * 1/64 = 6/512$	
Normalized numbers	0	0000	111	-6	$7/8 * 1/64 = 7/512$	largest denorm
	0	0001	000	-6	$8/8 * 1/64 = 8/512$	smallest norm
	0	0001	001	-6	$9/8 * 1/64 = 9/512$	
	...					
	0	0110	110	-1	$14/8 * 1/2 = 14/16$	
Normalized numbers	0	0110	111	-1	$15/8 * 1/2 = 15/16$	closest to 1 below
	0	0111	000	0	$8/8 * 1 = 1$	
	0	0111	001	0	$9/8 * 1 = 9/8$	closest to 1 above
	0	0111	010	0	$10/8 * 1 = 10/8$	
	...					
	0	1110	110	7	$14/8 * 128 = 224$	
	0	1110	111	7	$15/8 * 128 = 240$	largest norm
	0	1111	000	n/a	inf	

Special Properties of Encoding

- ❖ Floating point zero (0^+) exactly the same bits as integer zero
 - All bits = 0

- ❖ Can (Almost) Use Unsigned Integer Comparison
 - Must first compare sign bits
 - Must consider $0^- = 0^+ = 0$
 - NaNs problematic
 - Will be greater than any other values
 - What should comparison yield?
 - Otherwise OK
 - Denorm vs. normalized
 - Normalized vs. infinity