

The Stack & Procedures

CSE 351 Spring 2019

Instructor:

Ruth Anderson

Teaching Assistants:

Gavin Cai

Britt Henderson

Sophie Tian

Casey Xing

Jack Eggleston

Richard Jiang

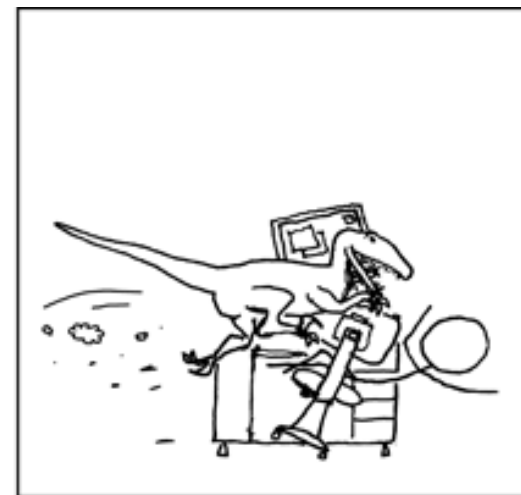
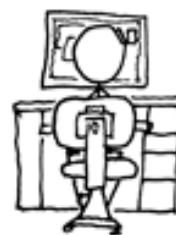
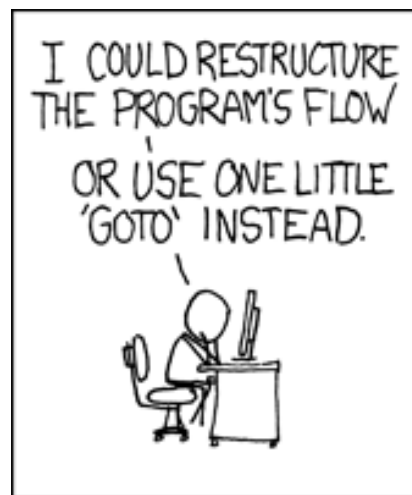
Connie Wang

Chin Yeoh

John Feltrup

Jack Skalitzyk

Sam Wolfson



<http://xkcd.com/571/>

Administrivia

- ❖ Homework 2 due TONIGHT Wednesday (4/24)
- ❖ Lab 2 (x86-64) due Wednesday (5/01)
 - Ideally want to finish well before the midterm
- ❖ Homework 3, coming soon
 - On midterm material, but due after the midterm
- ❖ Section tomorrow on Assembly and GDB
 - Bring your laptops!
- ❖ **Midterm** (Fri 5/03, 4:30-5:30pm in KNE 130)

Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

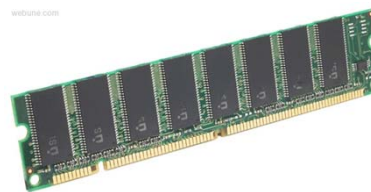
Assembly
language:

```
get_mpg:
    pushq    %rbp
    movq    %rsp, %rbp
    ...
    popq    %rbp
    ret
```

Machine
code:

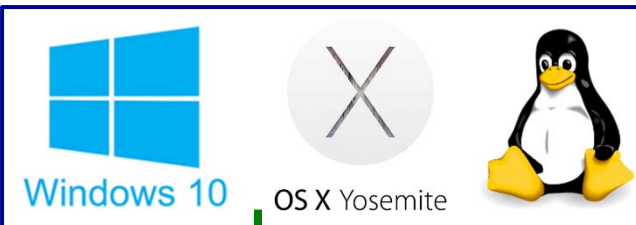
```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

Computer
system:



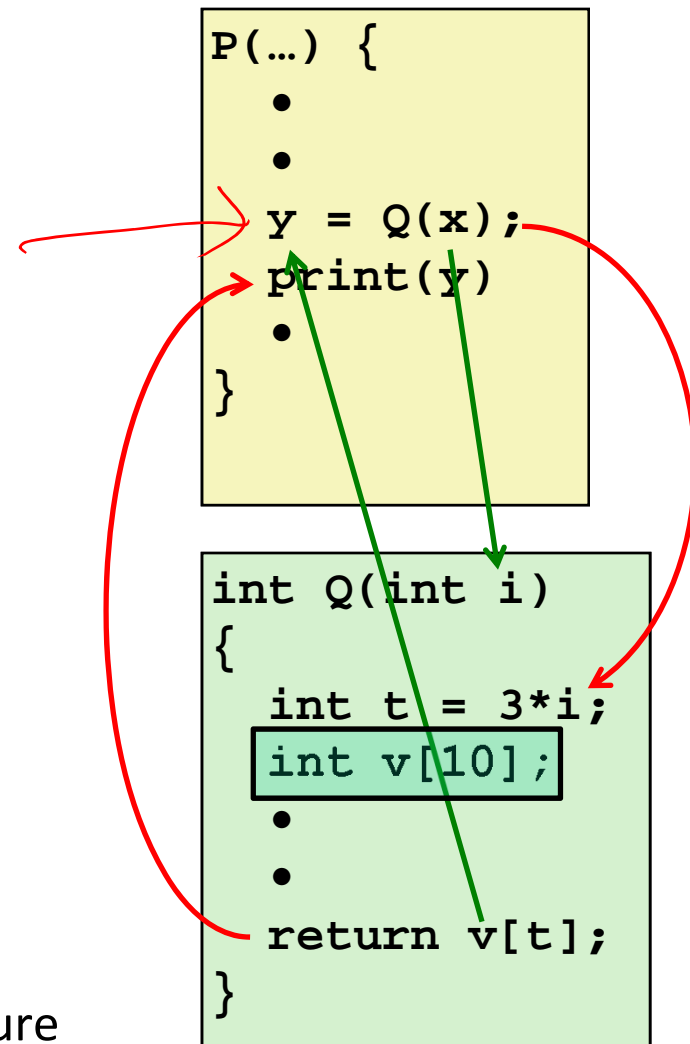
Memory & data
Integers & floats
x86 assembly
Procedures & stacks
Executables
Arrays & structs
Memory & caches
Processes
Virtual memory
Memory allocation
Java vs. C

OS:



Mechanisms required for *procedures*

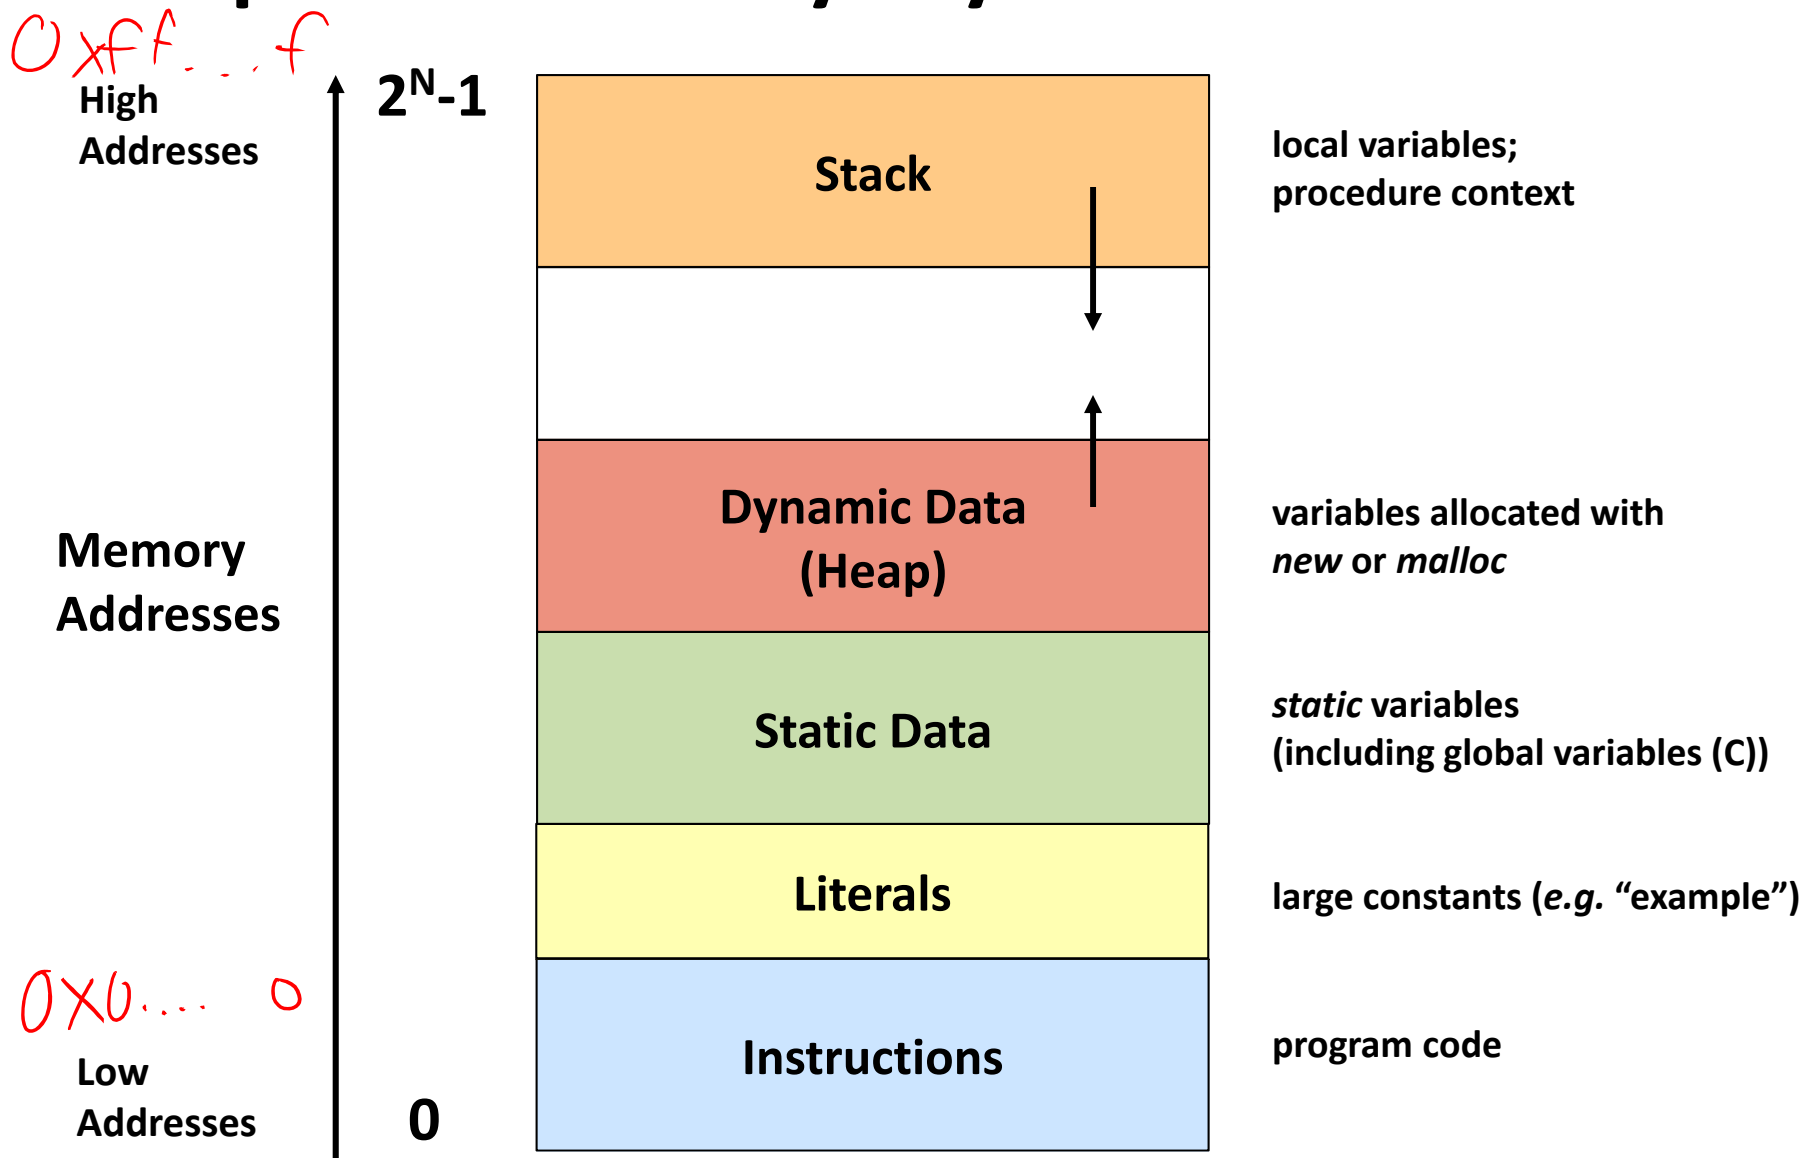
- 1) Passing control
 - To beginning of procedure code
 - Back to return point
- 2) Passing data
 - Procedure arguments
 - Return value
- 3) Memory management
 - Allocate during procedure execution
 - Deallocate upon return
- ❖ All implemented with machine instructions!
 - An x86-64 procedure uses only those mechanisms required for that procedure



Procedures

- ❖ **Stack Structure**
- ❖ Calling Conventions
 - Passing control
 - Passing data
 - Managing local data
- ❖ Register Saving Conventions
- ❖ Illustration of Recursion

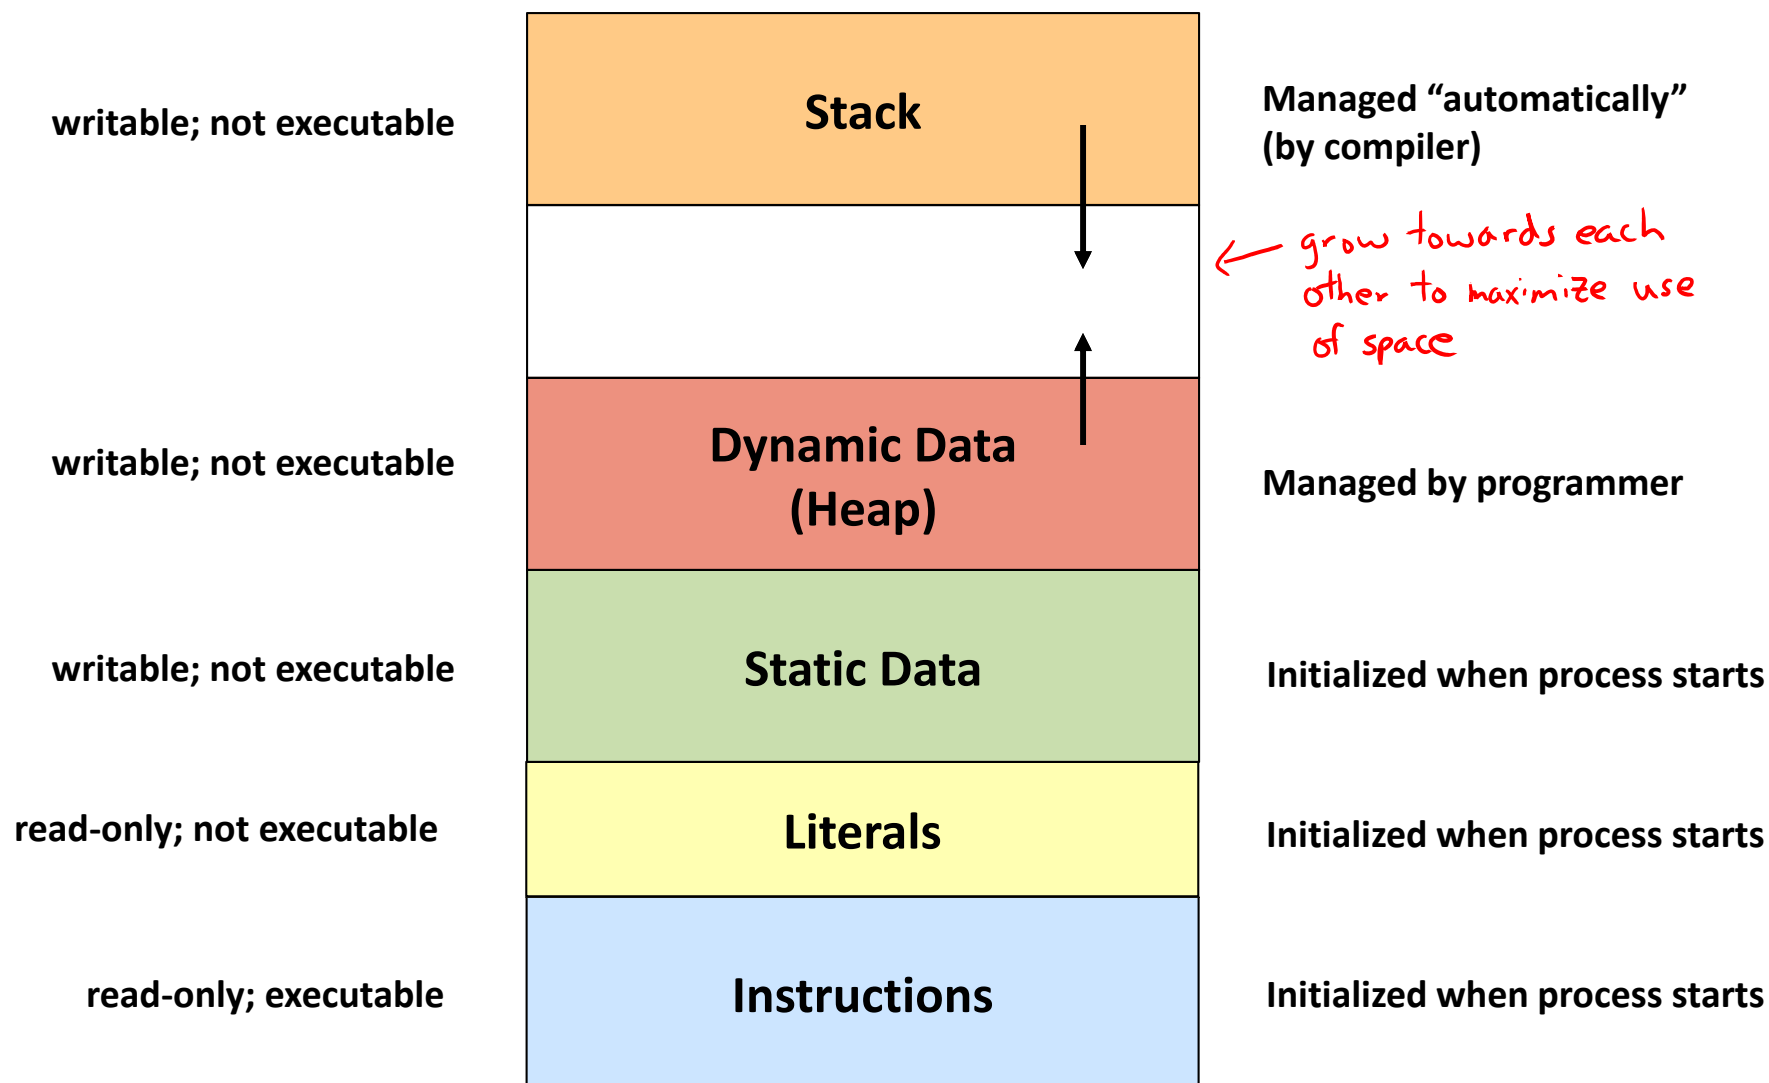
Simplified Memory Layout



Memory Permissions

segmentation faults?

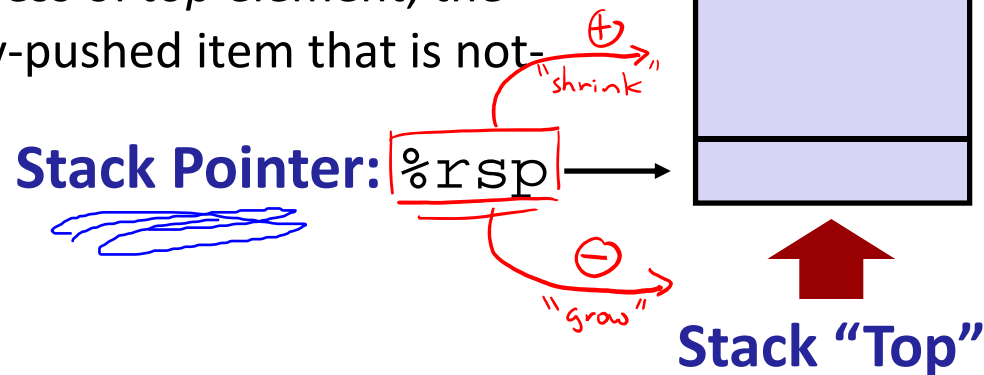
accessing memory in a way that you are not allowed to



x86-64 Stack *Last In, First Out (LIFO)*

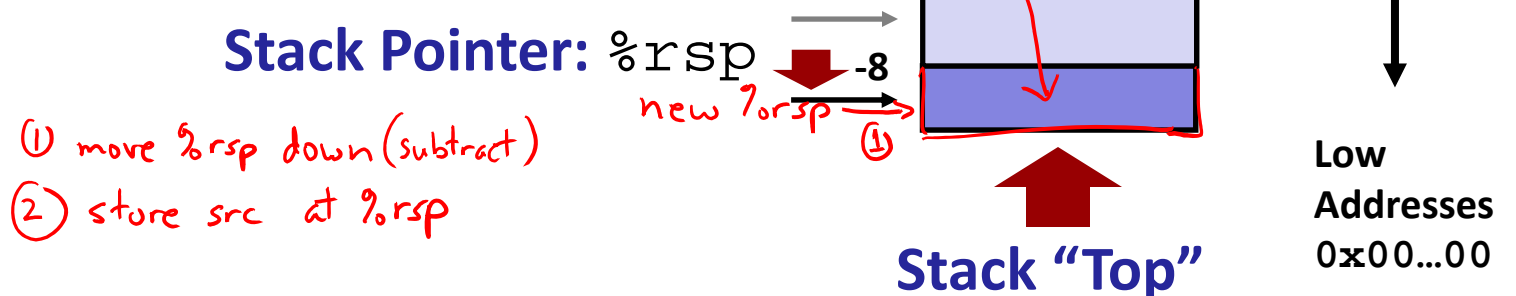
- ❖ Region of memory managed with stack “discipline”
 - Grows toward lower addresses
 - Customarily shown “upside-down”
- ❖ Register `%rsp` contains *lowest* stack address
 - `%rsp` = address of *top* element, the most-recently-pushed item that is not yet-popped

Stack Pointer: `%rsp`



x86-64 Stack: Push

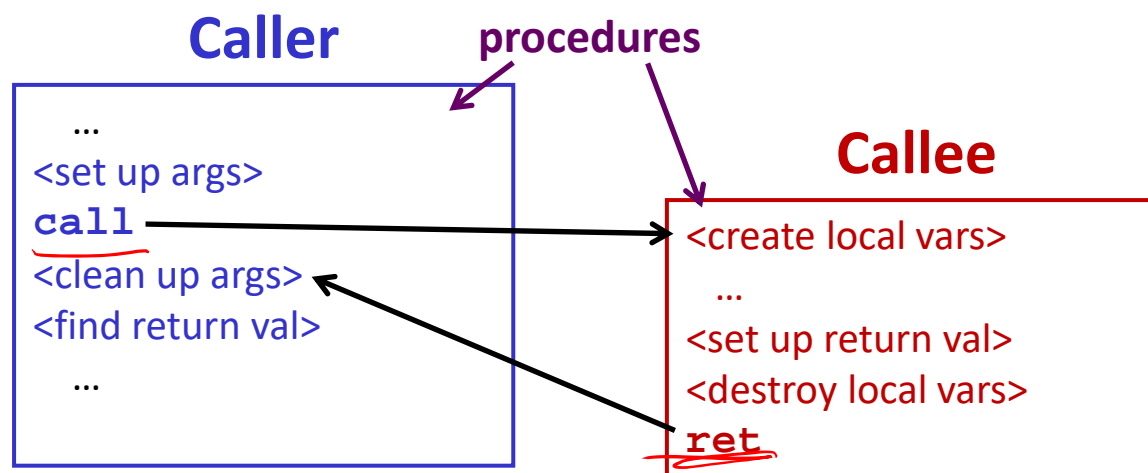
- ❖ pushq src
 - Fetch operand at *src*
 - *Src* can be reg, memory, immediate
 - **Decrement** `%rsp` by 8
 - Store value at address given by `%rsp`
- ❖ Example:
 - pushq %rcx
 - Adjust `%rsp` and store contents of `%rcx` on the stack



Procedures

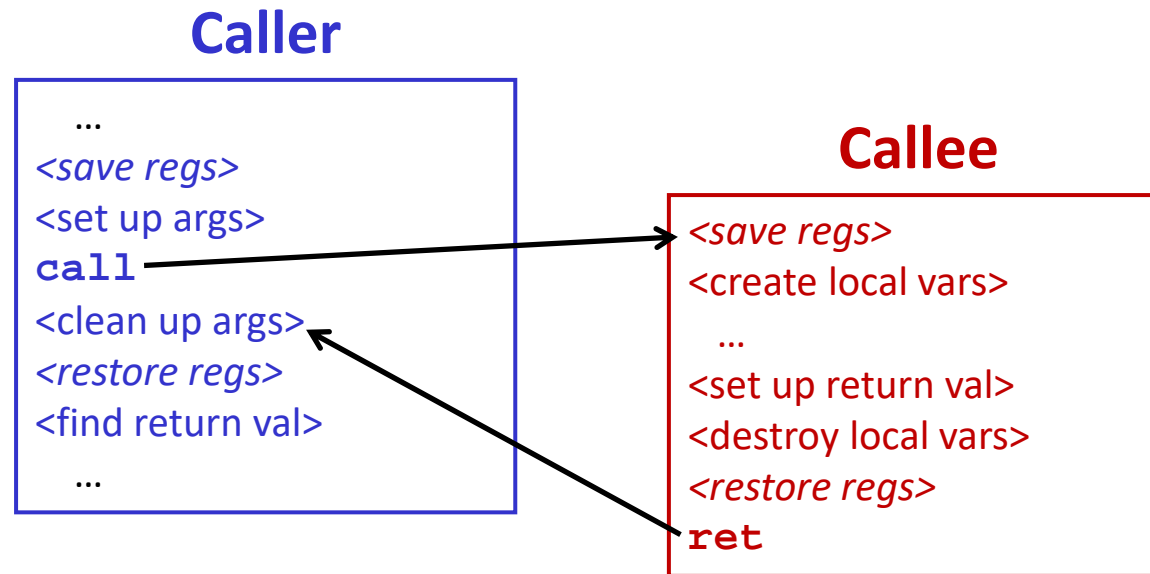
- ❖ Stack Structure
- ❖ **Calling Conventions**
 - **Passing control**
 - Passing data
 - Managing local data
- ❖ Register Saving Conventions
- ❖ Illustration of Recursion

Procedure Call Overview



- ❖ **Callee** must know where to find args
- ❖ **Callee** must know where to find *return address*
- ❖ **Caller** must know where to find *return value*
- ❖ **Caller** and **Callee** run on same CPU, so use the same registers
 - How do we deal with register reuse?
- ❖ Unneeded steps can be skipped (e.g. no arguments)

Procedure Call Overview



- ❖ The convention of where to leave/find things is called the calling convention (or procedure call linkage)
 - Details vary between systems
 - We will see the convention for x86-64/Linux in detail
 - What could happen if our program didn't follow these conventions?

If our program didn't follow these conventions, we might have a case where a callee overwrites some of the variables in the registers that the caller was going to use after the return statement.

Code Example (Preview)

```
void multstore
(long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

Caller

Compiler Explorer:

<https://godbolt.org/g/cKKDZn>

by moving to a register reserved for caller (see textbook page 223)

by pushing it onto the top of the stack

executable disassembly

```
0000000000400540 <multstore>:
400540: push    %rbx           # Save %rbx
400541: movq    %rdx,%rbx      # Save dest
400544: call   400550 <mult2>  # mult2(x,y)
400549: movq    %rax,(%rbx)    # Save at dest
40054c: pop     %rbx           # Restore %rbx
40054d: ret                        # Return
```

Callee

these are instruction addresses

address where instructions for <mult2> are stored

```
long mult2
(long a, long b)
{
    long s = a * b;
    return s;
}
```

```
0000000000400550 <mult2>:
400550: movq    %rdi,%rax      # a
400553: imulq   %rsi,%rax      # a * b
400557: ret                        # Return
```

Procedure Control Flow

❖ Use stack to support procedure call and return

❖ Procedure call: `call label` (special push)

1) Push return address on stack (*why? which address?*)

2) Jump to label

→ ① move `%rsp` down
→ ② store ret addr at `%rsp` (on the stack)
③ `label` → `%rip`

Procedure Control Flow

- ❖ Use stack to support procedure call and return

- ❖ **Procedure call:** `call label` (special push)

1) Push return address on stack (*why? which address?*)

2) Jump to *label*

→ ① move `%rsp` down
→ ② store ret addr at `%rsp`
③ `label` → `%rip`

- ❖ Return address:

- Address of instruction immediately after `call` instruction

- Example from disassembly:

```
400544: call    400550 <mult2>
400549: movq    %rax, (%rbx)
```

Return address = 0x400549

- ❖ **Procedure return:** `ret` (special pop)

1) Pop return address from stack ① read ret addr at `%rsp` into `%rip`

2) Jump to address

② move `%rsp` up

next instruction
happens to be a move,
but could be anything

Procedure Call Example (step 1)

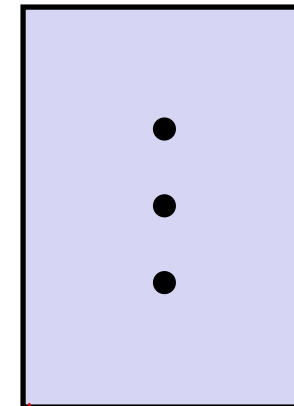
```
0000000000400540 <multstore>:
.
.
400544: call 400550 <mult2>
400549: movq %rax, (%rbx)
.
.
```

return value;
e.g. where
the instructions
are supposed to
pick up again
after call finishes

```
0000000000400550 <mult2>:
400550: movq %rdi, %rax
.
.
400557: ret
```

Memory

0x130
0x128
0x120



return value is
pushed onto
the stack

%rsp

0x120
0x118

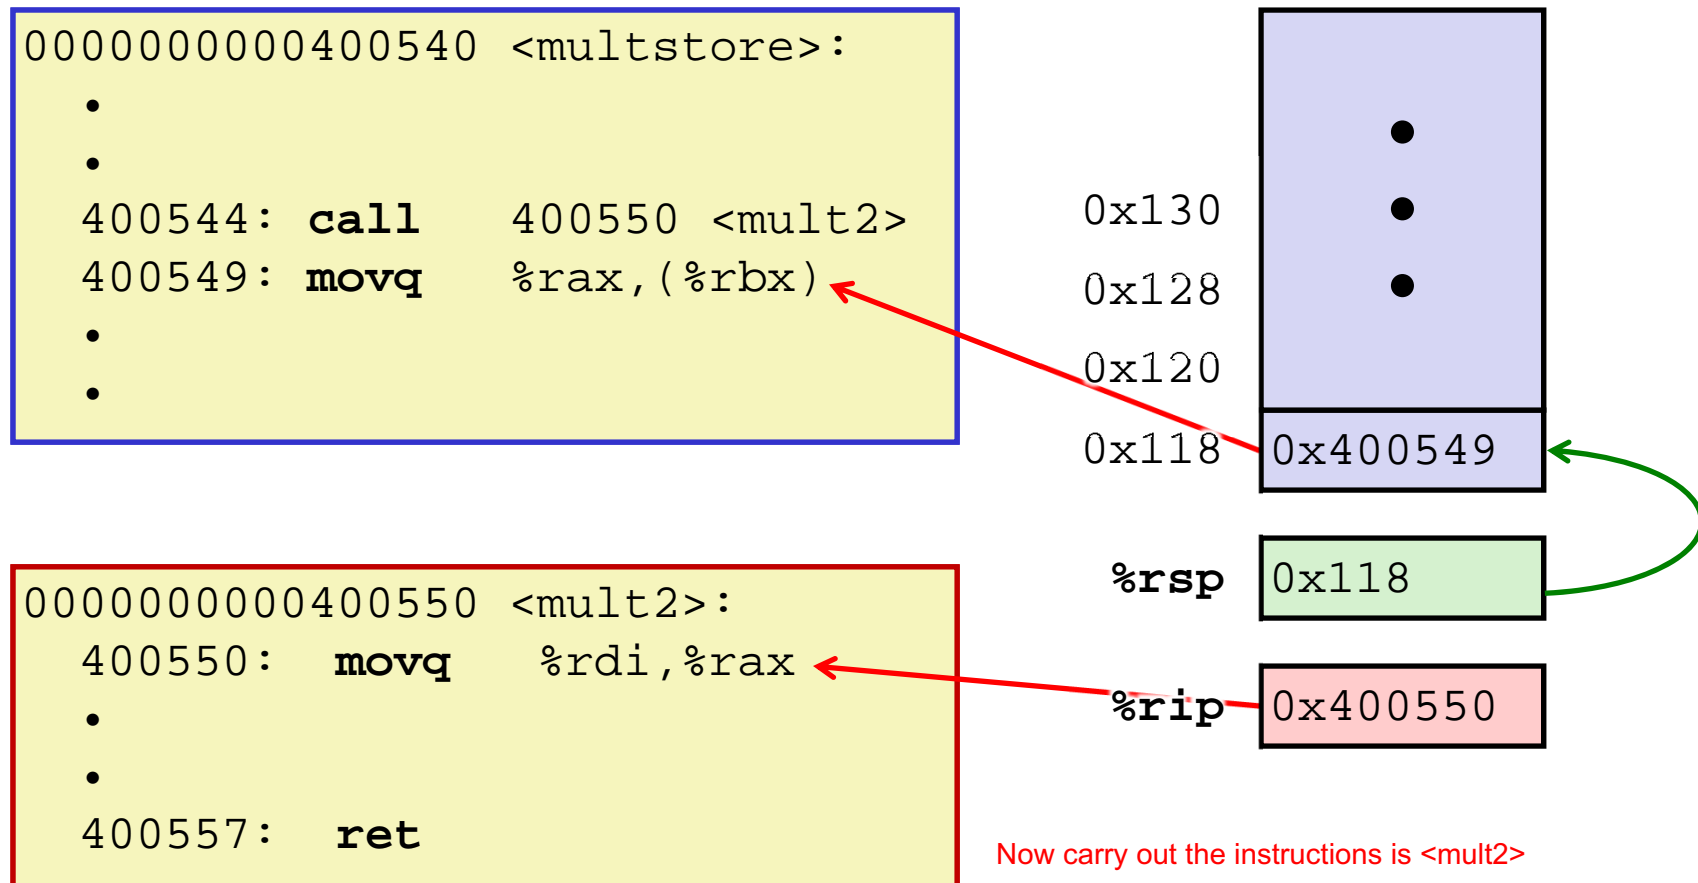
decrement the stack
pointer to now point
at the new top of stack

%rip

0x400544
400550

Adjust the instruction
pointer rip to point to
the new set of instructions
(e.g. where the mult2
instructions start)

Procedure Call Example (step 2)

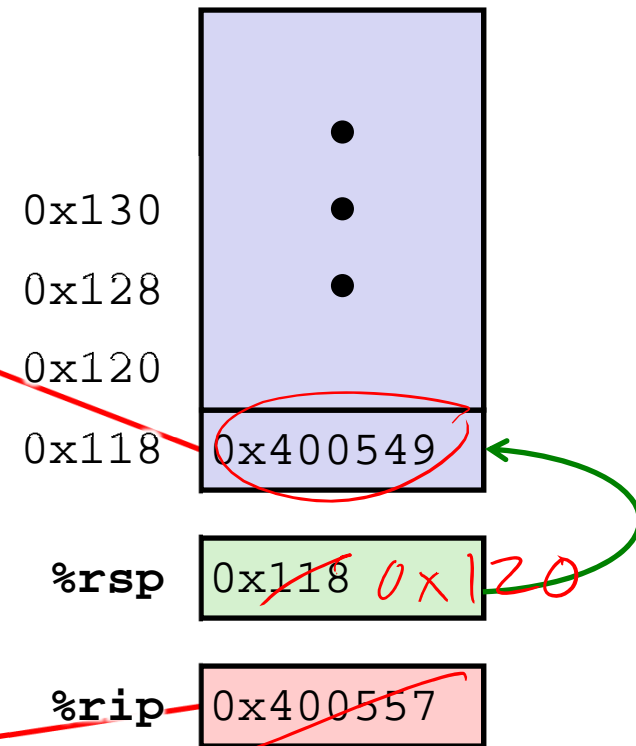


Procedure Return Example (step 1)

```
0000000000400540 <multstore>:  
.  
.  
400544: call    400550 <mult2>  
400549: movq    %rax, (%rbx)  
.  
.
```

```
0000000000400550 <mult2>:  
400550: movq    %rdi, %rax  
.  
.  
400557: ret
```

Memory



400549

Once we complete all the instructions in `<mult2>` (e.g. once we reach the `ret` statement) reset the instruction pointer `%rip` to the next command following the `call` command (e.g. instruction 400549 above). `<mult2>` can find this value as it was pushed onto the stack by `call`.

Procedure Return Example (step 2)

```
00000000000400540 <multstore>:  
.  
.  
400544: call    400550 <mult2>  
400549: movq    %rax, (%rbx)  
.  
.
```

```
00000000000400550 <mult2>:  
400550: movq    %rdi, %rax  
.  
.  
400557: ret
```

Memory

0x130

0x128

0x120

%rsp

0x120

%rip

0x400549

Procedures

- ❖ Stack Structure
- ❖ **Calling Conventions**
 - Passing control
 - **Passing data**
 - Managing local data
- ❖ Register Saving Conventions
- ❖ Illustration of Recursion

So we have seen how we can use call/ret to move the instruction pointer around. This is how we pass control to different functions/instructions stored in memory. Now, how do we pass data (e.g. the arguments of a function)?

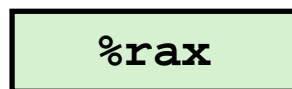
Procedure Data Flow

Registers (**NOT** in Memory)

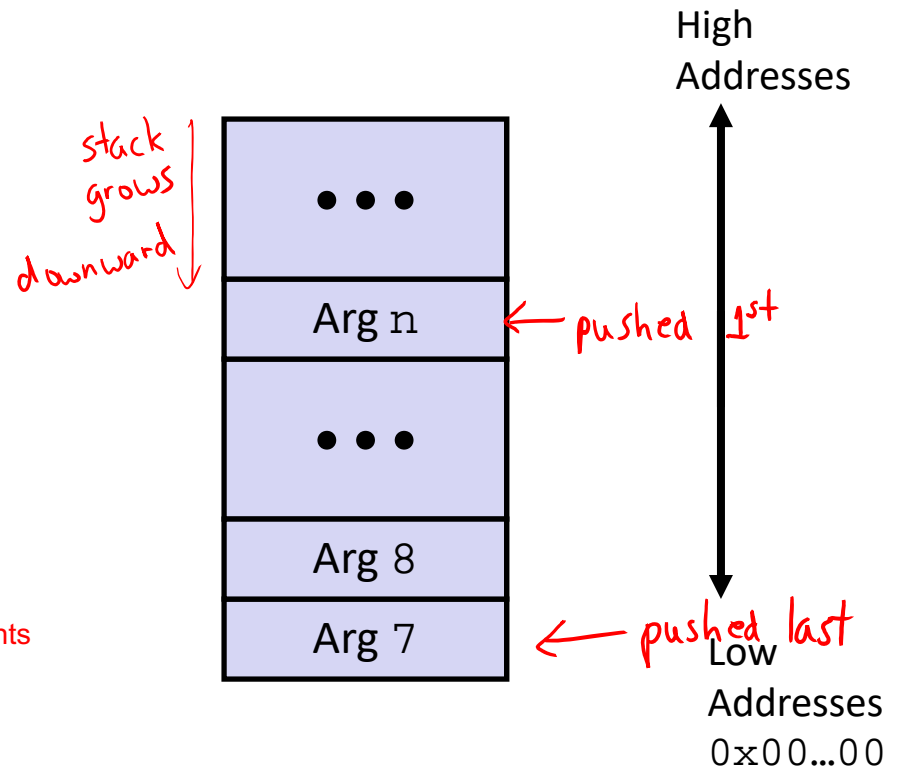
❖ First 6 arguments



❖ Return value



Stack (**M**emory)



- Only allocate stack space when needed

x86-64 Return Values

- ❖ By convention, values returned by procedures are placed in %rax

- Choice of %rax is arbitrary

- 1) **Caller** must make sure to save the contents of %rax before calling a **callee** that returns a value

save to the stack

because the callee is going to overwrite that stuff to store the return value there

- Part of register-saving convention

- 2) **Callee** places return value into %rax

- Any type that can fit in 8 bytes – integer, float, pointer, etc.
- For return values greater than 8 bytes, best to return a *pointer* to them

- 3) Upon return, **caller** finds the return value in %rax

Data Flow Examples

caller

```
void multstore
(long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

lined up nicely so we didn't have
to manipulate arguments

```
00000000000400540 <multstore>:
    # x in %rdi, y in %rsi, dest in %rdx
    ...
400541: movq    %rdx,%rbx        # "Save" dest
400544: call    400550 <mult2>    # mult2(x,y)
    # t in %rax
400549: movq    %rax,(%rbx)      # Save at dest
    ...
```

(will
explain
later)

callee

```
long mult2
(long a, long b)
{
    long s = a * b;
    return s;
}
```

```
00000000000400550 <mult2>:
    # a in %rdi, b in %rsi
400550: movq    %rdi,%rax        # a
400553: imulq   %rsi,%rax        # a * b
    # s in %rax
400557: ret                      # Return
```


Procedures

- ❖ Stack Structure
- ❖ **Calling Conventions**
 - Passing control
 - Passing data
 - **Managing local data**
- ❖ Register Saving Conventions
- ❖ Illustration of Recursion

Stack-Based Languages

- ❖ Languages that support recursion
 - e.g. C, Java, most modern languages
 - Code must be re-entrant
 - Multiple simultaneous instantiations of single procedure
 - Need some place to store *state* of each instantiation
 - Arguments, local variables, return pointer *address*
- ❖ Stack allocated in frames
 - State for a single procedure instantiation
- ❖ Stack discipline
 - State for a given procedure needed for a limited time
 - Starting from when it is called to when it returns
 - Callee always returns before caller does

Call Chain Example

```
yoo (...)  
{  
  •  
  •  
  who ( ) ;  
  •  
  •  
}
```

```
who (...)  
{  
  •  
  ① amI ( ) ;  
  •  
  ② amI ( ) ;  
  •  
}
```

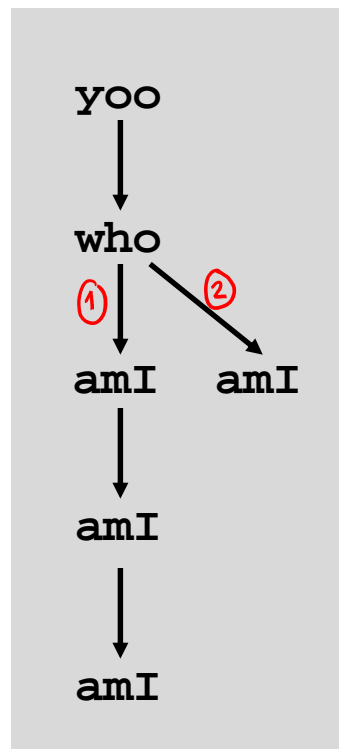
```
amI (...)  
{  
  •  
  if (...) {  
    amI ()  
  }  
  •  
}
```

1st call recurses twice

2nd call
doesn't recurse

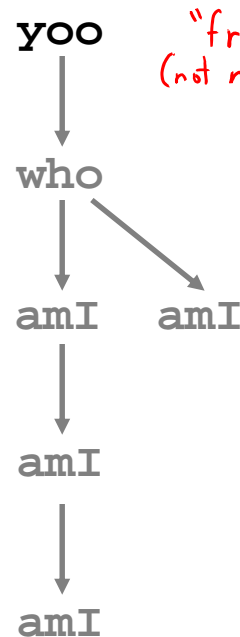
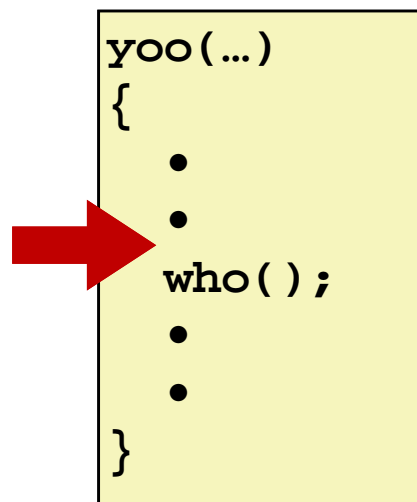
based on
condition

Example
Call Chain



Procedure amI is recursive
(calls itself)

1) Call to yoo



"frame pointer"
(not necessary)

`%rbp`

`%rsp`

Stack

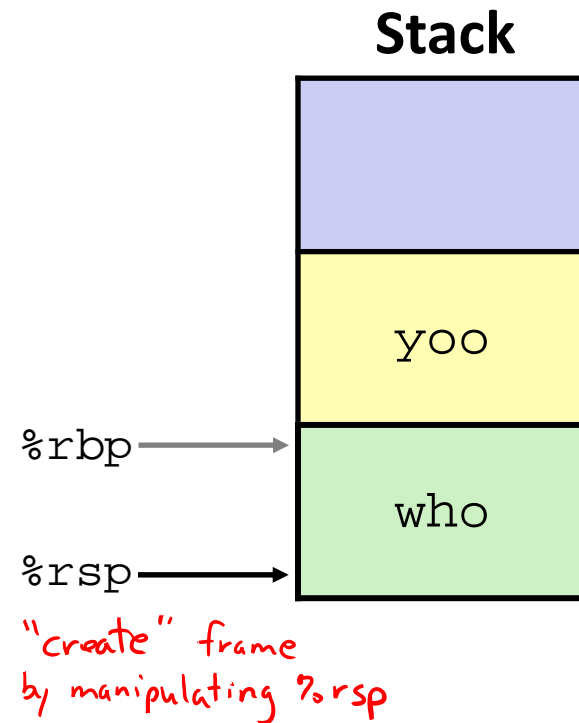
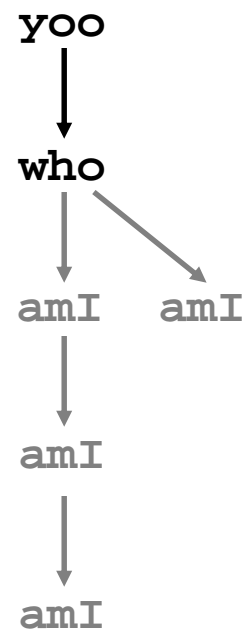
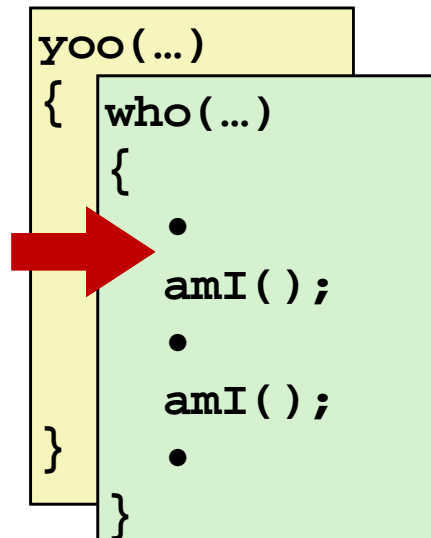


high
addr

could be any
procedure
that calls
yoo

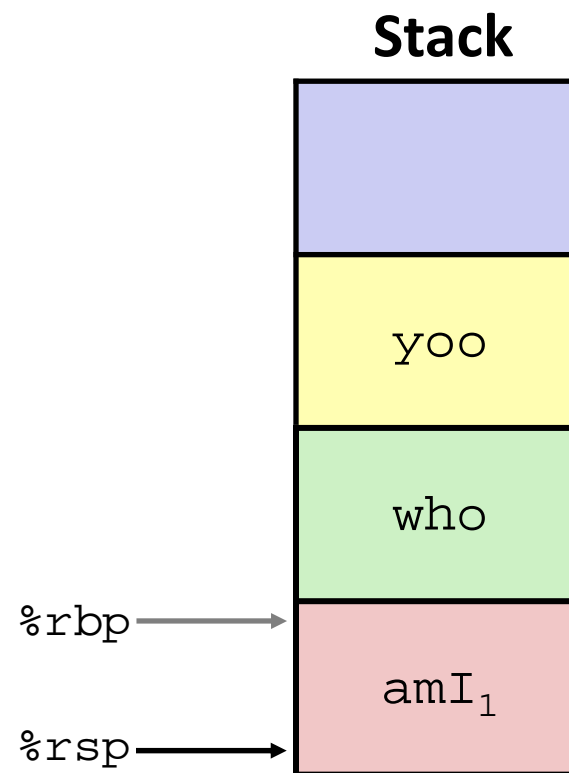
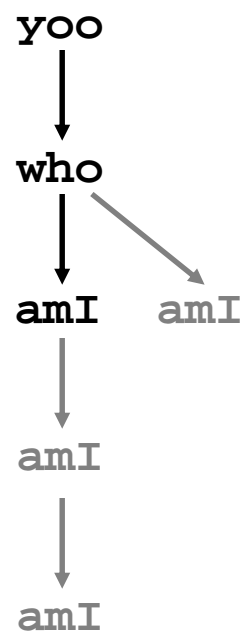
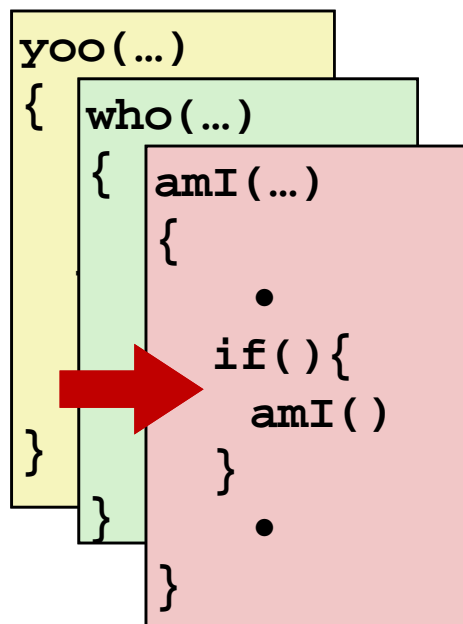
low
addr

2) Call to who

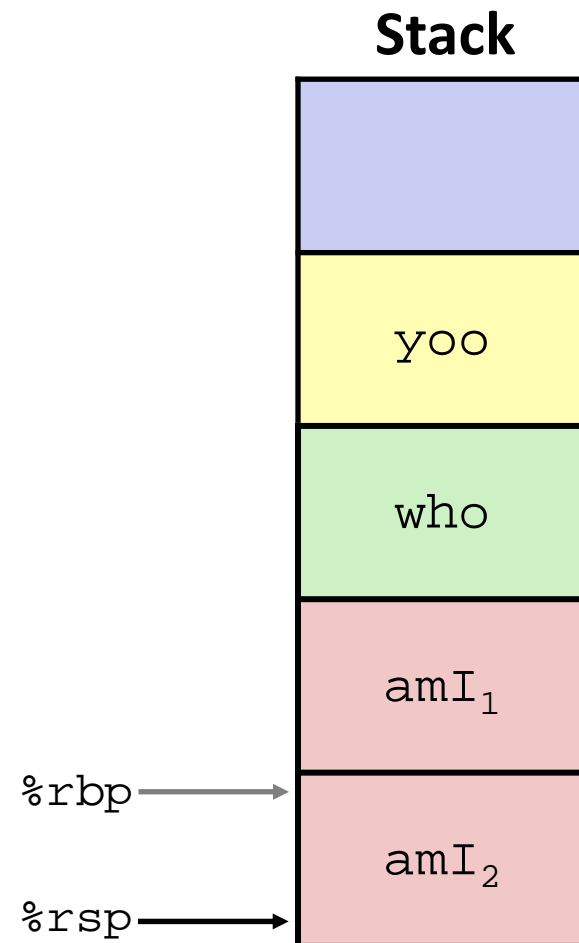
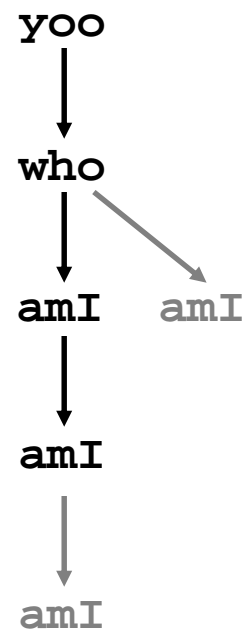
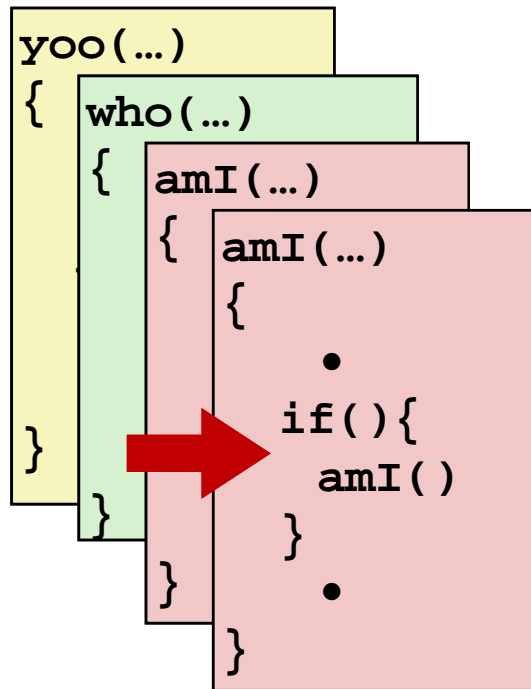


This frame will be allocated for storing local information about the who procedure

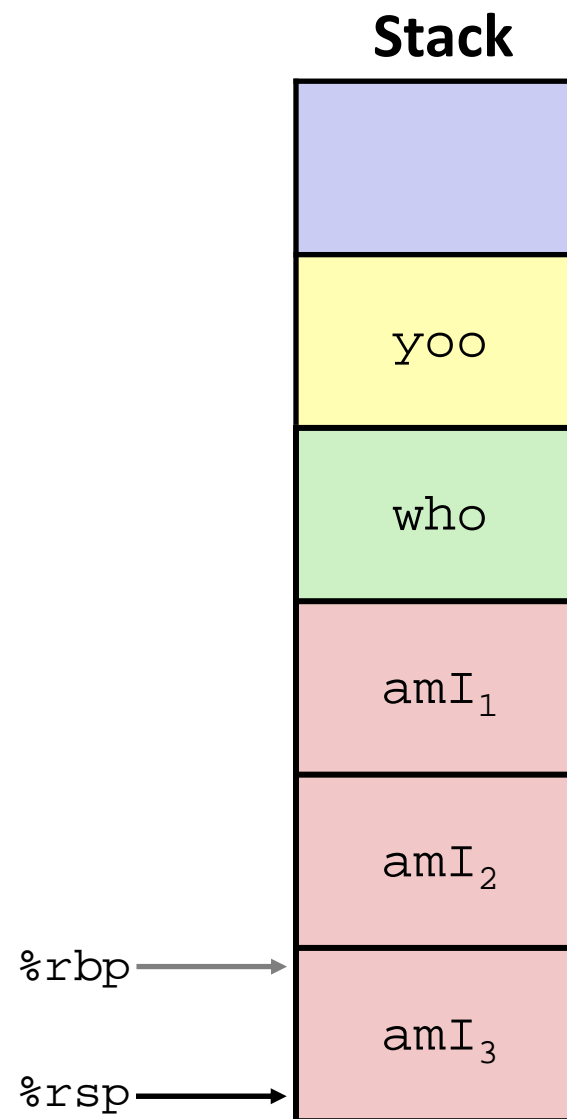
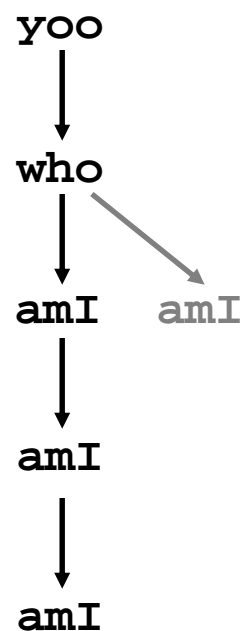
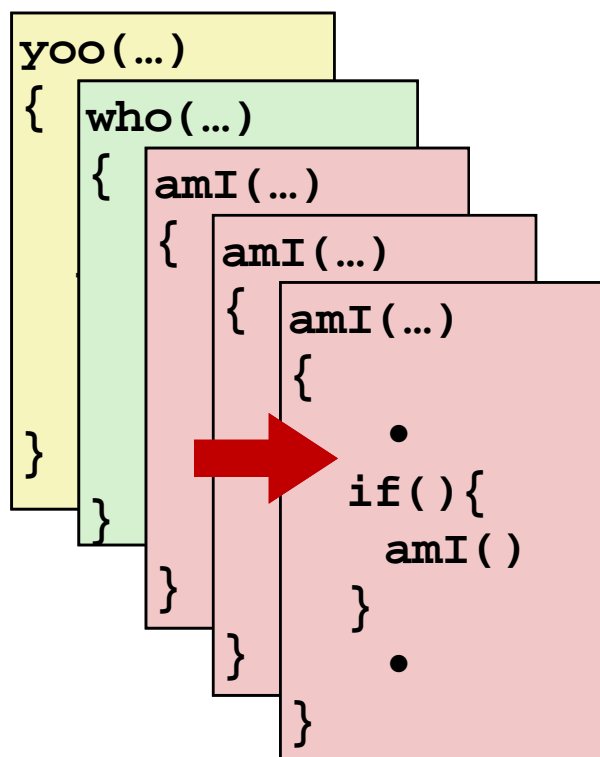
3) Call to amI (1)



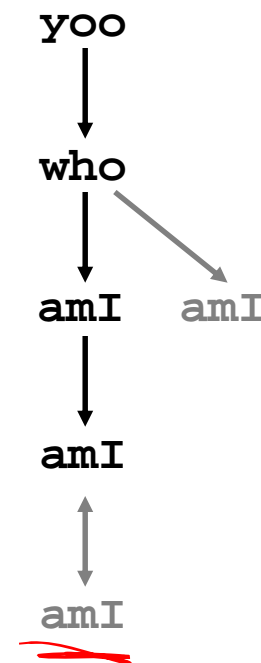
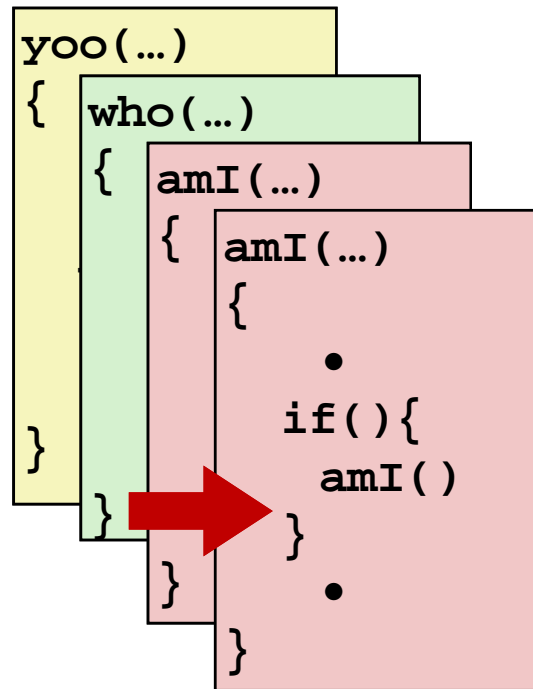
4) Recursive call to amI (2)



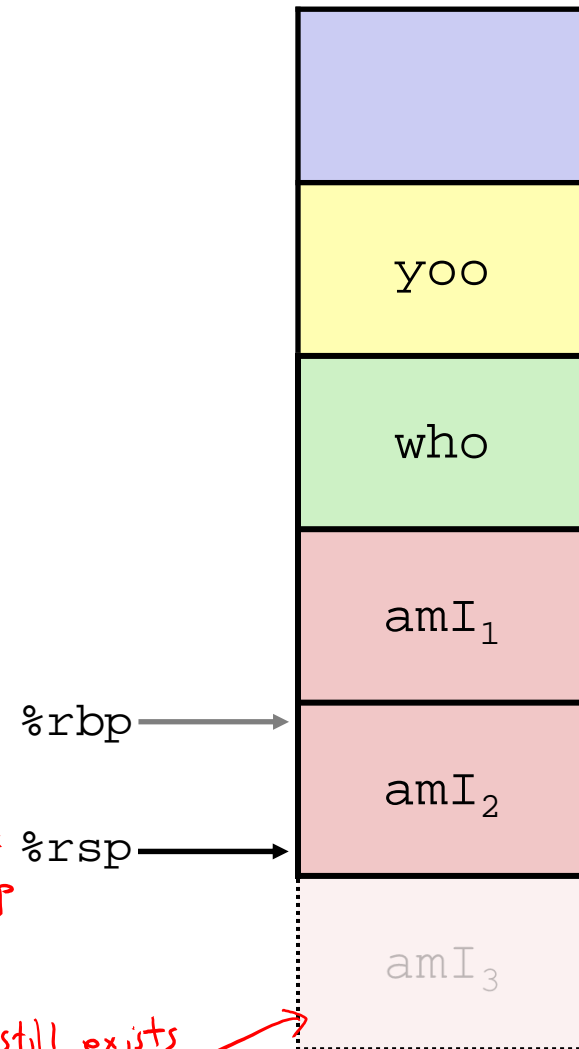
5) (another) Recursive call to amI (3)



6) Return from (another) recursive call to amI



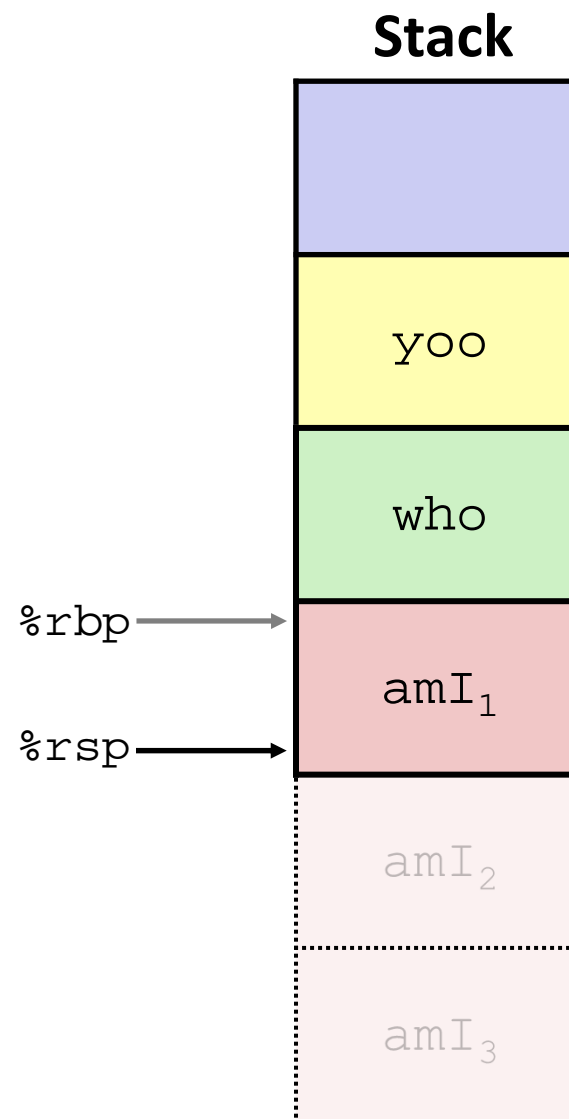
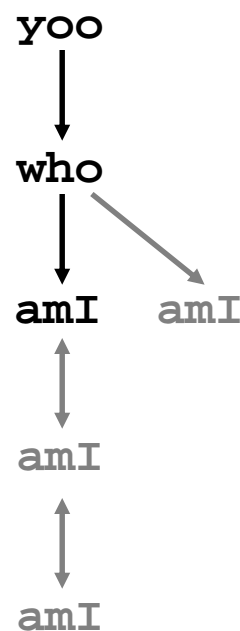
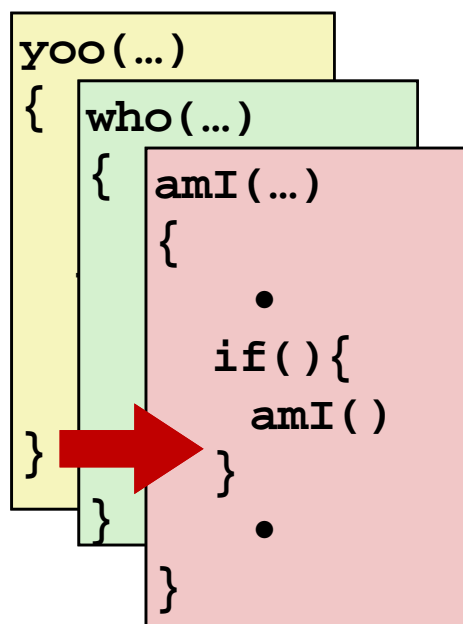
Stack



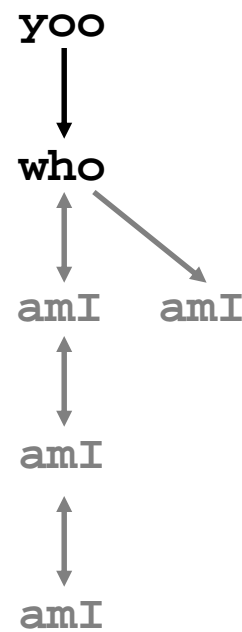
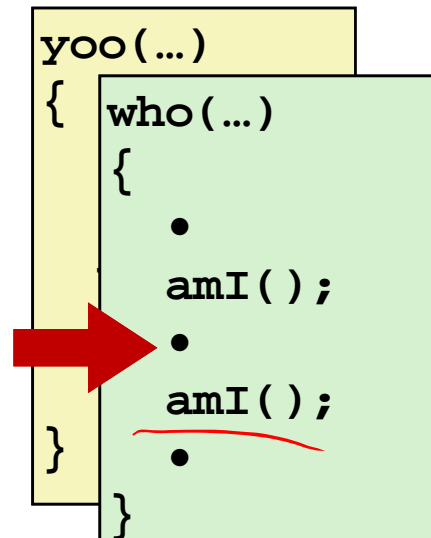
"deallocate" stack frame by moving `%rsp` back up

data still exists, but you shouldn't use it because at any point of time we could overwrite it; since we plan on overwriting it anyways, no point in just "zeroing out" all the data here.

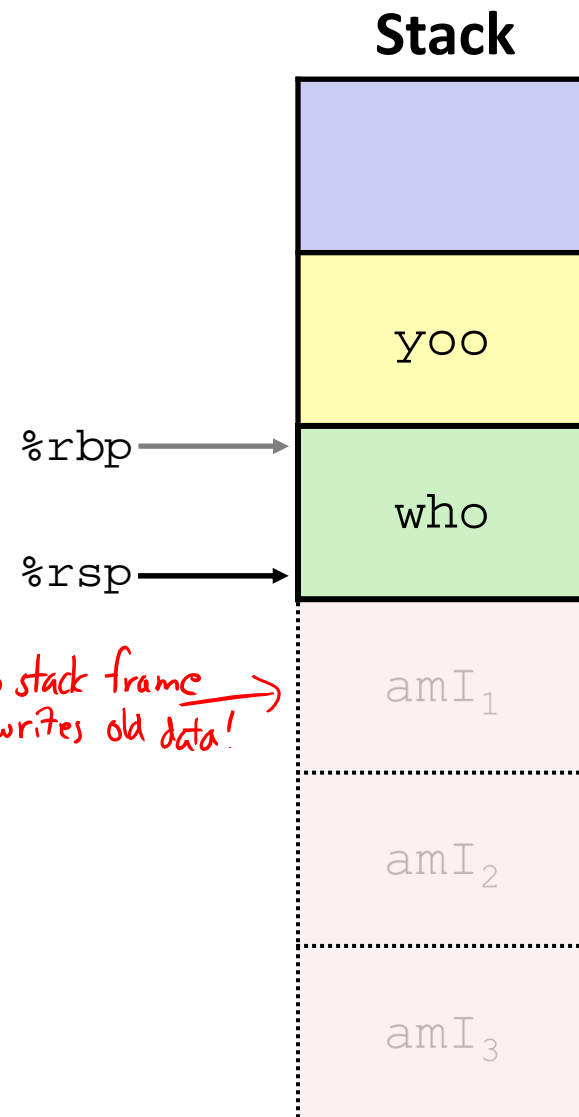
7) Return from recursive call to amI



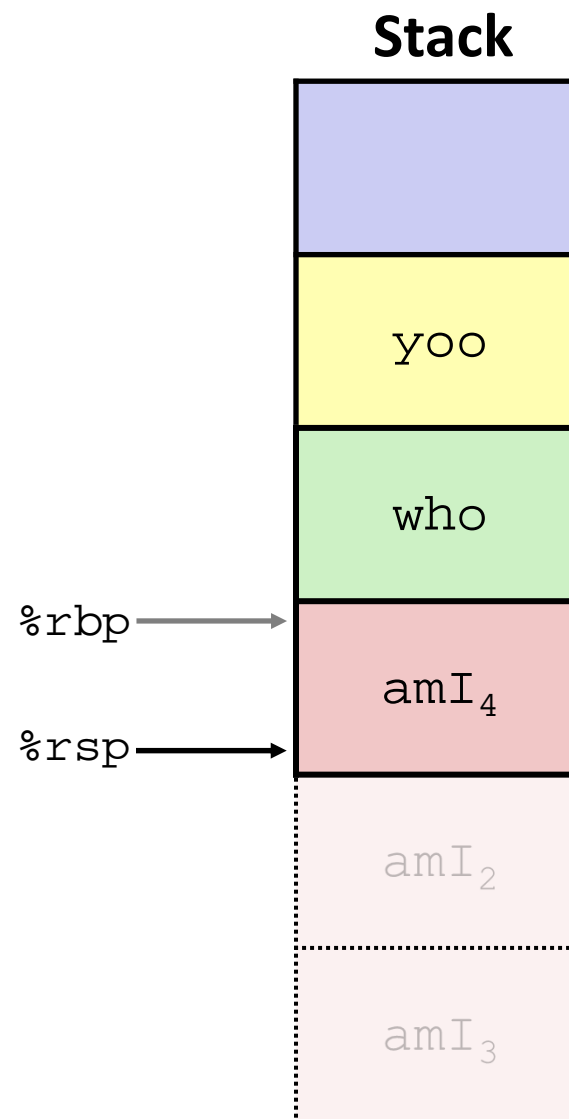
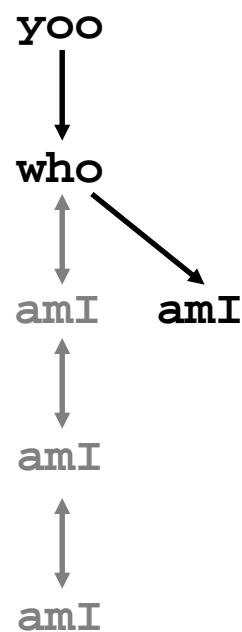
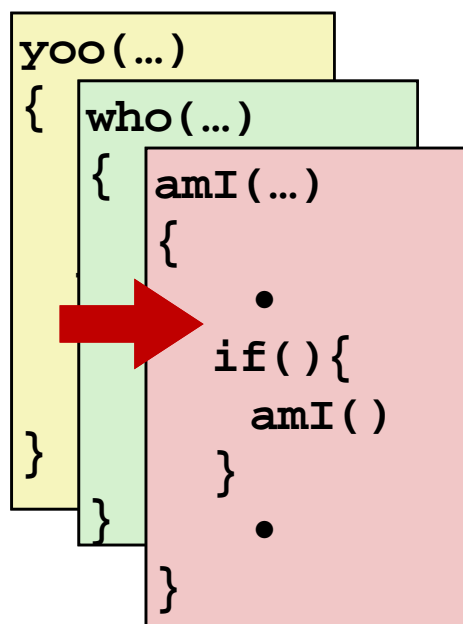
8) Return from call to amI



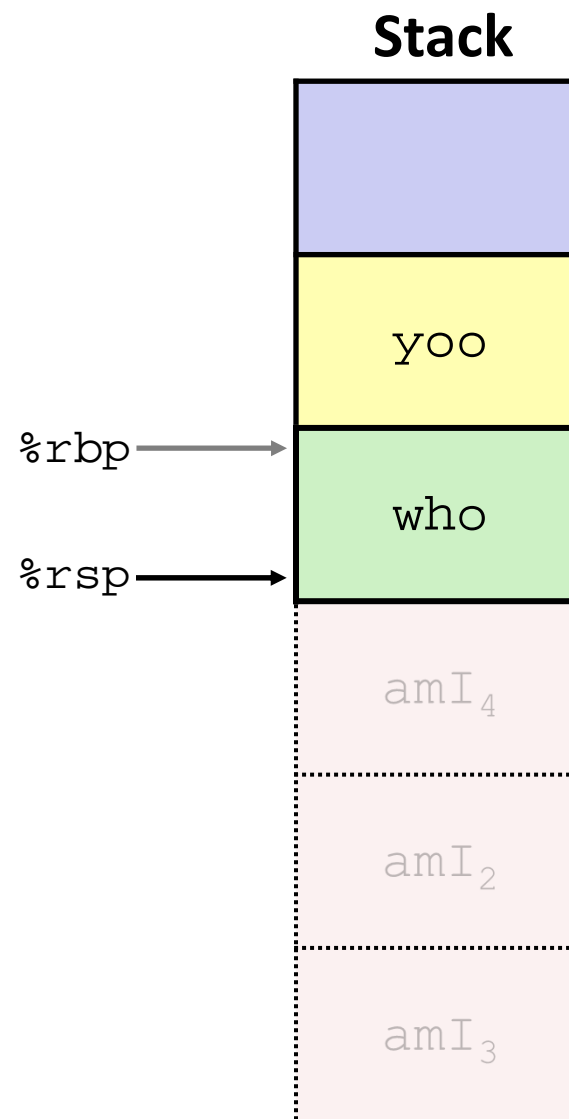
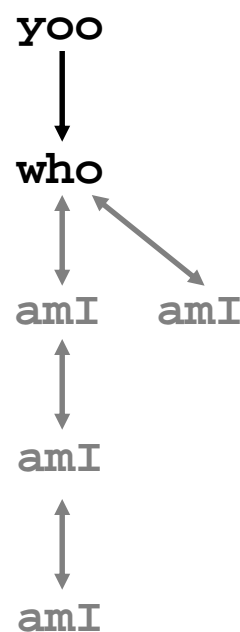
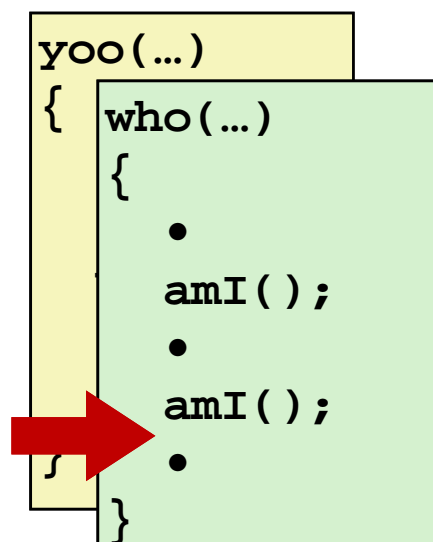
*new stack frame
overwrites old data!*



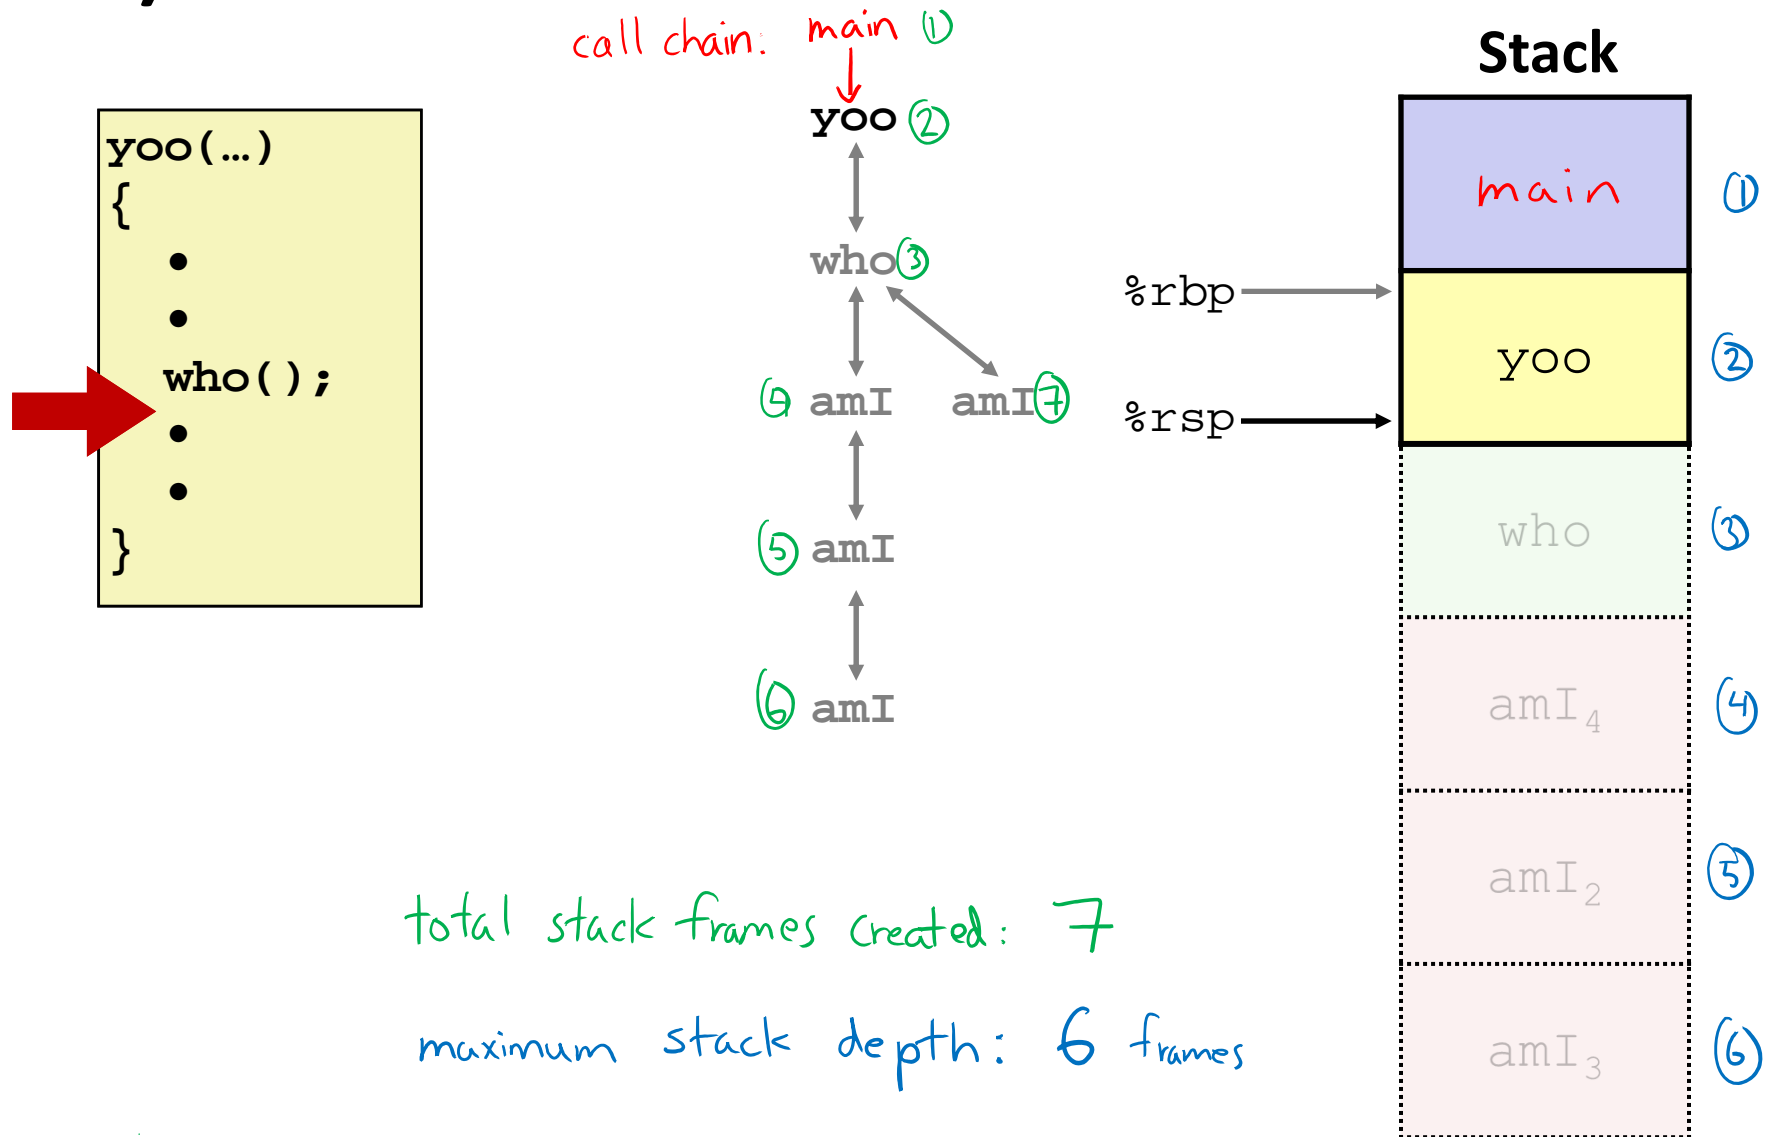
9) (second) Call to amI (4)



10) Return from (second) call to amI



11) Return from call to who



x86-64/Linux Stack Frame

❖ Caller's Stack Frame

remember some of the arguments can be stored in the registers

- Extra arguments (if > 6 args) for this call

❖ Current/Callee Stack Frame

- Return address
 - Pushed by `call` instruction
- Old frame pointer (optional)
- Saved register context (when reusing registers)
- Local variables (If can't be kept in registers)
- "Argument build" area (If callee needs to call another function - parameters for function about to call, if needed)

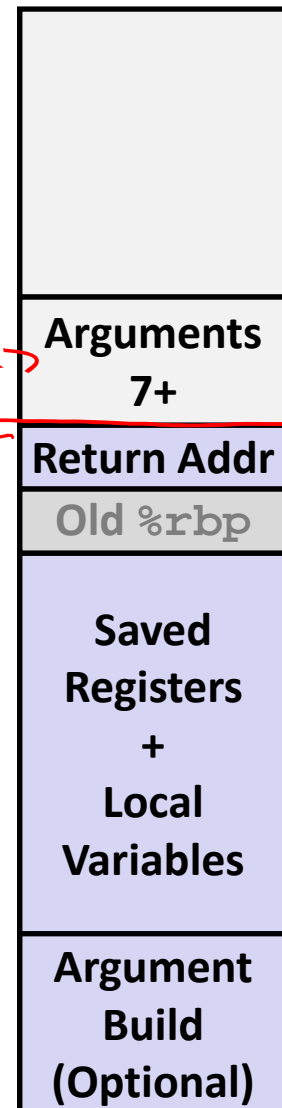
all frames set up similarly
(this is the "argument build")

Frame pointer
`%rbp`
(Optional)

Caller Frame

Callee Frame

Stack pointer
`%rsp`



e.g. so that we don't overwrite arguments stored in the register by the caller (we save these to stack so we can replace the values in register at the end)

Peer Instruction Question

Vote only on 3rd question at

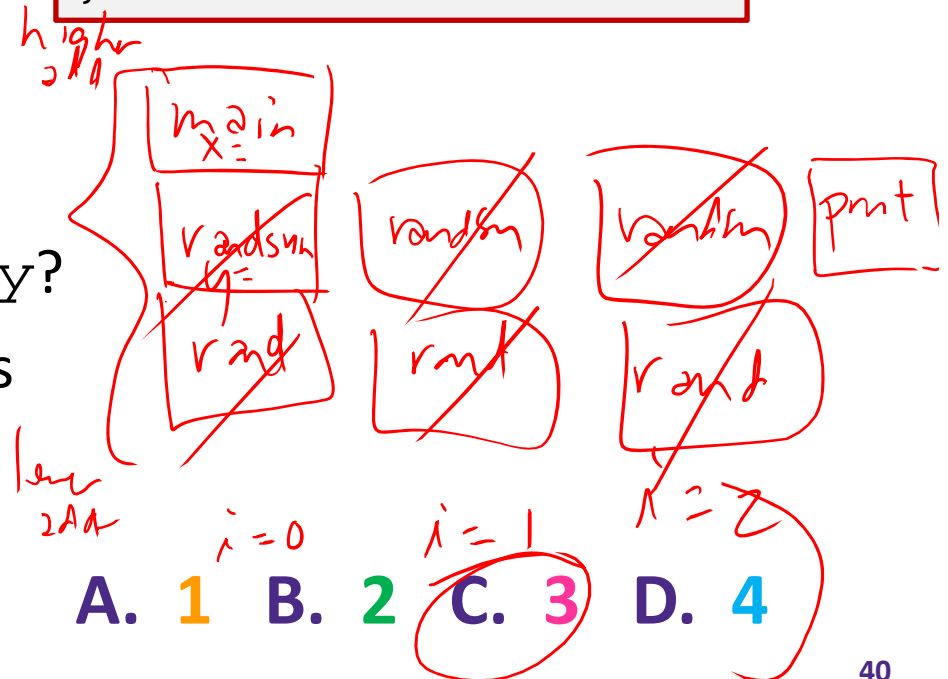
<http://pollev.com/rea>

- ❖ Answer the following questions about when `main()` is run (assume `x` and `y` stored on the Stack):

```
int main() {
    int i, x = 0;
    for(i = 0; i < 3; i++)
        x = randSum(x);
    printf("x = %d\n", x);
    return 0;
}
```

```
int randSum(int n) {
    int y = rand() % 20;
    return n + y;
}
```

- Higher/larger address: `x` or `y`?
- How many total stack frames are *created*? 8
- What is the maximum *depth* (# of frames) of the Stack?



A. 1 B. 2 C. 3 D. 4