

Executables & Arrays

CSE 351 Spring 2019

Instructor:

Ruth Anderson

Teaching Assistants:

Gavin Cai

Jack Eggleston

John Feltrup

Britt Henderson

Richard Jiang

Jack Skaltzky


Sophie Tian

Connie Wang

Sam Wolfson

Casey Xing

Chin Yeoh



<http://xkcd.com/1270/>

Administrivia

- ❖ Lab 2 (x86-64) due Wednesday (5/01)
- ❖ Homework 3, due Wednesday (5/8)
 - On midterm material, but due after the midterm
- ❖ **Midterm** (Fri 5/03, 4:30-5:30pm in KNE 130)
 - No lecture on Friday 5/03
 - Ruth will hold office hours instead
 - Fri 11:30am-12:30pm in CSE 460
 - Fri 2:30-3:30pm in CSE 460

Roadmap

C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

Memory & data
Integers & floats
x86 assembly
Procedures & stacks
Executables
Arrays & structs
Memory & caches
Processes
Virtual memory
Memory allocation
Java vs. C

Assembly
language:

```
get_mpg:  
    pushq    %rbp  
    movq     %rsp, %rbp  
    ...  
    popq     %rbp  
    ret
```


Machine
code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
110000011111101000011111
```

OS:



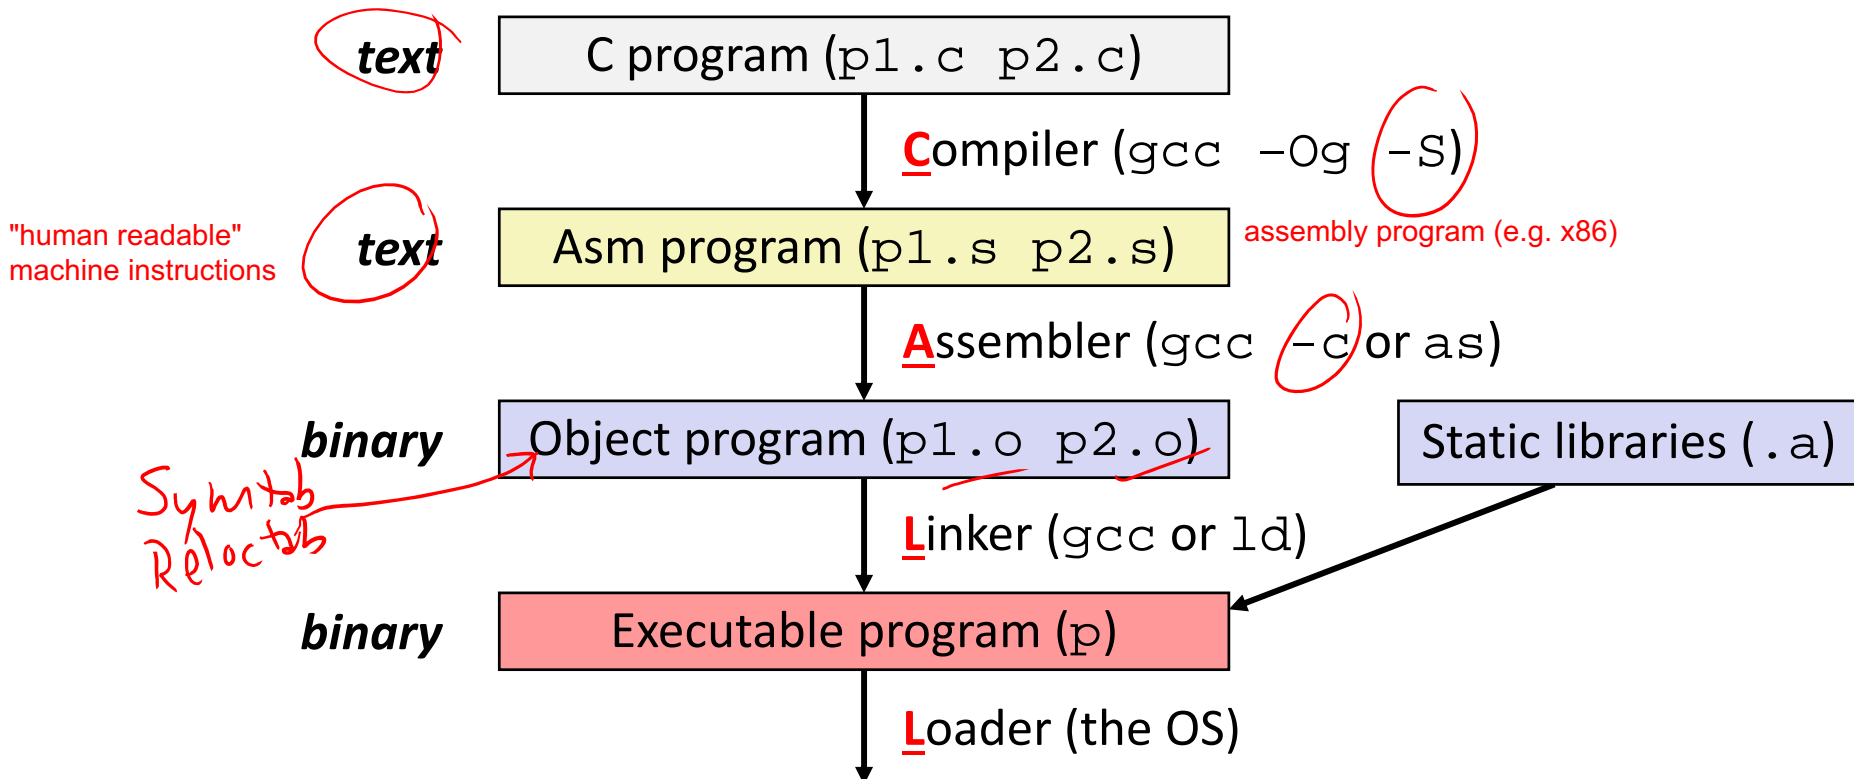
Computer
system:




Building an Executable from a C File

- ❖ Code in files p1.c p2.c
- ❖ Compile with command: `gcc -Og p1.c p2.c -o p`
 - Put resulting machine code in file p
- ❖ Run with command: `./p`

multiple code files that will be linked together



Compiler

- ❖ **Input:** Higher-level language code (*e.g.* C, Java)
 - `foo.c`
- ❖ **Output:** Assembly language code (*e.g.* x86, ARM, MIPS)
 - `foo.s` 
- ❖ First there's a preprocessor step to handle `#directives`
 - Macro substitution, plus other specialty directives
 - If curious/interested: <http://tigrcc.ticalc.org/doc/cpp.html>
- ❖ Super complex, whole courses devoted to these!
- ❖ Compiler optimizations
 - “Level” of optimization specified by capital ‘O’ flag (*e.g.* `-Og`, `-O3`)
 - Options: <https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>

Compiling Into Assembly

❖ C Code (sum.c)

```
void sumstore(long x, long y, long *dest) {  
    long t = x + y;  
    *dest = t;  
}
```

❖ x86-64 assembly (gcc -Og -S sum.c)

```
sumstore(long, long, long*):  
    addq    %rdi, %rsi  
    movq    %rsi, (%rdx)  
    ret
```

Warning: You may get different results with other versions of gcc and different compiler settings

Assembler

- ❖ **Input:** Assembly language code (*e.g.* x86, ARM, MIPS)
 - `foo.s`
- ❖ **Output:** Object files (*e.g.* ELF, COFF)
 - `foo.o`
 - Contains *object code* and *information tables*
- ❖ Reads and uses *assembly directives*
 - *e.g.* `.text`, `.data`, `.quad`
 - x86: https://docs.oracle.com/cd/E26502_01/html/E28388/eoiyg.html
- ❖ Produces “machine language”
 - ★ Does its best, but object file is *not* a completed binary
- ❖ Example: `gcc -c foo.s`

Producing Machine Language

Assembler

- ❖ **Simple cases:** *addq %rdi, %rsi* arithmetic and logical operations, shifts, etc.
 - All necessary information is contained in the instruction itself
- ❖ What about the following?
 - Conditional jump *addr/label*
 - Accessing static data (e.g. global var or jump table)
 - call *addr/label*
- ❖ **Addresses and labels are problematic because the final executable hasn't been constructed yet!**
 - So how do we deal with these in the meantime?

Object File Information Tables

Assembler
puts in
object file

- ❖ **Symbol Table** holds list of “items” that may be used by other files
“What I have”
 - *Non-local labels* – function names for `call`
 - *Static Data* – variables & literals that might be accessed across files
- ❖ **Relocation Table** holds list of “items” that this file needs the address of later (currently undetermined)
 - Any *label* or piece of *static data* referenced in an instruction in this file
 - Both internal and external“What I need”
- ❖ Each file has its own symbol and relocation tables

Object File Format

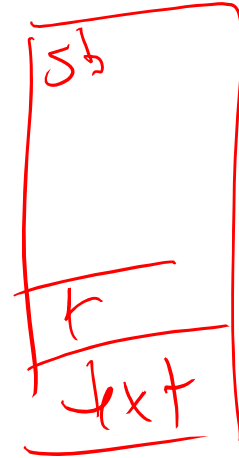
- 1) object file header: size and position of the other pieces of the object file *"table of contents"*
- 2) text segment: the machine code *(Instructions)*
- 3) data segment: data in the source file (binary) *(Static Data & Literals)*
- 4) relocation table: identifies lines of code that need to be "handled"
- 5) symbol table: list of this file's labels and data that can be referenced
- 6) debugging information *(info for GDB)*

❖ More info: ELF format

- http://www.skyfree.org/linux/references/ELF_Format.pdf

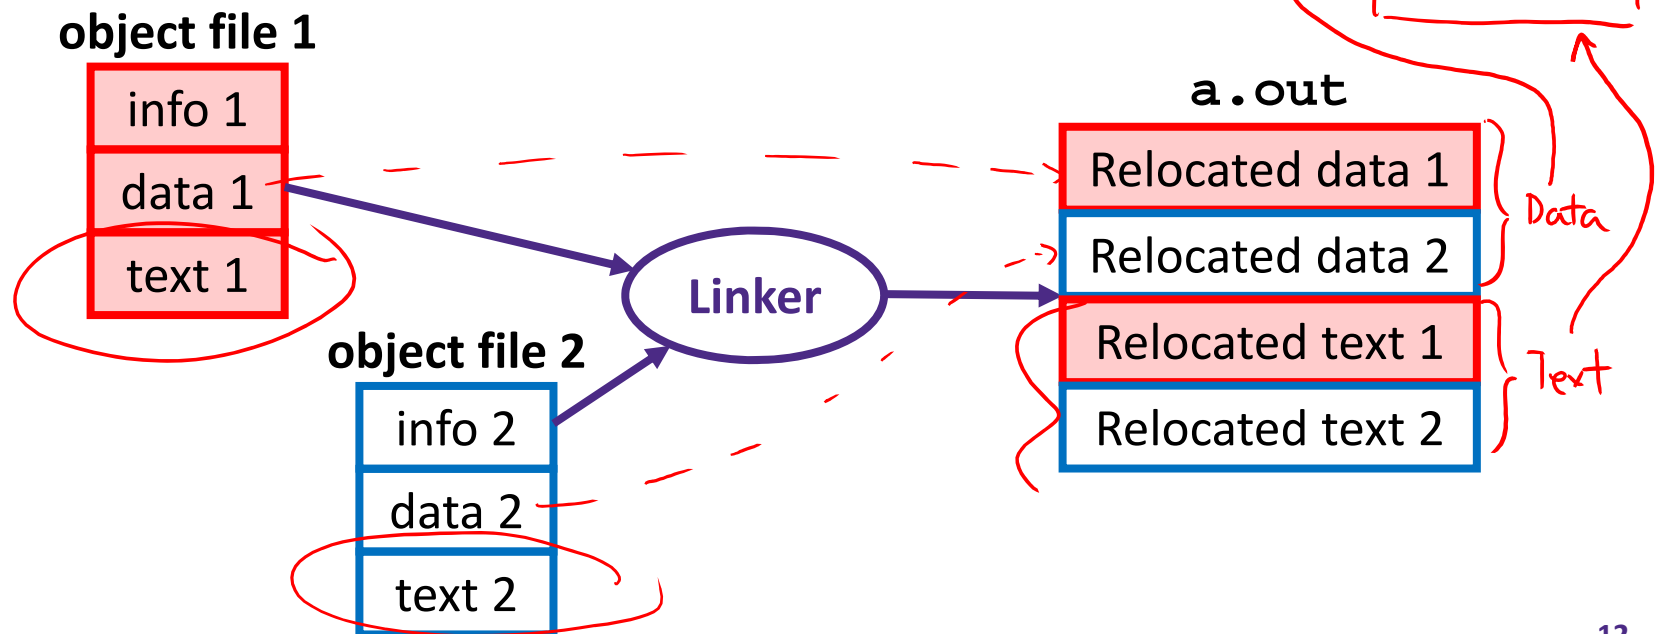
Linker

- ❖ **Input:** Object files (e.g. ELF, COFF)
 - `foo.o`
- ❖ **Output:** executable binary program
 - `a.out`
- ❖ Combines several object files into a single executable (*linking*)
- ❖ Enables separate compilation/assembling of files
 - Changes to one file do not require recompiling of whole program



Linking

- 1) Take text segment from each .o file and put them together
- 2) Take data segment from each .o file, put them together, and concatenate this onto end of text segments
- 3) Resolve References
 - Go through Relocation Table; handle each entry



Disassembling Object Code

❖ Disassembled:

00000000000400536	<sumstore>:
400536:	48 01 fe add %rdi,%rsi
400539:	48 89 32 mov %rsi, (%rdx)
40053c:	c3 retq

address of instruction object code bytes (hex) interpreted assembly instructions

❖ Disassembler (objdump -d sum)

- Useful tool for examining object code (man 1 objdump)
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can run on either a .out (complete executable) or .o file

What Can be Disassembled?

```
% objdump -d WINWORD.EXE
```

```
WINWORD.EXE:      file format pei-i386
```

```
No symbols in "WINWORD.EXE".
```

```
Disassembly of section .text:
```

```
30001000 <.text>:
```

```
30001000:
```

```
30001001:
```

```
30001003:
```

```
30001005:
```

```
3000100a:
```

**Reverse engineering forbidden by
Microsoft End User License Agreement**

- ❖ Anything that can be interpreted as executable code
- ❖ Disassembler examines bytes and attempts to reconstruct assembly source

Loader

- ❖ **Input:** executable binary program, command-line arguments
 - `./a.out arg1 arg2`
- ❖ **Output:** <program is run>
- ❖ Loader duties primarily handled by OS/kernel
 - More about this when we learn about processes
- ❖ Memory sections (Instructions, Static Data, Stack) are set up
- ❖ Registers are initialized

Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

Assembly
language:

```
get_mpg:
    pushq    %rbp
    movq     %rsp, %rbp
    ...
    popq     %rbp
    ret
```

Machine
code:

```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

Computer
system:



Memory & data
Integers & floats
x86 assembly
Procedures & stacks
Executables
Arrays & structs
Memory & caches
Processes
Virtual memory
Memory allocation
Java vs. C

OS:



Data Structures in Assembly

❖ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

❖ Structs

- Alignment


~~❖ Unions~~

Array Allocation


❖ Basic Principle

- **T** A[N] ; → array of data type **T** and length N
- ★ *Contiguously* allocated region of $N * \text{sizeof}(\mathbf{T})$ bytes
- Identifier A returns address of array (type **T***)


`char msg[12];`




`int val[5];`



`double a[3];`



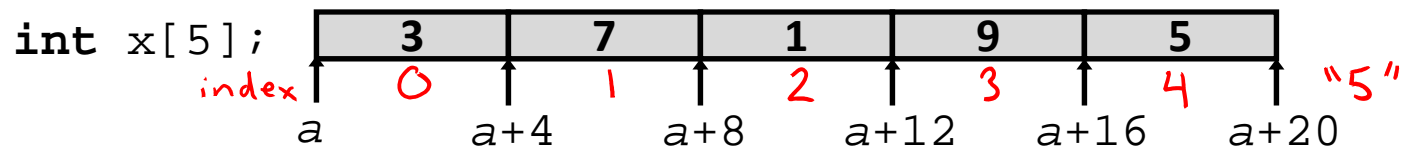
`char* p[3];`
(or `char *p[3];`)



Array Access

❖ Basic Principle

- **T** A[N] ; → array of data type **T** and length N
- Identifier A returns address of array (type **T***)



<u>Reference</u>	<u>Type</u>	<u>Value</u>
<code>x[4]</code>	<code>int</code>	5
<code>x</code>	<code>int*</code>	a
<code>x+1</code> ← ptr arithmetic	<code>int*</code>	a + 4
<code>&x[2]</code>	<code>int*</code>	a + 8
<code>x[5]</code>	<code>int</code>	?? (whatever's in memory at addr x+20)
<code>*(x+1)</code>	<code>int</code>	7
<code>x+i</code>	<code>int*</code>	a + 4*i

Array Example

```
typedef int zip_dig[5];  
zip_dig cmu = { 1, 5, 2, 1, 3 };  
zip_dig uw = { 9, 8, 1, 9, 5 };  
zip_dig ucb = { 9, 4, 7, 2, 0 };
```

typedef unsigned long int uli
old data type new, equivalent data type

initialization

- ❖ typedef: Declaration "zip_dig uw" equivalent to "int uw[5]"

Array Example

```
typedef int zip_dig[5];
```

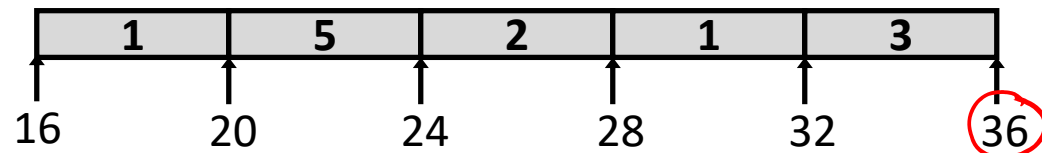
```
zip_dig cmu = { 1, 5, 2, 1, 3 };
```

```
zip_dig uw = { 9, 8, 1, 9, 5 };
```

```
zip_dig ucb = { 9, 4, 7, 2, 0 };
```


20 B each

zip_dig cmu;




no gap in this example

zip_dig uw;



zip_dig ucb;




- ❖ Example arrays happened to be allocated in successive 20 byte blocks

- Not guaranteed to happen in general

(could have allocated variables in-between)

```
typedef int zip_dig[5];
```

Array Accessing Example



```
int get_digit(zip_dig z, int digit)
{
    return z[digit];
}
```

```
get_digit:
    movl (%rdi,%rsi,4), %eax # z[digit]
```

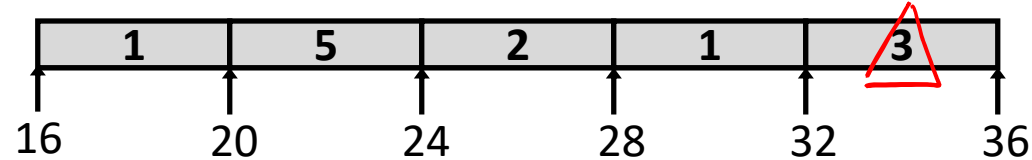
Annotations for the assembly code:

- `Rb` (base reg) points to `%rdi`
- `Ri` (index reg) points to `%rsi`
- `S: scale factor (sizeof)` points to `4`
- `size of the data type` points to `4`
- `memory operand dereferences address` points to the entire instruction `movl (%rdi,%rsi,4), %eax`

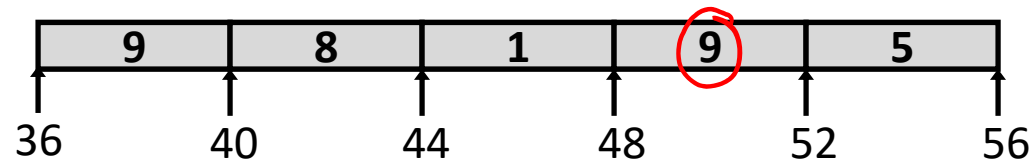
- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at `%rdi+4*%rsi`, so use memory reference `(%rdi,%rsi,4)`

Referencing Examples

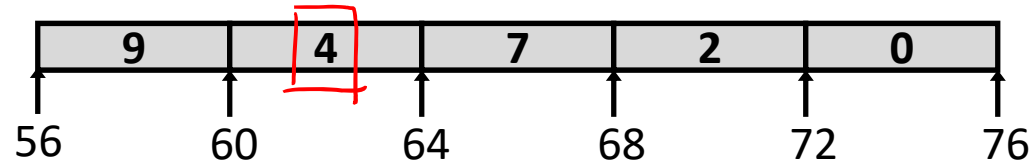
zip_dig cmu;



zip_dig uw;



zip_dig ucb;



Rb Ri S
uw 3 4

Reference	Address	Value	Guaranteed?
uw[3]	$36 + 3 * 4 = 48$	9	Yes
uw[6]	$36 + 6 * 4 = 60$	4	No
uw[-1]	$36 + (-1) * 4 = 32$	3	No
cmu[15]	$16 + 15 * 4 = 76$?	No

- ❖ No bounds checking
- ❖ Example arrays happened to be allocated in successive 20 byte blocks
 - Not guaranteed to happen in general

C Details: Arrays and Pointers

- ❖ Arrays are (almost) identical to pointers
 - `char *string` and `char string []` are nearly identical declarations
 - Differ in subtle ways: initialization, `sizeof ()`, etc.
- ❖ An array variable looks like a pointer to the first (0th) element
 - `ar[0]` same as `*ar`; `ar[2]` same as `*(ar+2)`
- ❖ An array variable is read-only (no assignment)
 - Cannot use "ar = <anything>"

C Details: Arrays and Functions

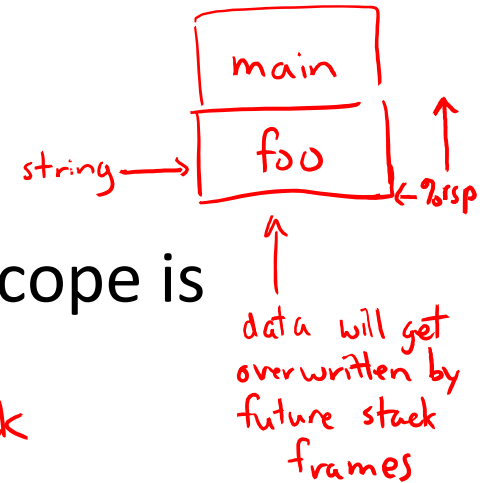
- ❖ Declared arrays only allocated while the scope is valid:

```
char* foo() {  
    char string[32]; ...;  
    return string;  
}
```

array is allocated on stack

BAD!

returns stack addr that is < %rsp



- ❖ An array is passed to a function as a pointer:
 - Array size gets lost!

```
int foo(int ar[], unsigned int size) {  
    ... ar[size-1] ...  
}
```

*Really int *ar (%rdi can only fit 8 bytes)*

Must explicitly pass the size!

Data Structures in Assembly

❖ Arrays

- One-dimensional
- **Multi-dimensional (nested)**
- Multi-level

❖ Structs

- Alignment

❖ ~~Unions~~

Nested Array Example

```
typedef int zip_dig[5];
```

```
zip_dig sea[4] =  
{ { 9, 8, 1, 9, 5 },  
  { 9, 8, 1, 0, 5 },  
  { 9, 8, 1, 0, 3 },  
  { 9, 8, 1, 1, 5 } };
```

2D array

Remember, $\mathbf{T} \ A[N]$ is
an array with elements
of type \mathbf{T} , with length N

same as:

```
int sea[4][5];
```

What is the layout in memory?

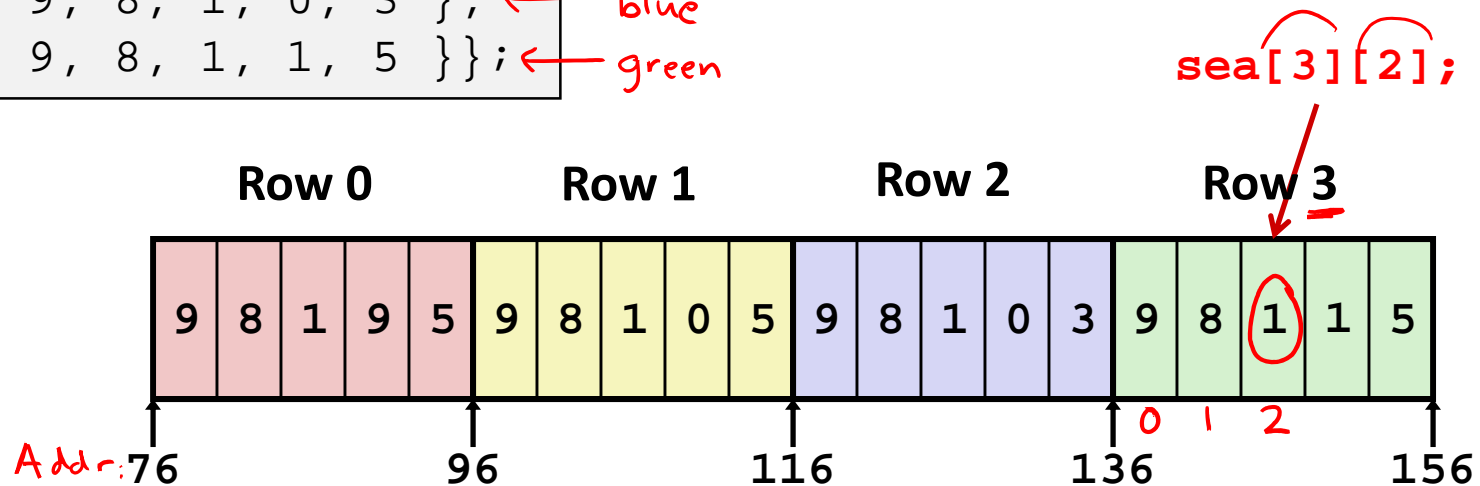
```
typedef int zip_dig[5];
```

Nested Array Example

```
zip_dig sea[4] =  
  {{ 9, 8, 1, 9, 5 },  
   { 9, 8, 1, 0, 5 },  
   { 9, 8, 1, 0, 3 },  
   { 9, 8, 1, 1, 5 }};
```

← red
← yellow
← blue
← green

Remember, $\mathbf{T} \ A[N]$ is
an array with elements
of type \mathbf{T} , with length N



- ❖ “Row-major” ordering of all elements
- ❖ Elements in the same row are contiguous
- ❖ Guaranteed (in C)

Two-Dimensional (Nested) Arrays

❖ Declaration: ~~T~~ ^{int} A[⁴~~R~~][⁵~~C~~];

- 2D array of data type T
- R rows, C columns
- Each element requires `sizeof(T)` bytes

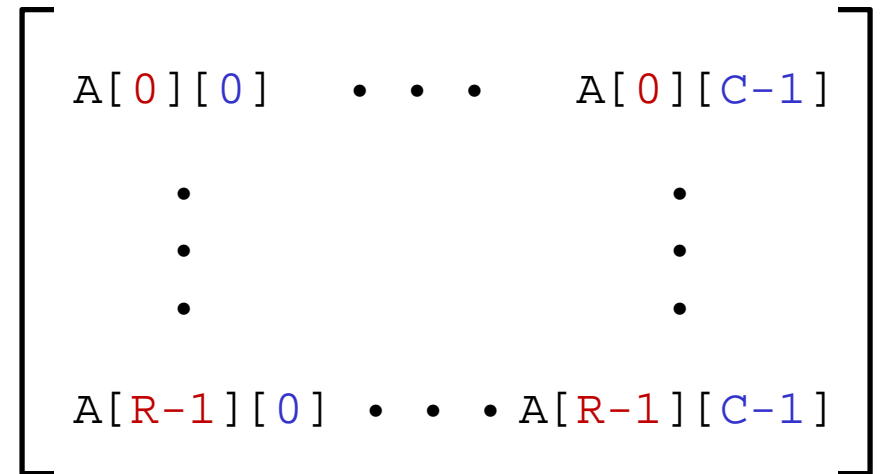
$$\begin{bmatrix} A[\underline{0}][\underline{0}] & \cdot & \cdot & \cdot & A[\underline{0}][\overline{C-1}] \\ \cdot & & & & \cdot \\ \cdot & & & & \cdot \\ \cdot & & & & \cdot \\ A[\underline{R-1}][\underline{0}] & \cdot & \cdot & \cdot & A[\underline{R-1}][\overline{C-1}] \end{bmatrix}$$

❖ Array size?

Two-Dimensional (Nested) Arrays

❖ Declaration: **T** A[R][C];

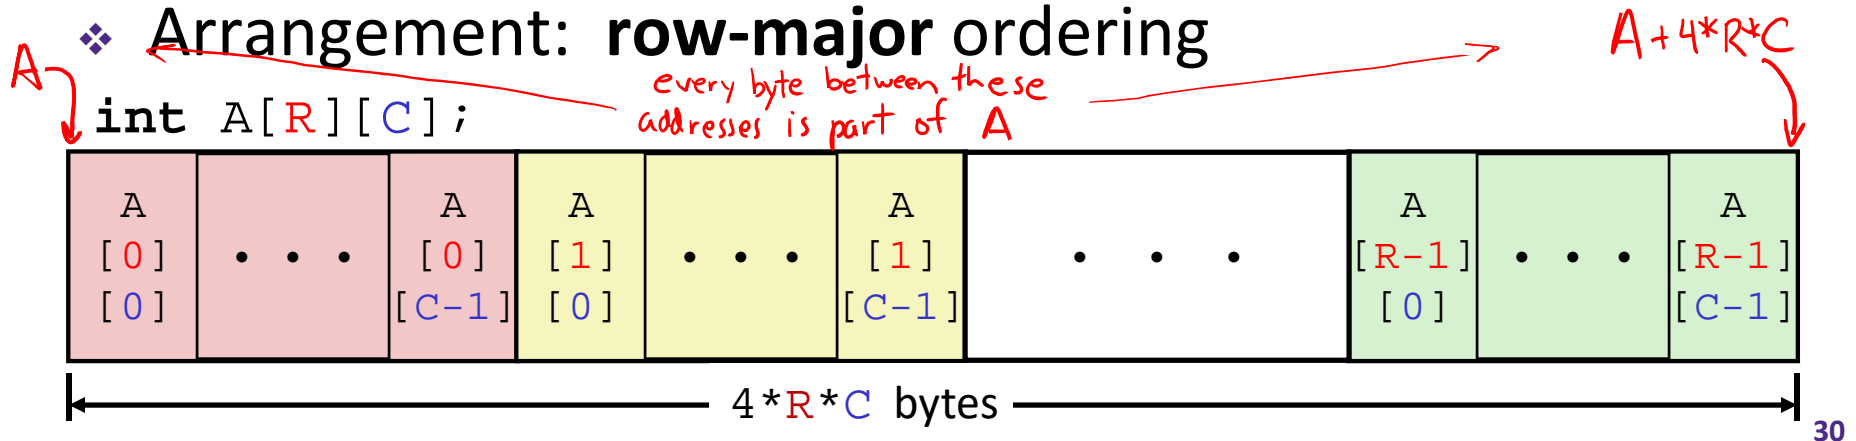
- 2D array of data type T
- R rows, C columns
- Each element requires **sizeof(T)** bytes



❖ Array size:

- $R * C * \text{sizeof}(T)$ bytes

❖ Arrangement: **row-major** ordering

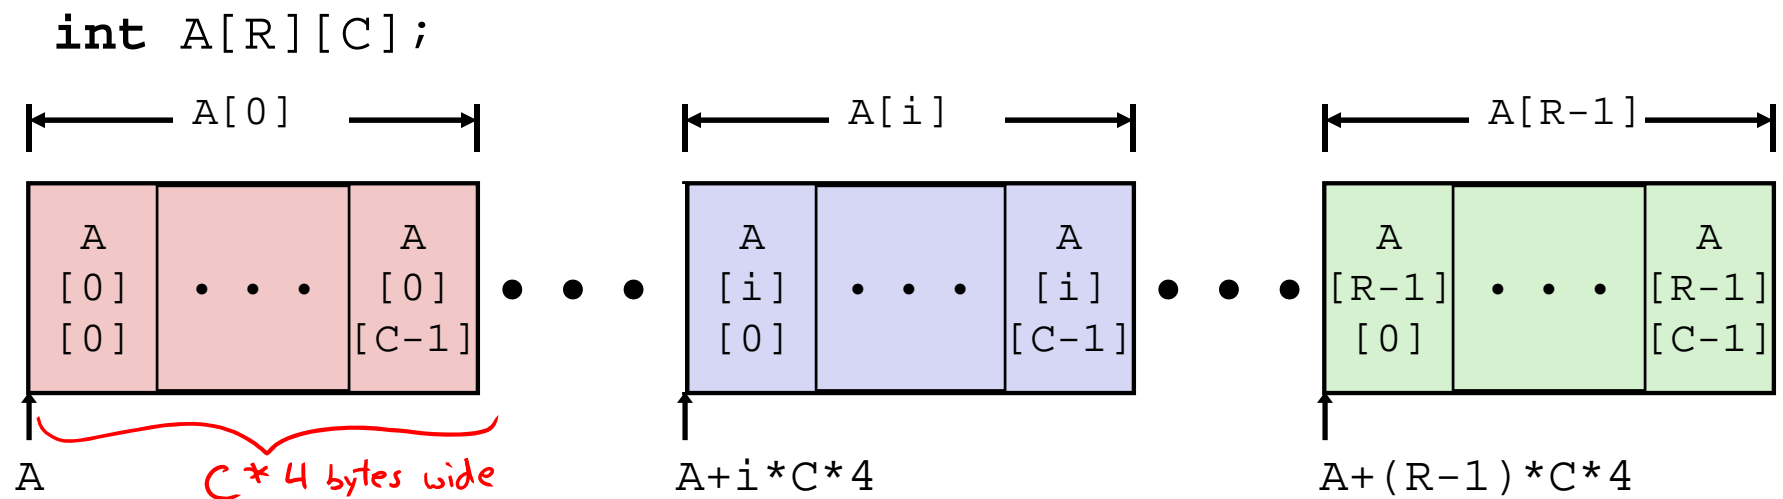


Nested Array Row Access

❖ Row vectors

■ Given $\mathbf{T} \ A[R][C]$,

- $A[i]$ is an array of C elements ("row i ") \rightarrow just an address!
- A is address of array
- Starting address of row $i = A + i * (C * \text{sizeof}(\mathbf{T}))$



Nested Array Row Access Code

```
int* get_sea_zip(int index)
{
    return sea[index];
}
```

```
int sea[4][5] =
{{ 9, 8, 1, 9, 5 },
 { 9, 8, 1, 0, 5 },
 { 9, 8, 1, 0, 3 },
 { 9, 8, 1, 1, 5 }};
```

```
get_sea_zip(int):
    movslq    %edi, %rdi
    leaq      (%rdi,%rdi,4), %rax
    leaq      sea(,%rax,4), %rax
    ret
```

address of array → sea:

```
.long 9
.long 8
.long 1
.long 9
.long 5
.long 9
.long 8
```

...

ends up in memory!

Nested Array Row Access Code

```
int* get_sea_zip(int index)
{
    return sea[index];
}
```

R C

```
int sea[4][5] =
    {{ 9, 8, 1, 9, 5 },
     { 9, 8, 1, 0, 5 },
     { 9, 8, 1, 0, 3 },
     { 9, 8, 1, 1, 5 }};
```

- What data type is `sea[index]`? address
- What is its value? $A + C * \text{sizeof}(T) * i \rightarrow \text{sea} + 5 * 4 * \text{index}$

```
# %rdi = index
leaq (%rdi,%rdi,4),%rax
leaq sea(,%rax,4),%rax
```

Translation?

using a label as D

Nested Array Row Access Code

```
int* get_sea_zip(int index)
{
    return sea[index];
}
```

```
int sea[4][5] =
    {{ 9, 8, 1, 9, 5 },
     { 9, 8, 1, 0, 5 },
     { 9, 8, 1, 0, 3 },
     { 9, 8, 1, 1, 5 }};
```

```
# %rdi = index
leaq (%rdi,%rdi,4),%rax # 5 * index
leaq sea(,%rax,4),%rax  # sea + (20 * index)
```

just calculating an address, so no memory access

❖ Row Vector

- `sea[index]` is array of 5 ints
- Starting address = `sea+20*index`

❖ Assembly Code

- Computes and returns address
- Compute as: `sea+4*(index+4*index) = sea+20*index`


Nested Array Element Access

reminder: $ar[j] = *(ar + j)$

❖ Array Elements

- $A[i][j]$ is element of type T , which requires K bytes
- Address of $(A[i])[j]$ is $(A + i * C * \text{sizeof}(T)) + j * \text{sizeof}(T)$
address

```
int A[R][C];
```



Nested Array Element Access


❖ Array Elements

- $A[i][j]$ is element of type \mathbf{T} , which requires K bytes

- Address of $A[i][j]$ is

$$A + i*(C*K) + j*K == A + (i*C + j)*K$$

```
int A[R][C];
```



Nested Array Element Access Code

```
int get_sea_digit
(int index, int digit)
{
    return sea[index][digit];
}
```

```
int sea[4][5] =
{{ 9, 8, 1, 9, 5 },
 { 9, 8, 1, 0, 5 },
 { 9, 8, 1, 0, 3 },
 { 9, 8, 1, 1, 5 }};
```

```
leaq    (%rdi,%rdi,4), %rax    # 5*index
->addl   %rax, %rsi            # 5*index+digit
movl    sea(,%rsi,4), %eax    # *(sea + 4*(5*index+digit))
```

mov gets data

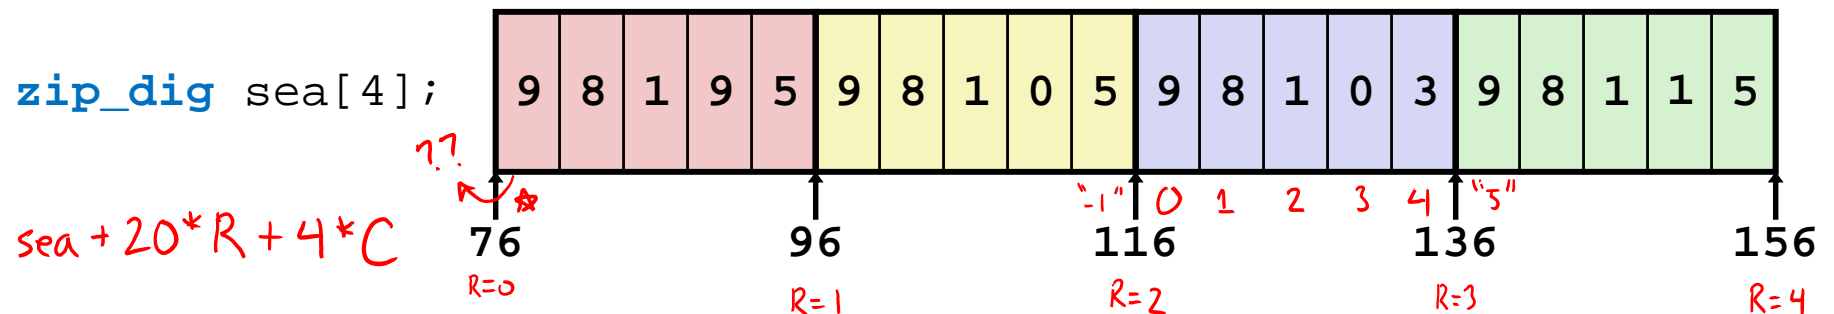
❖ Array Elements

- `sea[index][digit]` is an **int** (**sizeof(int)=4**)
- $\text{Address} = \text{sea} + 5 \cdot 4 \cdot \text{index} + 4 \cdot \text{digit}$
start of array *start of row* *column of interest*

❖ Assembly Code

- Computes address as: $\text{sea} + ((\text{index} + 4 \cdot \text{index}) + \text{digit}) \cdot 4$
- `movl` performs memory reference

Multi-Dimensional Referencing Examples



Reference	Address	Value	Guaranteed?
<code>sea[3][3]</code>	$76 + 20 * 3 + 4 * 3 = 148$	1	Yes
<code>sea[2][5]</code>	$76 + 20 * 2 + 4 * 5 = 136$	9	Yes
<code>sea[2][-1]</code>	$76 + 20 * 2 + 4 * (-1) = 112$	5	Yes
<code>sea[4][-1]</code>	$76 + 20 * 4 + 4 * (-1) = 152$	5	Yes
<code>sea[0][19]</code>	$76 + 20 * 0 + 4 * (19) = 152$	5	Yes
<code>sea[0][-1]</code>	$76 + 20 * 0 + 4 * (-1) = 72$??	No

- Code does not do any bounds checking
- Ordering of elements within array guaranteed

Data Structures in Assembly

❖ Arrays

- One-dimensional
- Multi-dimensional (nested)
- **Multi-level**

❖ Structs

- Alignment

~~❖ Unions~~

Multi-Level Array Example

Multi-Level Array Declaration(s):

```
int cmu[5] = { 1, 5, 2, 1, 3 };
int uw[5] = { 9, 8, 1, 9, 5 };
int ucb[5] = { 9, 4, 7, 2, 0 };
```

} could be
apart

4 arrays

```
int* univ[3] = {uw, cmu, ucb};
```

$univ[2][2] == 7$

Is a multi-level array the
same thing as a 2D array?

NO

2D Array Declaration:

```
zip_dig univ2D[3] = {
    { 9, 8, 1, 9, 5 },
    { 1, 5, 2, 1, 3 },
    { 9, 4, 7, 2, 0 }
};
```

} guaranteed
contiguous

$univ2D[2][2] == 7$

← 1 array

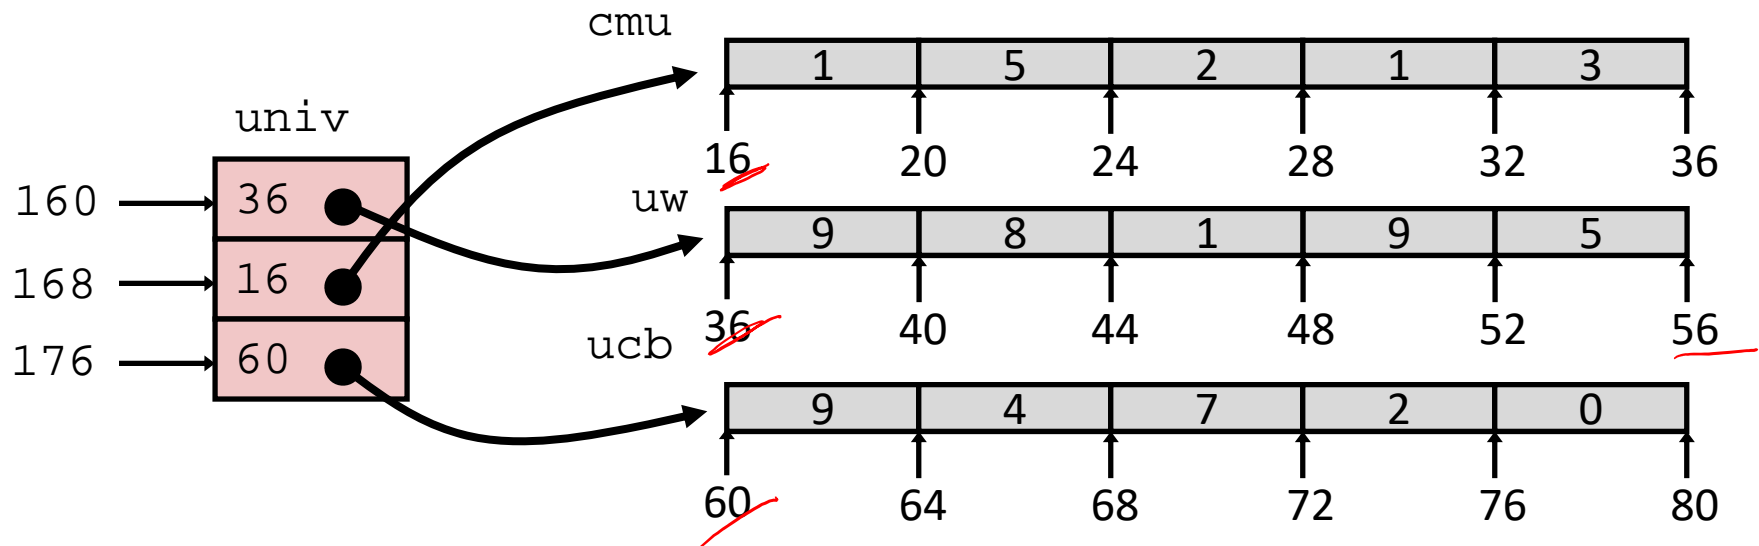
One array declaration = one contiguous block of memory

Multi-Level Array Example

```
int cmu[5] = { 1, 5, 2, 1, 3 };  
int uw[5] = { 9, 8, 1, 9, 5 };  
int ucb[5] = { 9, 4, 7, 2, 0 };
```

```
int* univ[3] = {uw, cmu, ucb};
```

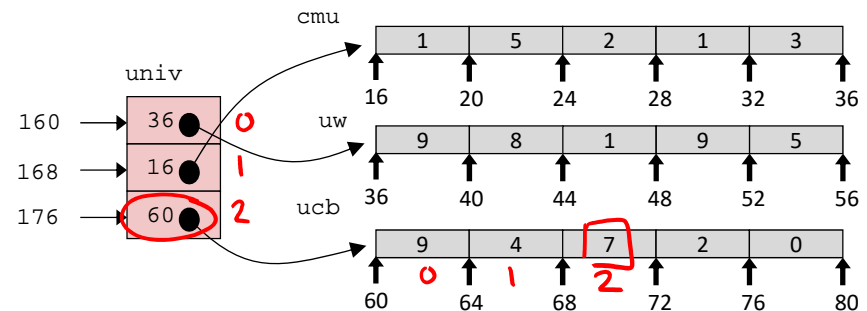
- ❖ Variable `univ` denotes array of 3 elements
- ❖ Each element is a pointer
 - 8 bytes each
- ❖ Each pointer points to array of `ints`



Note: this is how Java represents multi-dimensional arrays

Element Access in Multi-Level Array

```
int get_univ_digit
(int index, int digit)
{
    return univ[index][digit];
}
```



```
salq    $2, %rsi          # rsi = 4*digit
addq    univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax       # return *p
ret
```

❖ Computation

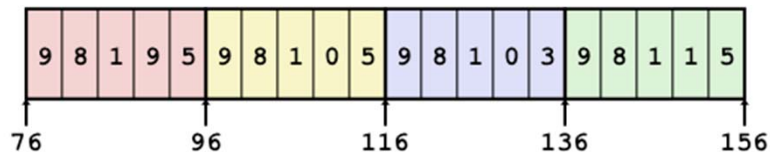
- Element access $\text{Mem}[\text{Mem}[\text{univ} + 8 * \text{index}] + 4 * \text{digit}]$
- Must do **two memory reads**
 - First get pointer to row array
 - Then access element within array
- But allows inner arrays to be different lengths (not in this example)

• also easier to "fit" smaller arrays in memory
 • also can "swap out" rows in multi-level

Array Element Accesses

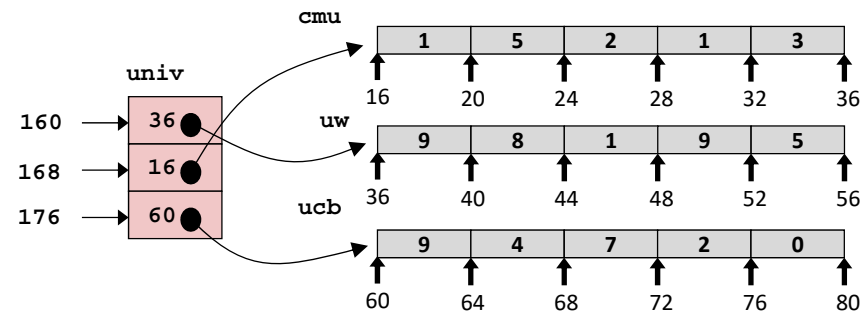
Nested array

```
int get_sea_digit
(int index, int digit)
{
    return sea[index][digit];
}
```



Multi-level array

```
int get_univ_digit
(int index, int digit)
{
    return univ[index][digit];
}
```

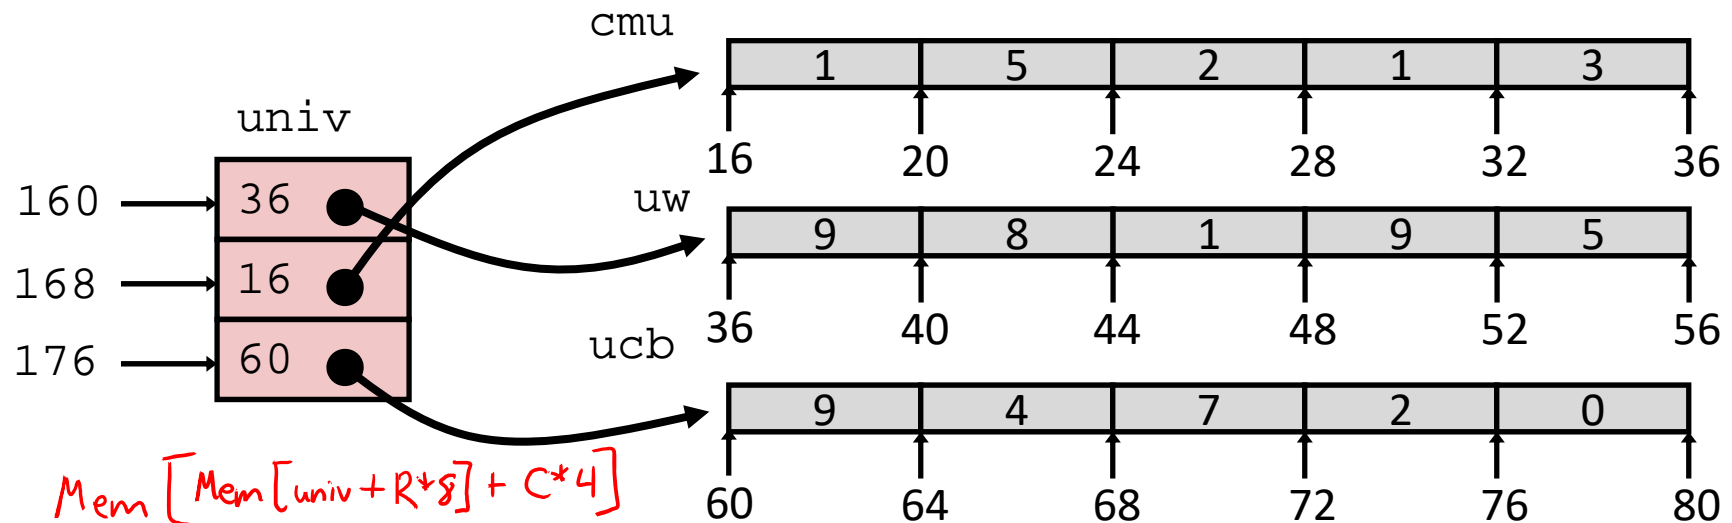


Access *looks* the same, but it isn't:

Mem[sea+20*index+4*digit]

Mem[**Mem**[univ+8*index]+4*digit]

Multi-Level Referencing Examples



Reference	Address	Value	Guaranteed?
<code>univ[2][3]</code>	$Mem[176] + 3 * 4 = 60 + 12 = 72$	2	Yes
<code>univ[1][5]</code>	$Mem[168] + 5 * 4 = 16 + 20 = 36$	9	No
<code>univ[2][-2]</code>	$Mem[176] + (-2) * 4 = 60 - 8 = 52$	5	No
<code>univ[3][-1]</code>	$Mem[184] + (-1) * 4 = ?? - 4 = ??$???	No
<code>univ[1][12]</code>	$Mem[168] + 12 * 4 = 16 + 48 = 64$	4	No

- C code does not do any bounds checking
- Location of each lower-level array in memory is *not* guaranteed

Summary

- ❖ Contiguous allocations of memory
- ❖ **No bounds checking** (and no default initialization)
- ❖ Can usually be treated like a pointer to first element
- ❖ **int** a[4][5]; → array of arrays
 - all levels in one contiguous block of memory
- ❖ **int*** b[4]; → array of pointers to arrays
 - First level in one contiguous block of memory
 - Each element in the first level points to another “sub” array
 - Parts anywhere in memory