

Structs & Alignment

CSE 351 Spring 2019

Instructor:

Ruth Anderson

Teaching Assistants:

Gavin Cai

Jack Eggleston

John Feltrup

Britt Henderson

Richard Jiang

Jack Skalitzy

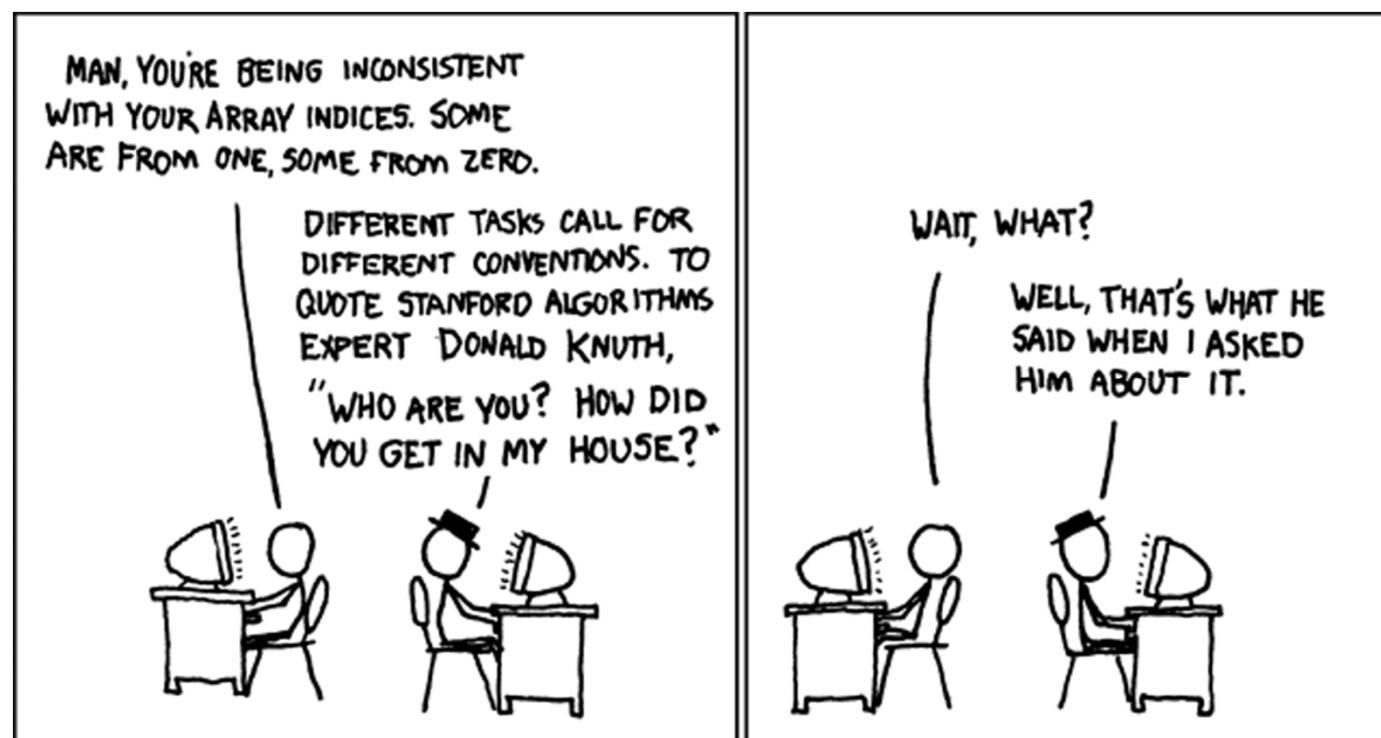
Sophie Tian

Connie Wang

Sam Wolfson

Casey Xing

Chin Yeoh



<http://xkcd.com/163/>

Administrivia

- ❖ Lab 2 (x86-64) due TONIGHT (5/01)
- ❖ Homework 3, due Wednesday (5/8)
 - On midterm material, but due after the midterm
- ❖ **Midterm** (Fri 5/03, 4:30-5:30pm in KNE 130)
 - Review Session: Thurs: 6:30-8:30pm in Sieg 134
 - No lecture on Friday 5/03
 - Ruth will hold office hours instead
 - Fri 11:30am-12:30pm in CSE 460
 - Fri 2:30-3:30pm in CSE 460

Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

Memory & data
Integers & floats
x86 assembly
Procedures & stacks
Executables
Arrays & structs
Memory & caches
Processes
Virtual memory
Memory allocation
Java vs. C

Assembly
language:

```
get_mpg:
    pushq    %rbp
    movq     %rsp, %rbp
    ...
    popq     %rbp
    ret
```


Machine
code:

```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

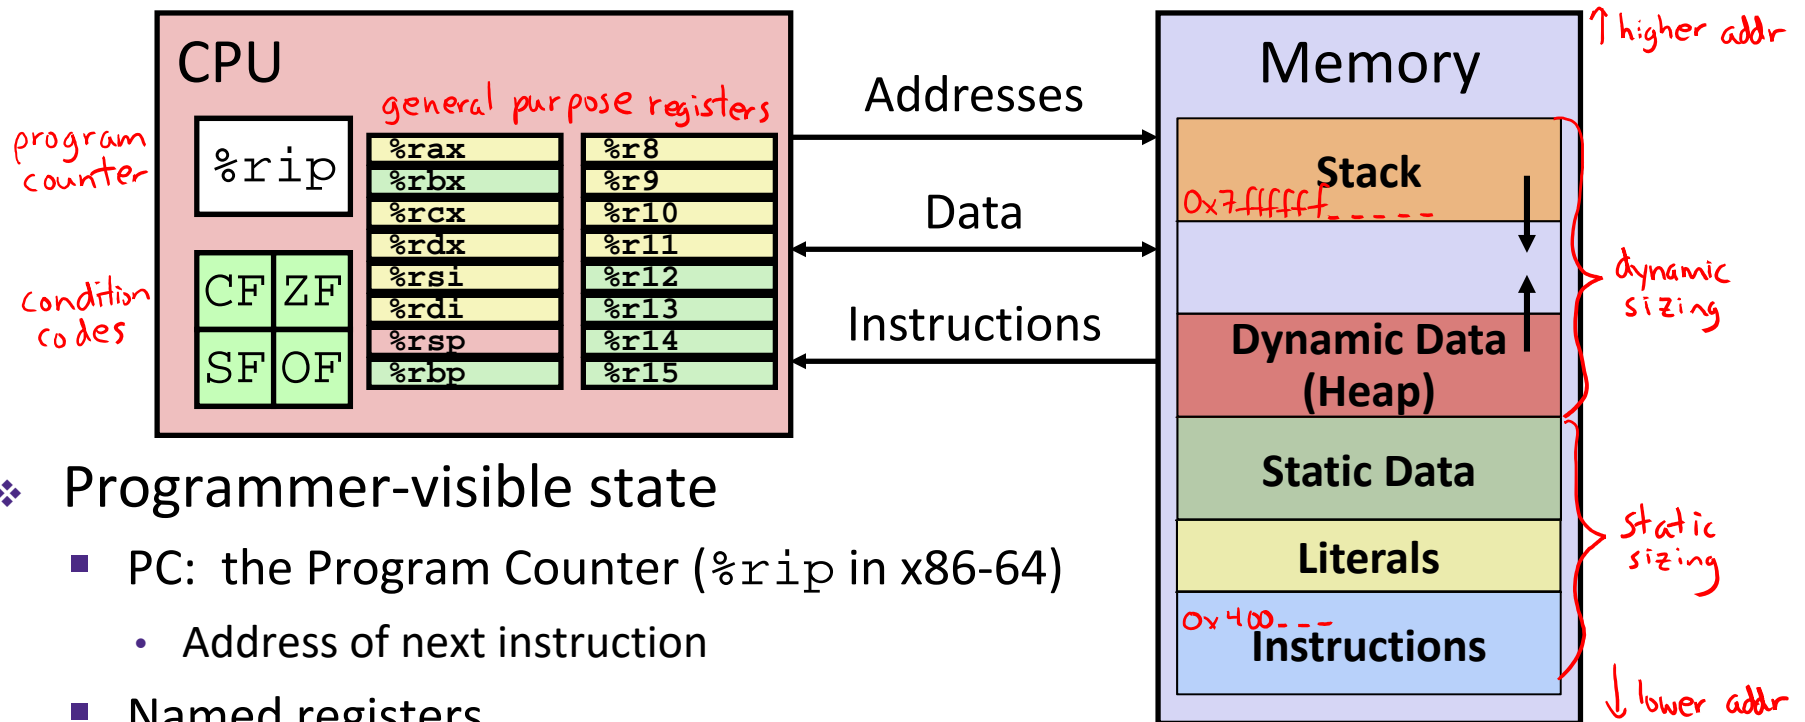
OS:



Computer
system:



Assembly Programmer's View



❖ Programmer-visible state

- **PC: the Program Counter (`%rip` in x86-64)**
 - Address of next instruction
- **Named registers**
 - Together in “register file”
 - Heavily used program data
- **Condition codes**
 - Store status information about most recent arithmetic operation
 - Used for conditional branching

❖ Memory

- Byte-addressable array
- Code and user data
- Includes *the Stack* (for supporting procedures)

x86-64 Instructions

size specifiers: b, w, l, q
1, 2, 4, 8 bytes

① ❖ Data movement

- mov, movs, movz, ...

operand types: Imm \$
Reg %
Mem ()

② ❖ Arithmetic

- add, sub, shl, sar, lea, ...

Labels are addresses

❖ Control flow

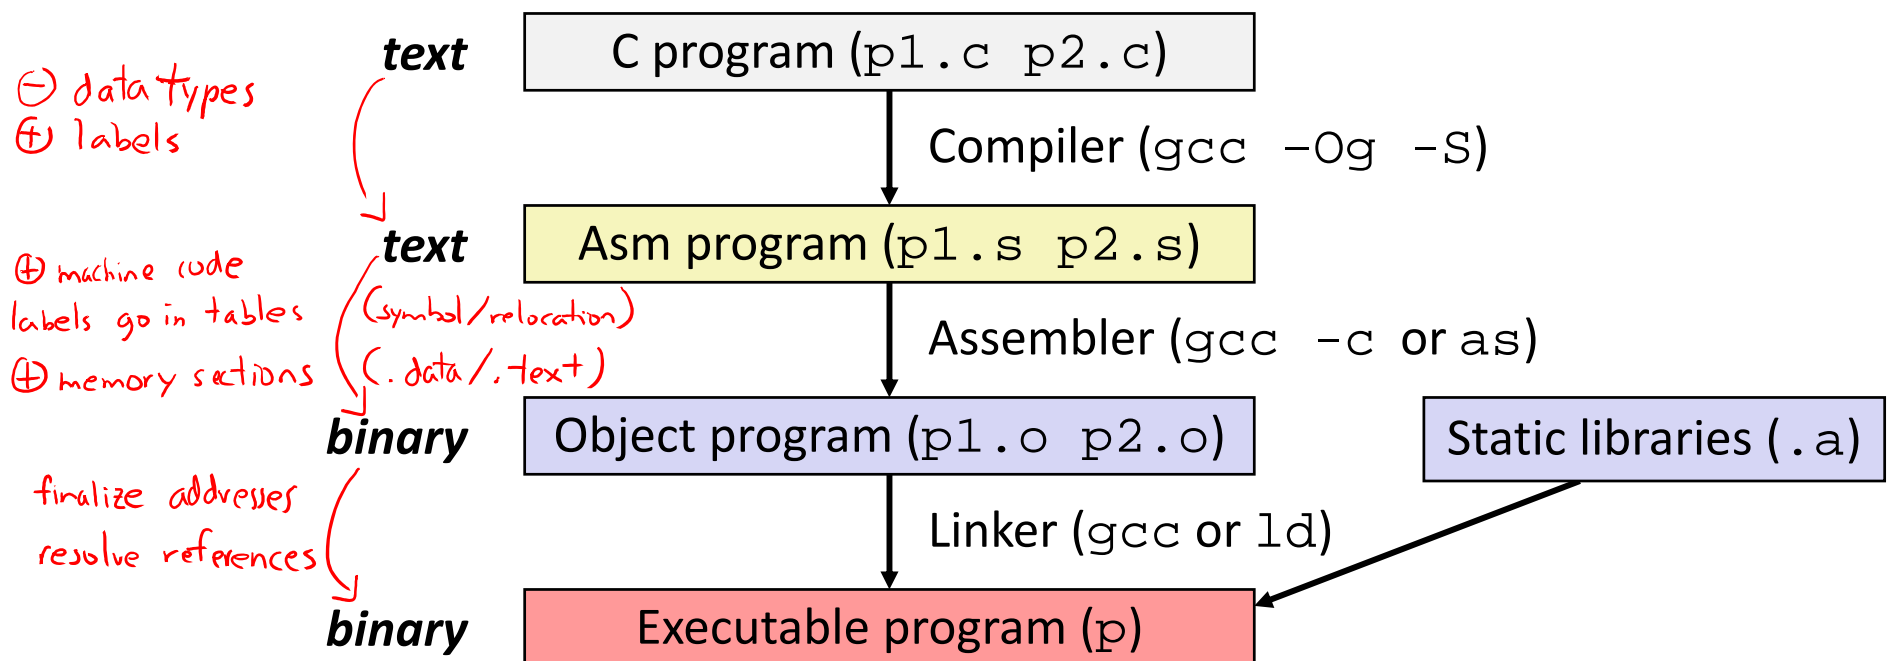
- cmp, test, j*, set*, ...

❖ Stack/procedures

- push, pop, call, ret, ...

Turning C into Object Code

- ❖ Code in files `p1.c` `p2.c`
- ❖ Compile with command: `gcc -Og p1.c p2.c -o p`
 - Use basic optimizations (`-Og`) [New to recent versions of GCC]
 - Put resulting machine code in file `p`



Assembling

- ❖ Executable has **addresses** (no more labels)

assembler →

```

00000000004004f6 <pcount_r>:
4004f6:  b8 00 00 00 00    mov     $0x0,%eax
4004fb:  48 85 ff          test    %rdi,%rdi
4004fe:  74 13            je      400513 <pcount_r+0x1d>
400500:  53              push    %rbx
400501:  48 89 fb          mov     %rdi,%rbx
400504:  48 d1 ef          shr     %rdi
400507:  e8 ea ff ff ff    callq   4004f6 <pcount_r>
40050c:  83 e3 01          and     $0x1,%ebx
40050f:  48 01 d8          add     %rbx,%rax
400512:  5b              pop     %rbx
400513:  f3 c3            rep     ret
  
```

used to be a label
(Exit: or .Lb:) →

pcount_r + 0x1d = 30 bytes after start of pcount_r

- gcc -g pcount.c -o pcount
- objdump -d pcount

A Picture of Memory (64-bit view)

00000000004004f6 <pcount_r>:

4004f6: **b8** 00 00 00 00

4004fb: 48 85 ff

4004fe: 74 13

400500: 53

400501: 48 89 fb

400504: 48 d1 ef

400507: e8 ea ff ff ff

40050c: 83 e3 01

40050f: 48 01 d8

400512: 5b

400513: f3 c3

mov \$0x0,%eax

test %rdi,%rdi

je 400513 <pcount_r+0x1d>

push %rbx

mov %rdi,%rbx

shr %rdi

callq 4004f6 <pcount_r>

and \$0x1,%ebx

add %rbx,%rax

pop %rbx

rep ret

instruction
addresses

stored bytes

unaligned, but
more compact

0|8 1|9 2|a 3|b 4|c 5|d 6|e 7|f

...							
						b8	00
00	00	00	48	85	ff	74	13
53	48	89	fb	48	d1	ef	e8
ea	ff	ff	ff	83	e3	01	48
01	d8	5b	f3	c3			

0x00

0x08

0x10

...

0x4004f0

0x4004f8

0x400500

0x400508

0x400510

Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```


Assembly
language:

```
get_mpg:
    pushq    %rbp
    movq     %rsp, %rbp
    ...
    popq     %rbp
    ret
```

Machine
code:

```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

Computer
system:



Memory & data
Integers & floats
x86 assembly
Procedures & stacks
Executables
Arrays & structs
Memory & caches
Processes
Virtual memory
Memory allocation
Java vs. C


OS:



Peer Instruction Question

❖ Which of the following statements is **FALSE**?

■ Vote at <http://pollev.com/rea>



A. `sea[4][-2]` is a *valid* array reference

Yes, returns 1

B. `sea[1][1]` makes *two* memory accesses

No, only single memory access

C. `sea[2][1]` will *always* be a higher address than `sea[1][2]`

Yes, because C is row-major

D. `sea[2]` is calculated using *only* `lea`

Yes, `sea[2]` returns address of array row

E. We're lost...

Data Structures in Assembly

❖ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

❖ Structs


- Alignment

❖ ~~Unions~~

Structs in C

- ❖ Way of defining compound data types
- ❖ A structured group of variables, possibly including other structs

```
typedef struct {  
    int lengthInSeconds;  
    int yearRecorded;  
} Song;  
  
Song song1;  
  
song1.lengthInSeconds = 213;  
song1.yearRecorded    = 1994;  
  
Song song2;  
  
song2.lengthInSeconds = 248;  
song2.yearRecorded    = 1988;
```



Struct Definitions

typedef *unsigned long int* *uli;*
 old var type *new name*

❖ Structure definition:

- Does NOT declare a variable
- Variable type is "struct name"

your choice

```
struct name {
    /* fields */
};
```

Easy to forget
semicolon!

```
struct name name1, *pn, name_ar[3];
```

pointer

instance

array

❖ Joint struct definition and typedef

- Don't need to give struct a name in this case

① define struct

```
struct nm {
    /* fields */
};
```

② typedef

```
typedef struct nm name;
name n1;
```


combined

unnamed!

```
typedef struct {
    /* fields */
} name;
name n1;
```

Scope of Struct Definition

- ❖ Why is placement of struct definition important?
 - What actually happens when you declare a variable?
 - Creating space for it somewhere!
 - Without definition, program doesn't know how much space



- ❖ Almost always define structs in global scope near the top of your C file
 - Struct definitions follow normal rules of scope

Accessing Structure Members

- ❖ Given a struct instance, access member using the `.` operator:

```
struct rec r1;  
r1.i = val;
```

- ❖ Given a *pointer* to a struct:

```
struct rec *r;
```

`r = &r1;` // or malloc space for `r` to point to

We have two options:

- Use `*` and `.` operators: `(*r).i = val;`
- Use `->` operator for short: `r->i = val;`

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```

① dereference (get instance)

② access field

equivalent

- ❖ **In assembly:** register holds address of the first byte

- Access members with offsets

$D(Rb, Ri, S)$

Java side-note

```
class Record { ... }  
Record x = new Record( );
```

- ❖ An instance of a class is like a *pointer to* a struct containing the fields
 - (Ignoring methods and subclassing for now)
 - So Java's $x.f$ is like C's $\underline{x} \rightarrow f$ or $(*x).f$
- ❖ In Java, almost everything is a pointer ("*reference*") to an object
 - Cannot declare variables or fields that are structs or arrays
 - Always a *pointer* to a struct or array
 - So every Java variable or field is ≤ 8 bytes (but can point to lots of data)


Structure Representation

struct definition

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
}
```

declare a pointer

```
*r;
```




❖ Characteristics

- Contiguously-allocated region of memory
- Refer to members within structure by names
- Members may be of different types

Structure Representation


```
struct rec {  
  ① int a[4];  
  ② long i;  
  ③ struct rec *next;  
} *r;
```



- ❖ Structure represented as block of memory
 - Big enough to hold all of the fields
- ❖ ~~Fields ordered according to declaration order~~
 - Even if another ordering would be more compact
- ❖ Compiler determines overall size + positions of fields
 - Machine-level program has no understanding of the structures in the source code

Accessing a Structure Member

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



- ❖ Compiler knows the *offset* of each member within a struct


- Compute as
 - $*(r + \text{offset})$
 - Referring to absolute offset, so no pointer arithmetic

```
long get_i(struct rec *r)  
{  
    return r->i;  
}
```

```
# r in %rdi, index in %rsi  
movq 16(%rdi), %rax  
ret
```

Exercise: Pointer to Structure Member

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



```
long* addr_of_i(struct rec *r)  
{  
    return &(r->i);  
}
```


```
# r in %rdi  
leaq    16(%rdi), %rax  
ret
```

```
struct rec** addr_of_next(struct rec *r)  
{  
    return &(r->next);  
}
```

```
# r in %rdi  
leaq    24(%rdi), %rax  
ret
```

Generating Pointer to Array Element

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



❖ Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Compute as:
 $r + 4 * \text{index}$

```
int* find_addr_of_array_elem  
    (struct rec *r, long index)  
{  
    return &r->a[index];  
}
```

\searrow
`&(r->a[index])`

```
# r in %rdi, index in %rsi  
leaq (%rdi,%rsi,4), %rax  
ret
```

Review: Memory Alignment in x86-64

- ❖ *Aligned* means that any primitive object of K bytes must have an address that is a multiple of K
- ❖ Aligned addresses for data types:

K	Type	Addresses
1	char	No restrictions
2	short	Lowest bit must be zero: $\dots 0_2$
4	int, float	Lowest 2 bits zero: $\dots \underline{00}_2$
8	long, double, *	Lowest 3 bits zero: $\dots 000_2$
16	long double	Lowest 4 bits zero: $\dots 0000_2$

lowest $\log_2(K)$
bits should be 0

"multiple of" means no remainder when you divide by.
since K is a power of 2, dividing by K is equivalent to $\gg \log_2(K)$.
No remainder means no weight is "lost" during the shift \rightarrow all zeros in lowest $\log_2(K)$ bits.

Alignment Principles

❖ Aligned Data


- Primitive data type requires K bytes
- Address must be multiple of K
- Required on some machines; advised on x86-64

❖ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of bytes (width is system dependent)
 - Inefficient to load or store value that spans quad word boundaries
 - Virtual memory trickier when value spans 2 pages (more on this later)
- Though x86-64 hardware will work regardless of alignment of data just possibly slower because we have to load/read more data

Structures & Alignment

❖ Unaligned Data



```

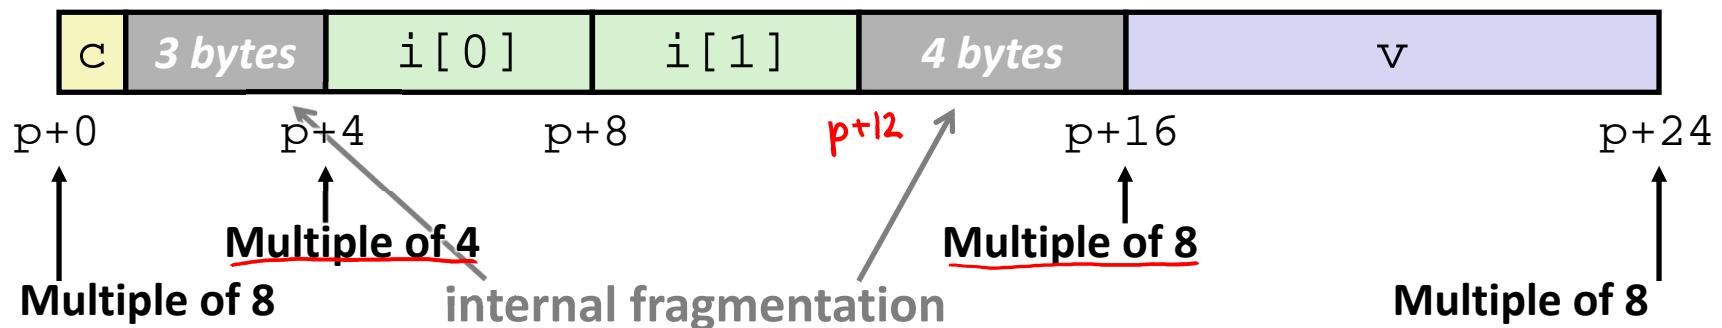
struct S1 {
  ① char c;      ←  $\frac{K}{1}$ 
  ② int i[2];    ← 4
  ③ double v;    ← 8
} *p;

```

❖ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K

24 B total



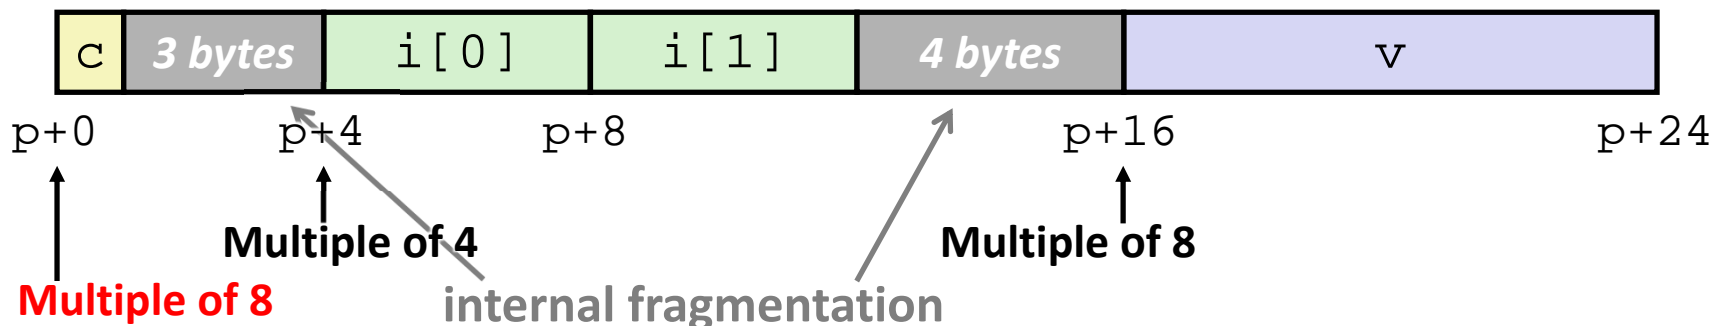
Satisfying Alignment with Structures (1)

- ❖ Within structure:
 - Must satisfy each element's alignment requirement
- ❖ Overall structure placement
 - Each structure has alignment requirement K_{\max}
 - K_{\max} = Largest alignment of any element
 - Counts array elements individually as elements

```
struct s1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

K
1
4
8
 $K_{\max} = 8$

- ❖ Example:
 - $K_{\max} = 8$, due to double element



Satisfying Alignment with Structures (2)


- ❖ Can find offset of individual fields using `offsetof()`
 - Need to `#include <stddef.h>`
 - Example: `offsetof(struct S2, c)` returns 16

```

struct S2 {
    double v;
    int i[2];
    char c;
} *p;

```


- ❖ For largest alignment requirement K_{\max} , **overall structure size must be multiple of K_{\max}**
 - Compiler will add padding **at end** of structure to meet overall structure alignment requirement



Arrays of Structures

- ❖ Overall structure length multiple of K_{max}
- ❖ Satisfy alignment requirement for every element in array

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```

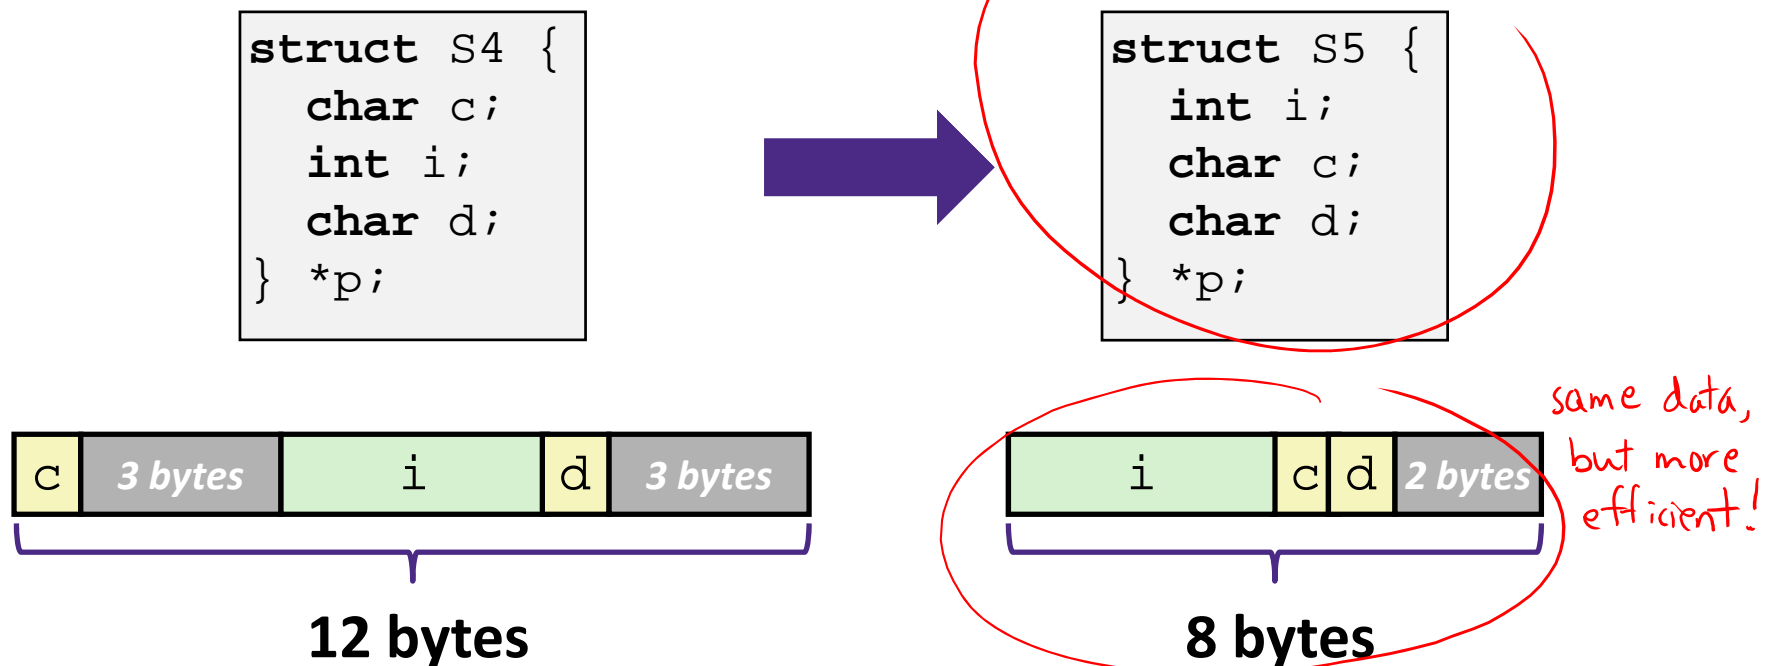


Alignment of Structs

- ❖ Compiler will do the following:
 - Maintains declared *ordering* of fields in struct
 - Each **field** must be aligned *within* the struct (*may insert padding*)
 - `offsetof` can be used to get actual field offset
 - Overall struct must be **aligned** according to largest field
 - Total struct **size** must be multiple of its alignment (*may insert padding*)
 - `sizeof` should be used to get true size of structs
- so that if we just to store an array of structures, they are all aligned (e.g. so we can store the array elements contiguously in memory)

How the Programmer Can Save Space

- ❖ Compiler must respect order elements are declared in
 - Sometimes the programmer can save space by declaring large data types first

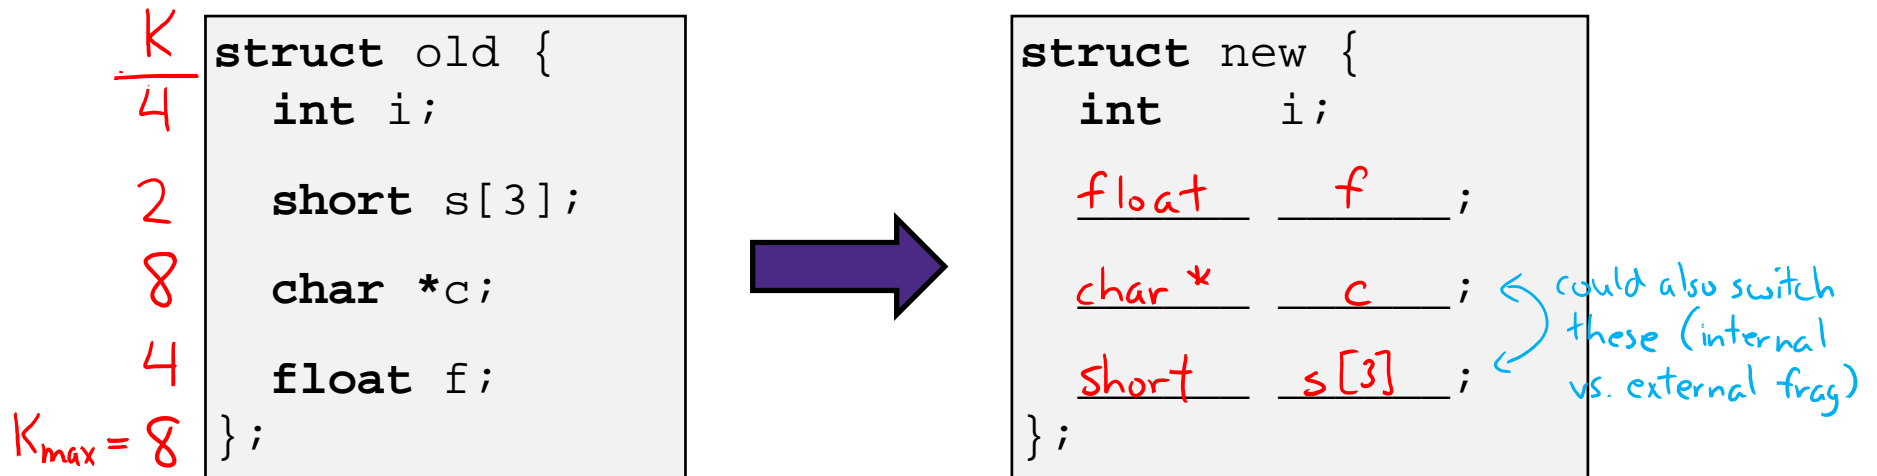


Peer Instruction Question

Vote on sizeof(struct old):

<http://pollev.com/rea>

- ❖ Minimize the size of the struct by re-ordering the vars



- ❖ What are the old and new sizes of the struct?

sizeof(struct old) = 32 B

sizeof(struct new) = 24 B


A. 16 bytes

B. 22 bytes

C. 28 bytes

D. 32 bytes

E. We're lost...



Summary

- ❖ Arrays in C
 - Aligned to satisfy every element's alignment requirement
- ❖ Structures
 - Allocate bytes in order declared
 - Pad in middle and at end to satisfy alignment